



Memory Saving Discrete Fourier Transformation on CUDA

Daniel Kauker Steffen Frey Harald Sanftmann

`kaukerdl@studi.informatik.uni-stuttgart.de`

`steffen.frey@vis.uni-stuttgart.de`

`harald.sanftmann@vis.uni-stuttgart.de`

Visualization and Interactive Systems Group

October 1, 2009



Motivation

- Fourier transformation is often used in image and video processing software (filtering, etc).
- Fast implementations available for graphics hardware, e.g. NVidias CUFFT library.
- Bring this technique to users, who might not have the latest high end graphics cards
- Problem: the created CUFFT-„Plans“ consume too much memory. Large images cannot be processed, as the operating system and other applications also consume video memory.
Example for 8 Megapixel image:
32 Megabytes data + ca. 121 Megabytes plan
but only ca. 120 Megabytes of memory available
on a 256 Megabytes device. (4-byte-float, one channel)



Table of Contents

1. Fourier Transformation
2. Flexible Fourier Transformation
3. Results



Fourier Transformation

Standard two-dimensional Fourier transformations:

- Discrete Fourier Transformation: $O(N^4)$
- Fast Fourier Transformation: $O(N^2 \cdot \log(N^2))$

for a $N \times N$ image.

Use separability to calculate the Fourier transformation for 2d images in $2 \cdot N$ 1d transformations.



Existing Libraries

- CPU
 - OpenCV:
Open source computer vision framework
 - Intel Performance Primitives:
Commercial, enhance Intel CPU architecture
 - libfftw:
Open source library
implementing sophisticated Fourier transformation algorithms
 - And several others
- Graphic Cards
 - NVidia's CUFFT:
Fast, but very memory consuming
 - GPUFFTW:
Power-of-Two FFT Library from UNC



Previous Work on GPU

- Basic Shader Model FFT implementations
[Kenneth Moreland and Edward Angel, 2003]
- FFT algorithms for CUDA aiming at efficiently exploiting shared memory
[Naga K. Govindaraju et al., 2008]
- ...

In this presentation:

Use CUFFT but exploit separability of the Fourier transformation.

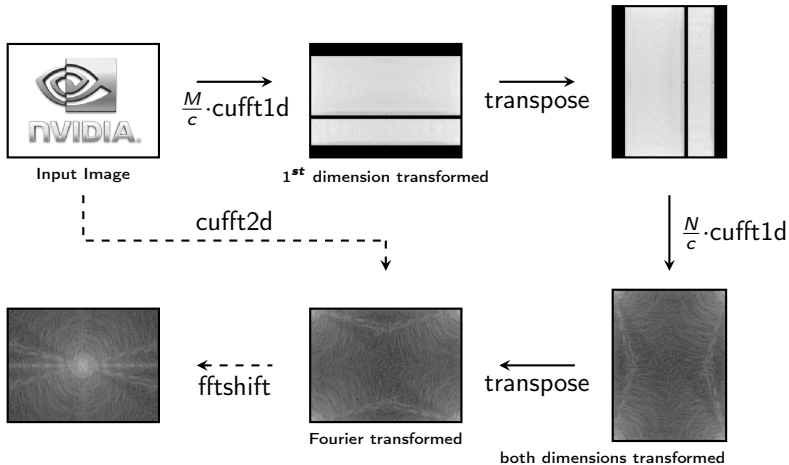


Flexible Fourier Transformation - Idea

- Separate Fourier transformation using one dimensional transforms
- This requires to transpose the intermediate results twice
- If the memory is not sufficient, transform the data in chunks



Flexible Fourier Transformation



c : chunk size, N : image width, M : image height




Flexible Fourier Transformation – Implementation

$$\begin{aligned} F &= FFT(f) \\ &= \text{cufft2d}(f, \text{width}, \text{height}) \\ &= \text{cufft1d}(\underbrace{\text{cufft1d}(f_x, \text{width}, \text{height times})^T}_{1^{\text{st}} \text{ dimension transformed}}, \text{height}, \text{width times})^T \\ &= \underbrace{\text{cufft1d}(F_x, \text{height}, \text{width times})^T}_{\text{fully transformed}} \end{aligned}$$



Memory Saving one dimensional Transformations



c : chunk size

Only one chunk is processed at a time. The chunk size should be as large as possible to exploit the memory on the card and avoid memory copy overhead.



Hardware

- Intel Core Quad Q9550 @ 2.83 GHz (March 2008)
- 8 GB RAM
- Graphic Cards:
 - NVidia GeForce 8600 GTS, 256MB (November 2006)
 - NVidia GeForce GTX 280, 1024MB (June 2008)
- Windows 7 RC1




Evaluated Image Sizes

| Width | Height | Megapixel | Comment |
|-------|--------|-----------|---|
| 6048 | 4032 | 24.38 | Large DSLR Camera Format |
| 4096 | 4096 | 16.78 | Large Power of 2 Example |
| 4752 | 3168 | 15.05 | Standard Digital Camera Format |
| 3456 | 2304 | 7.96 | Standard Digital Camera Format |
| 2560 | 1920 | 4.92 | Photo Mobile Phone Camera |
| 2048 | 2048 | 4.19 | Power of 2 Example |
| 2353 | 1568 | 3.69 | Standard Digital Camera Format |
| 1920 | 1080 | 2.07 | Full HD Video Format & Desktop Resolution (16 : 9) |
| 1280 | 1024 | 1.31 | Typical Desktop Resolution (4 : 3) |
| 720 | 480 | 0.35 | High Definition Video Format |
| 512 | 512 | 0.26 | Small Power of 2 Example |



Performance - GTS 8600 256MB



Transforming one channel

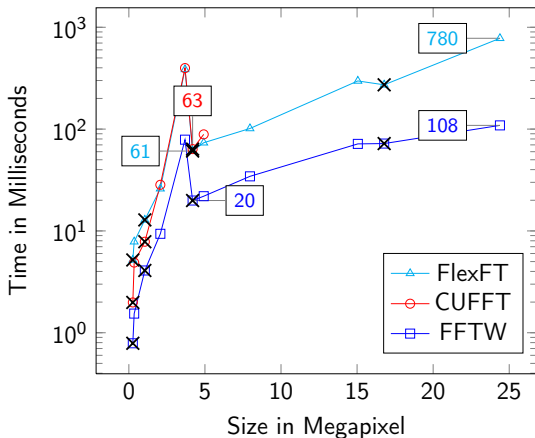
X marks power-of-2 sizes which perform best on FFT

Measuring

```
plan CUFFT
sync threads
start timer
for i = 1:k
    until all rows done
        copy to GPU
        1d transform
        copy to Host
transpose
until all rows done
    copy to GPU
    1d transform
    copy to Host
transpose
sync threads
t = stop timer / k
```



Performance - GTS 8600 256MB



Transforming one channel

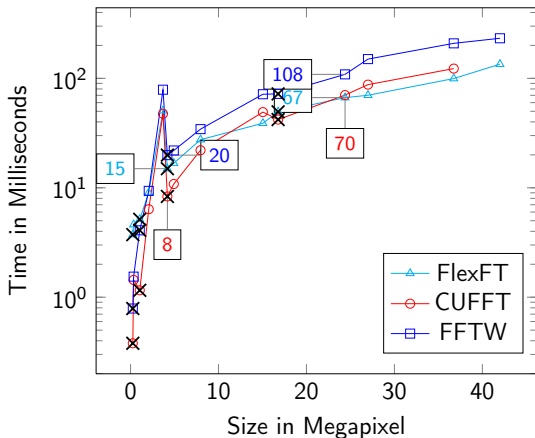
X marks power-of-2 sizes which perform best on FFT

Interpretation

- CUFFT and FlexFT achieve similar performance
- FFTW always faster, especially ≥ 5 MP due to memory copies in FlexFT and better plans (FFTW_PATIENT)

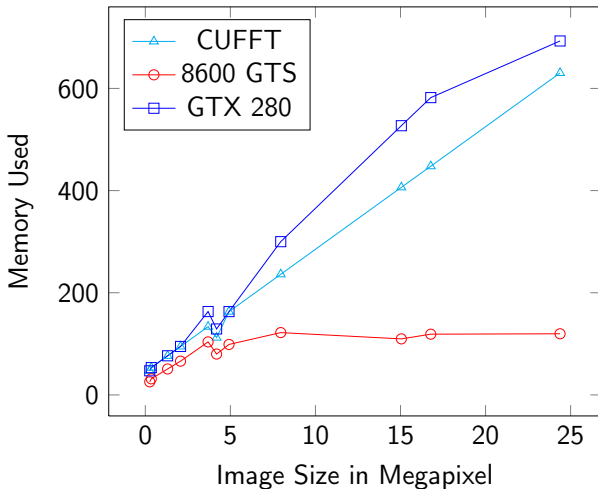


Performance - GTX 280 1024MB





Memory Usage CUFFT vs. FlexDFT





Memory Usage CUFFT vs. FlexDFT

GeForce 8600 GTS

Uses as much memory as possible, maximum of 90% of free memory.

Chunkifies the input to according to fitting partition sizes.

GeForce GTX 280

Allocates buffer for input and output on the device.

Transforms the data in chunks then, as there is not enough memory left for a single transform.

This approach has no internal CPU ↔ GPU memory copies and thus is much faster.



Conclusions

- Use separability saves memory on device without timing loss.
- Performance equal to 2D CUFFT but much larger transformations with little memory possible.
- Allows trading memory consumption for speed.
- Extensible for higher dimensions (e.g. for geological or medical 3d scans).



Thank you for your attention!

Questions?