

# Reduced Representations for In-Situ Visualization

Steffen Frey and Thomas Ertl - [Visualization Research Center, University of Stuttgart](#)

*Workshop on Software Frameworks for Scalable Scientific Simulations - ISC 2015*  
July 16th, 2015

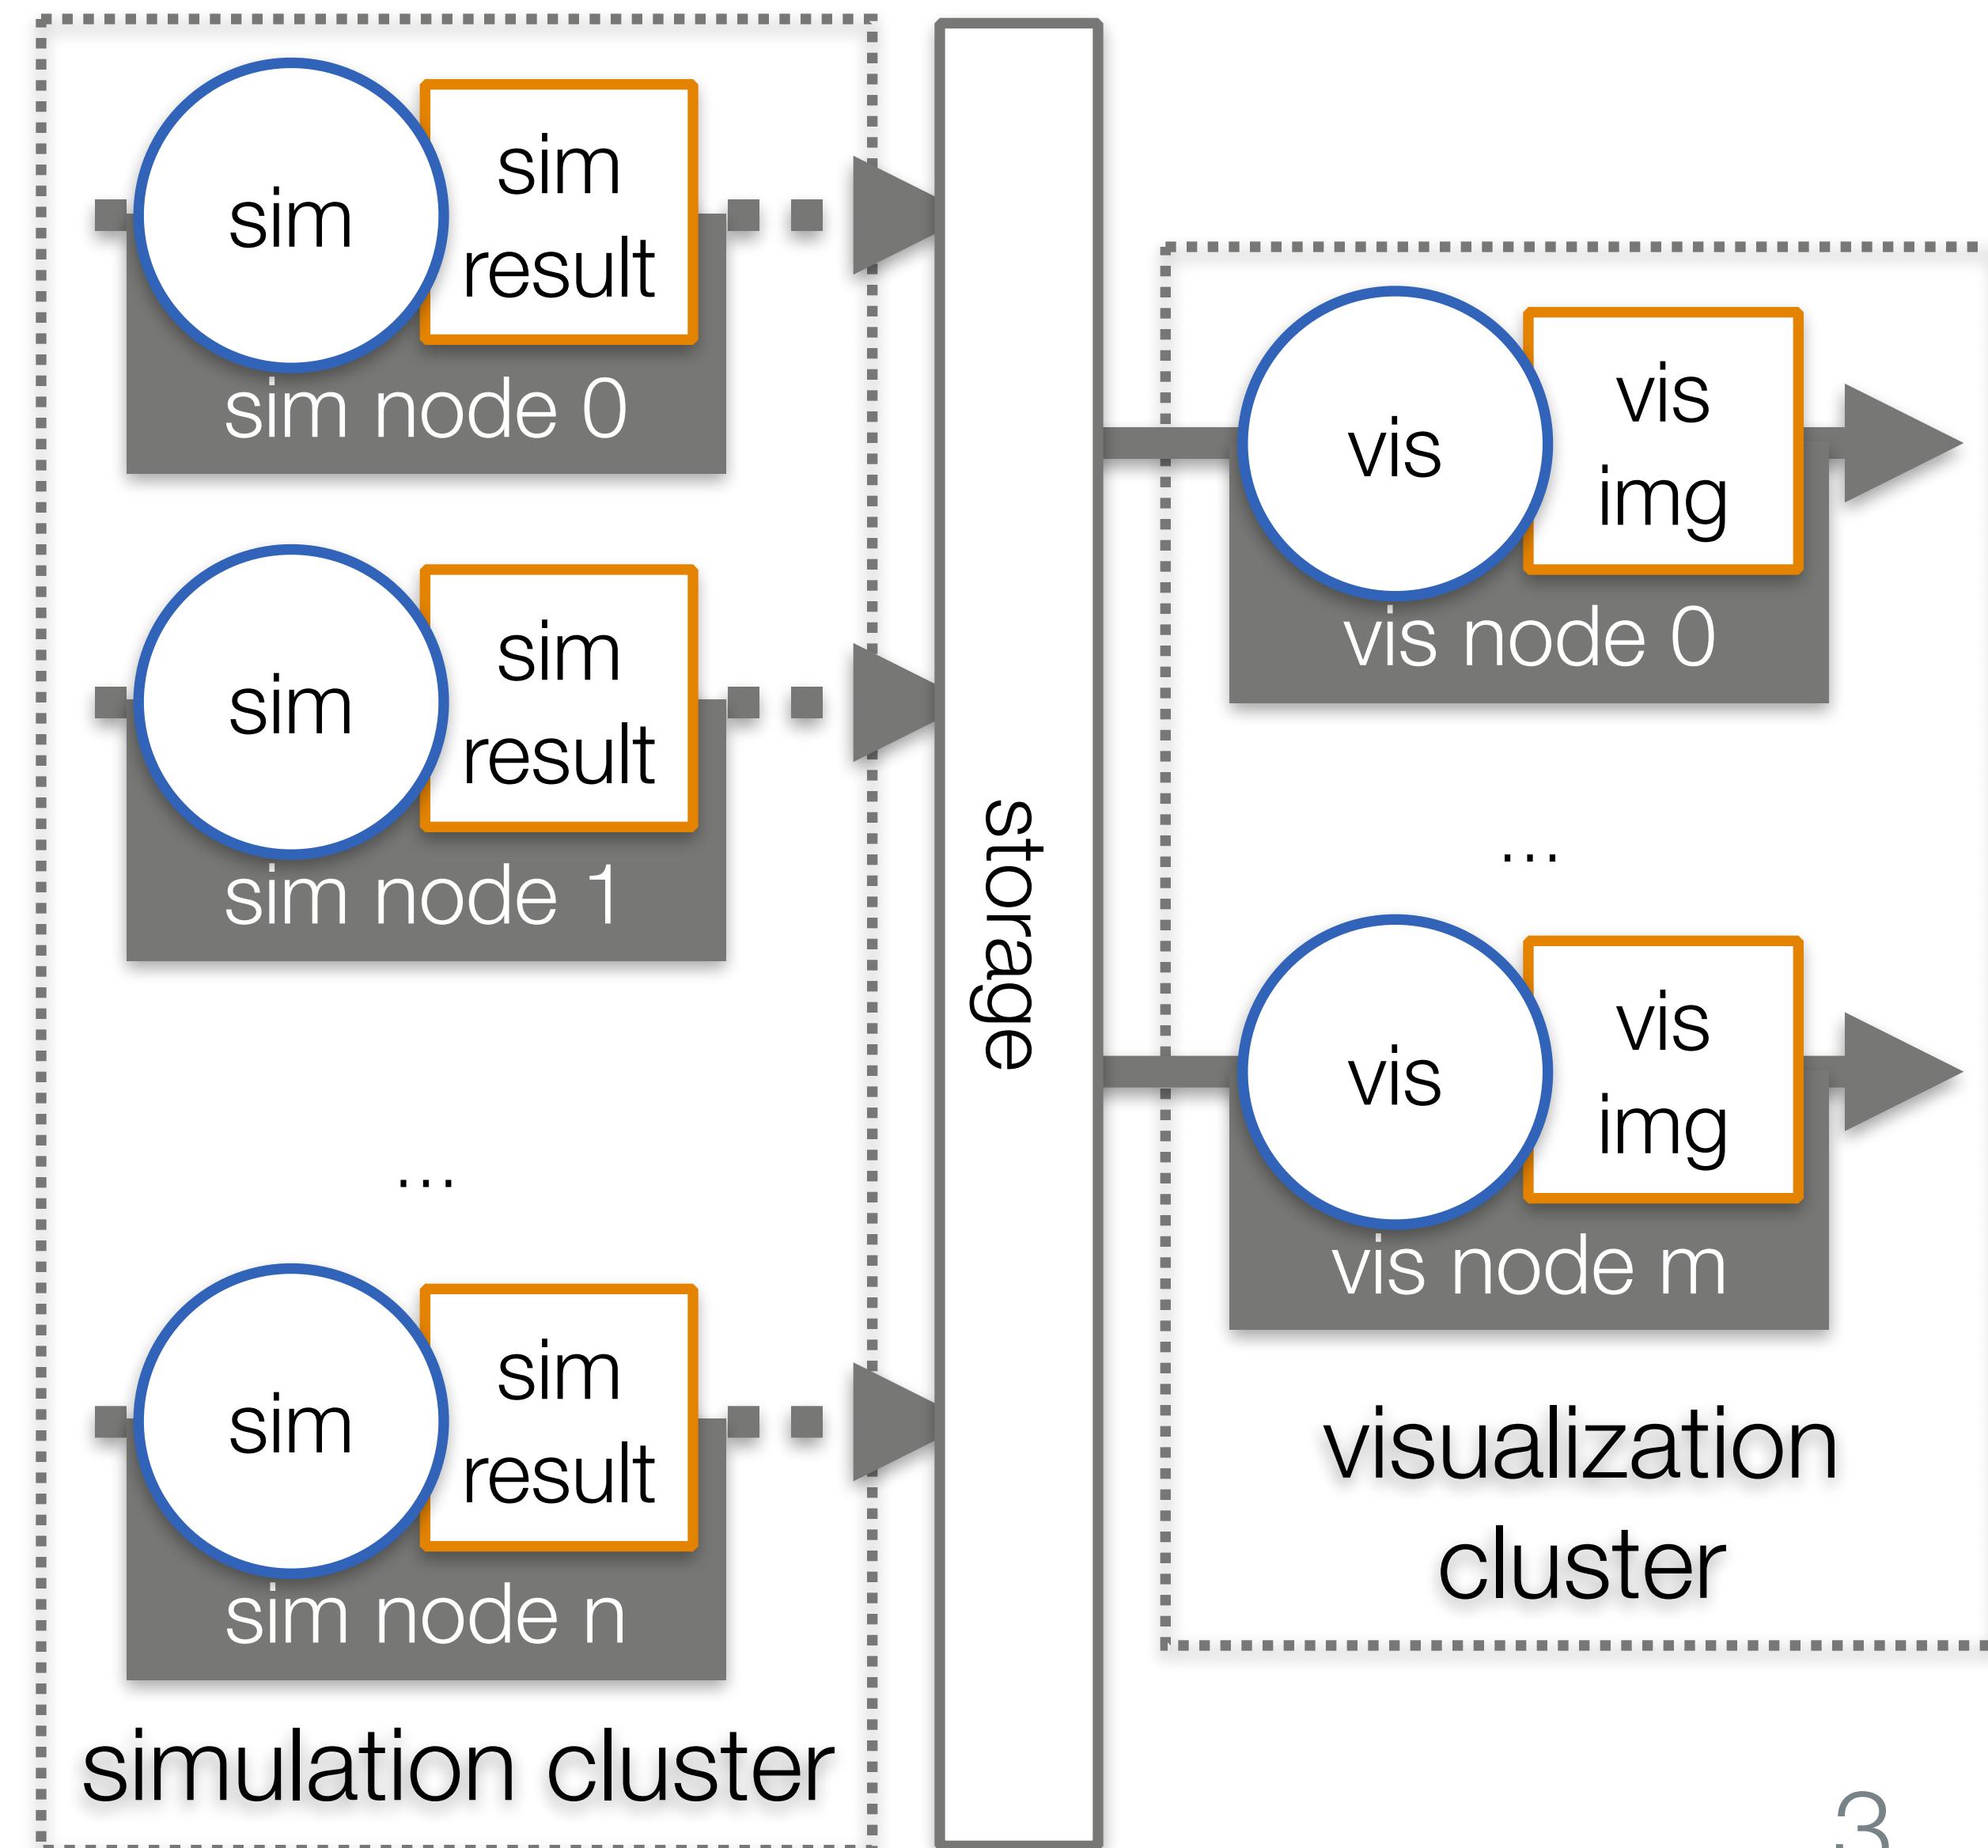
# Motivation

- increasing dataset sizes
  - large-scale supercomputer simulations
- storing (or moving) all generated raw data impossible/  
impractical
  - consumes a lot of space, time and power
  - just discarding data not desirable
- approach: in-situ (on site) reduction of data
- by transformation to visualization representation that
  - ... achieves a significant reduction in size
  - ... can be generated quickly/efficiently
  - ... enables the analysis the selected data of interest
  - ... allows for interactive exploration
  - ... can be used for monitoring and steering



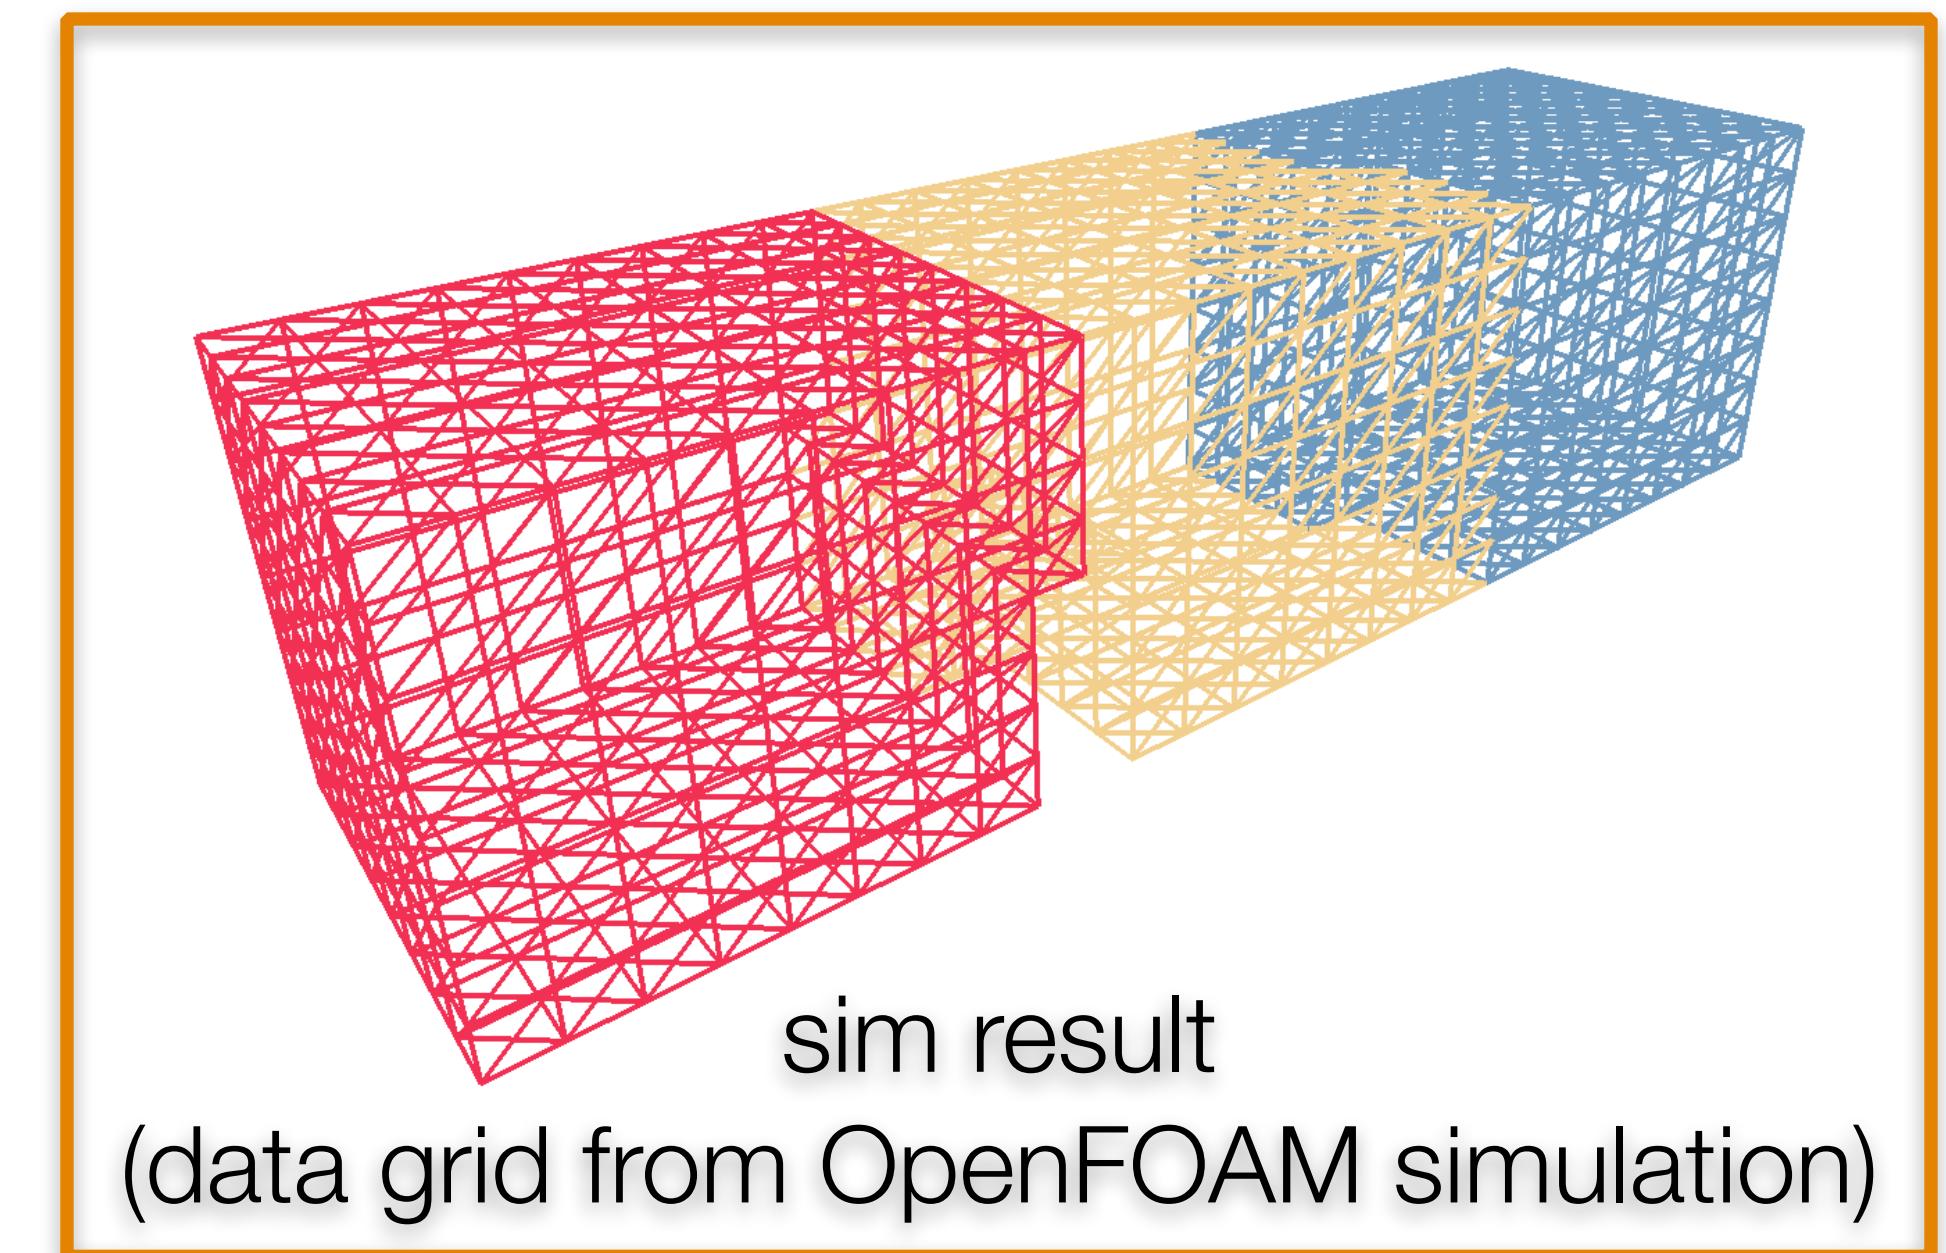
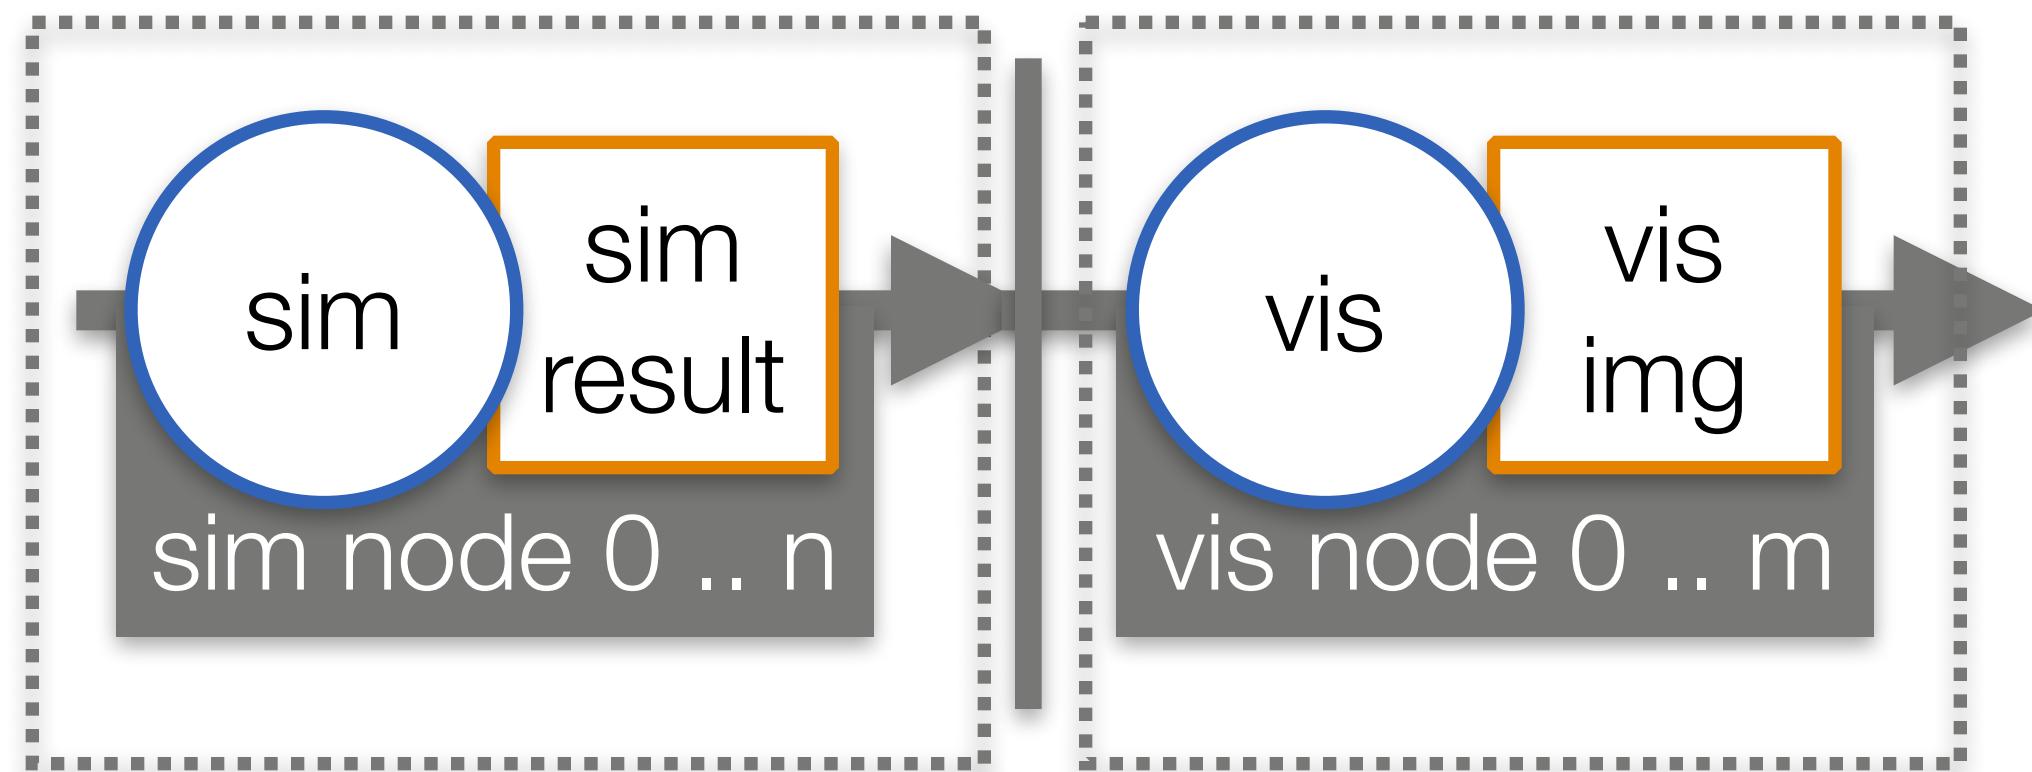
# Simulation-Visualization Coupling (1)

- *traditional simulation-visualization procedure*
  - decoupled **simulation** and **visualization**
    - vis runs independently after sim is completed
    - vis only needs to be able to read output data structure
  - sim result data stored directly → huge amounts of data need to be transferred
    - storage/time/energy constraints limit resolution of stored result data
      - typically w.r.t. time
  - ... but full flexibility preserved for exploring/analyzing stored data



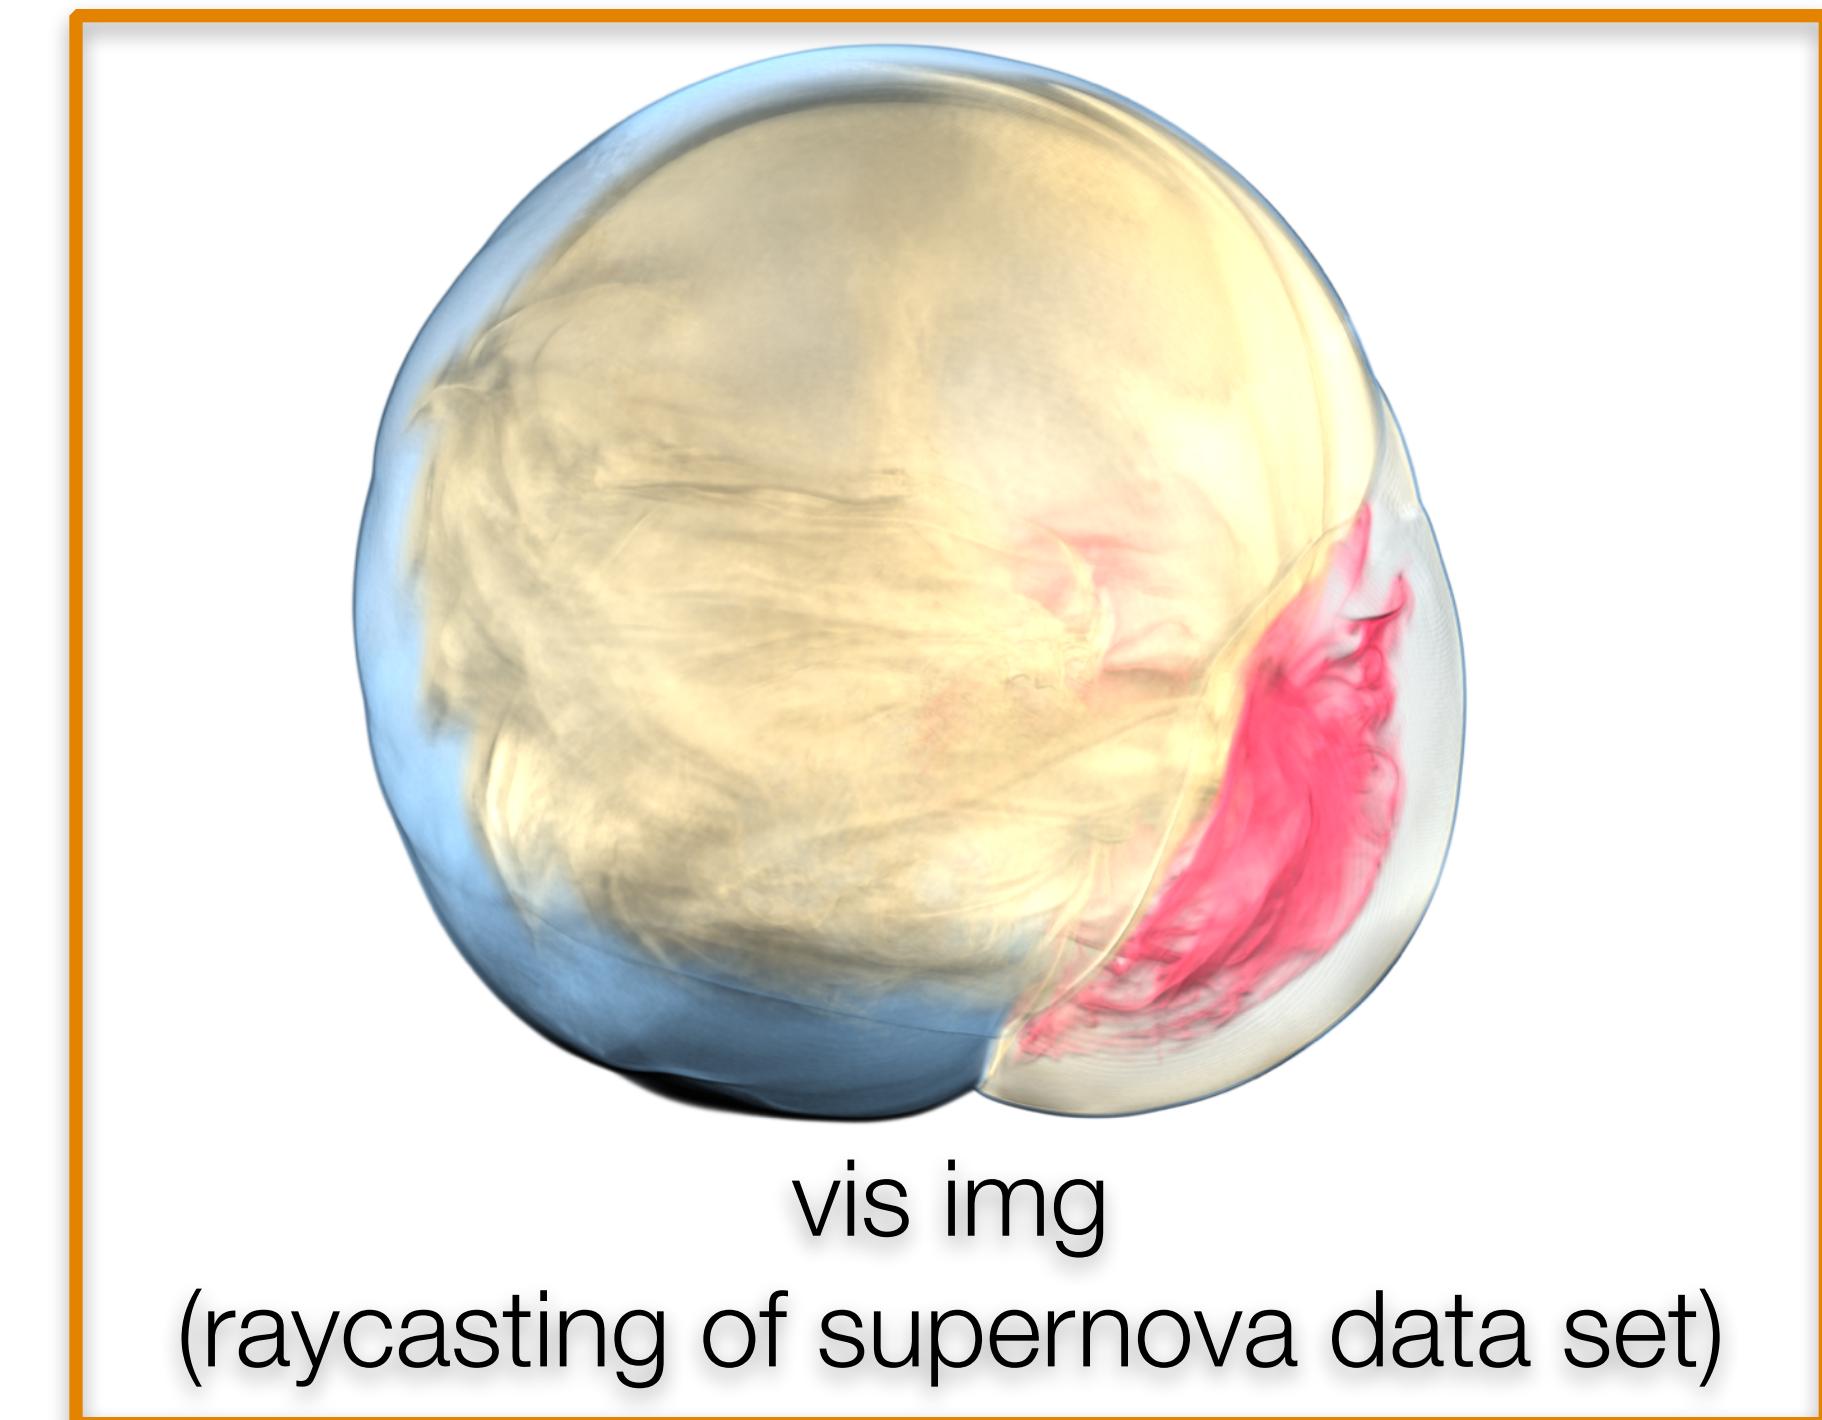
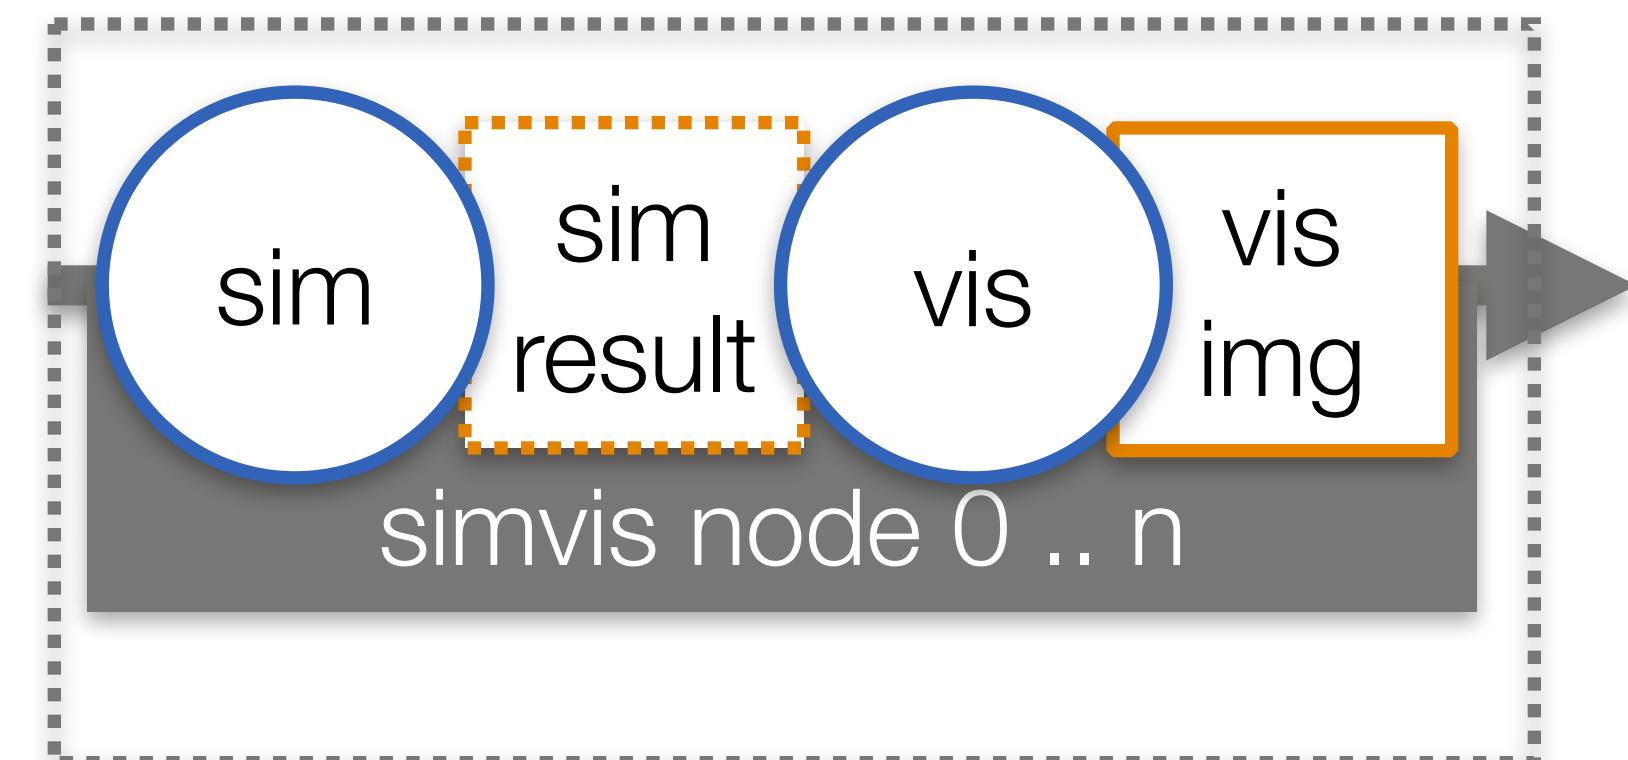
# Simulation-Visualization Coupling (2)

- *loose coupling*
  - simulation and visualization run concurrently
    - sim result not stored to disk anymore
    - amount of data that is stored permanently is reduced
      - vis result should be much smaller than sim result
    - results become available during simulation run
  - ... but on separate resources
    - high network load
    - forced data conversion
    - potentially redundant generation of data structures



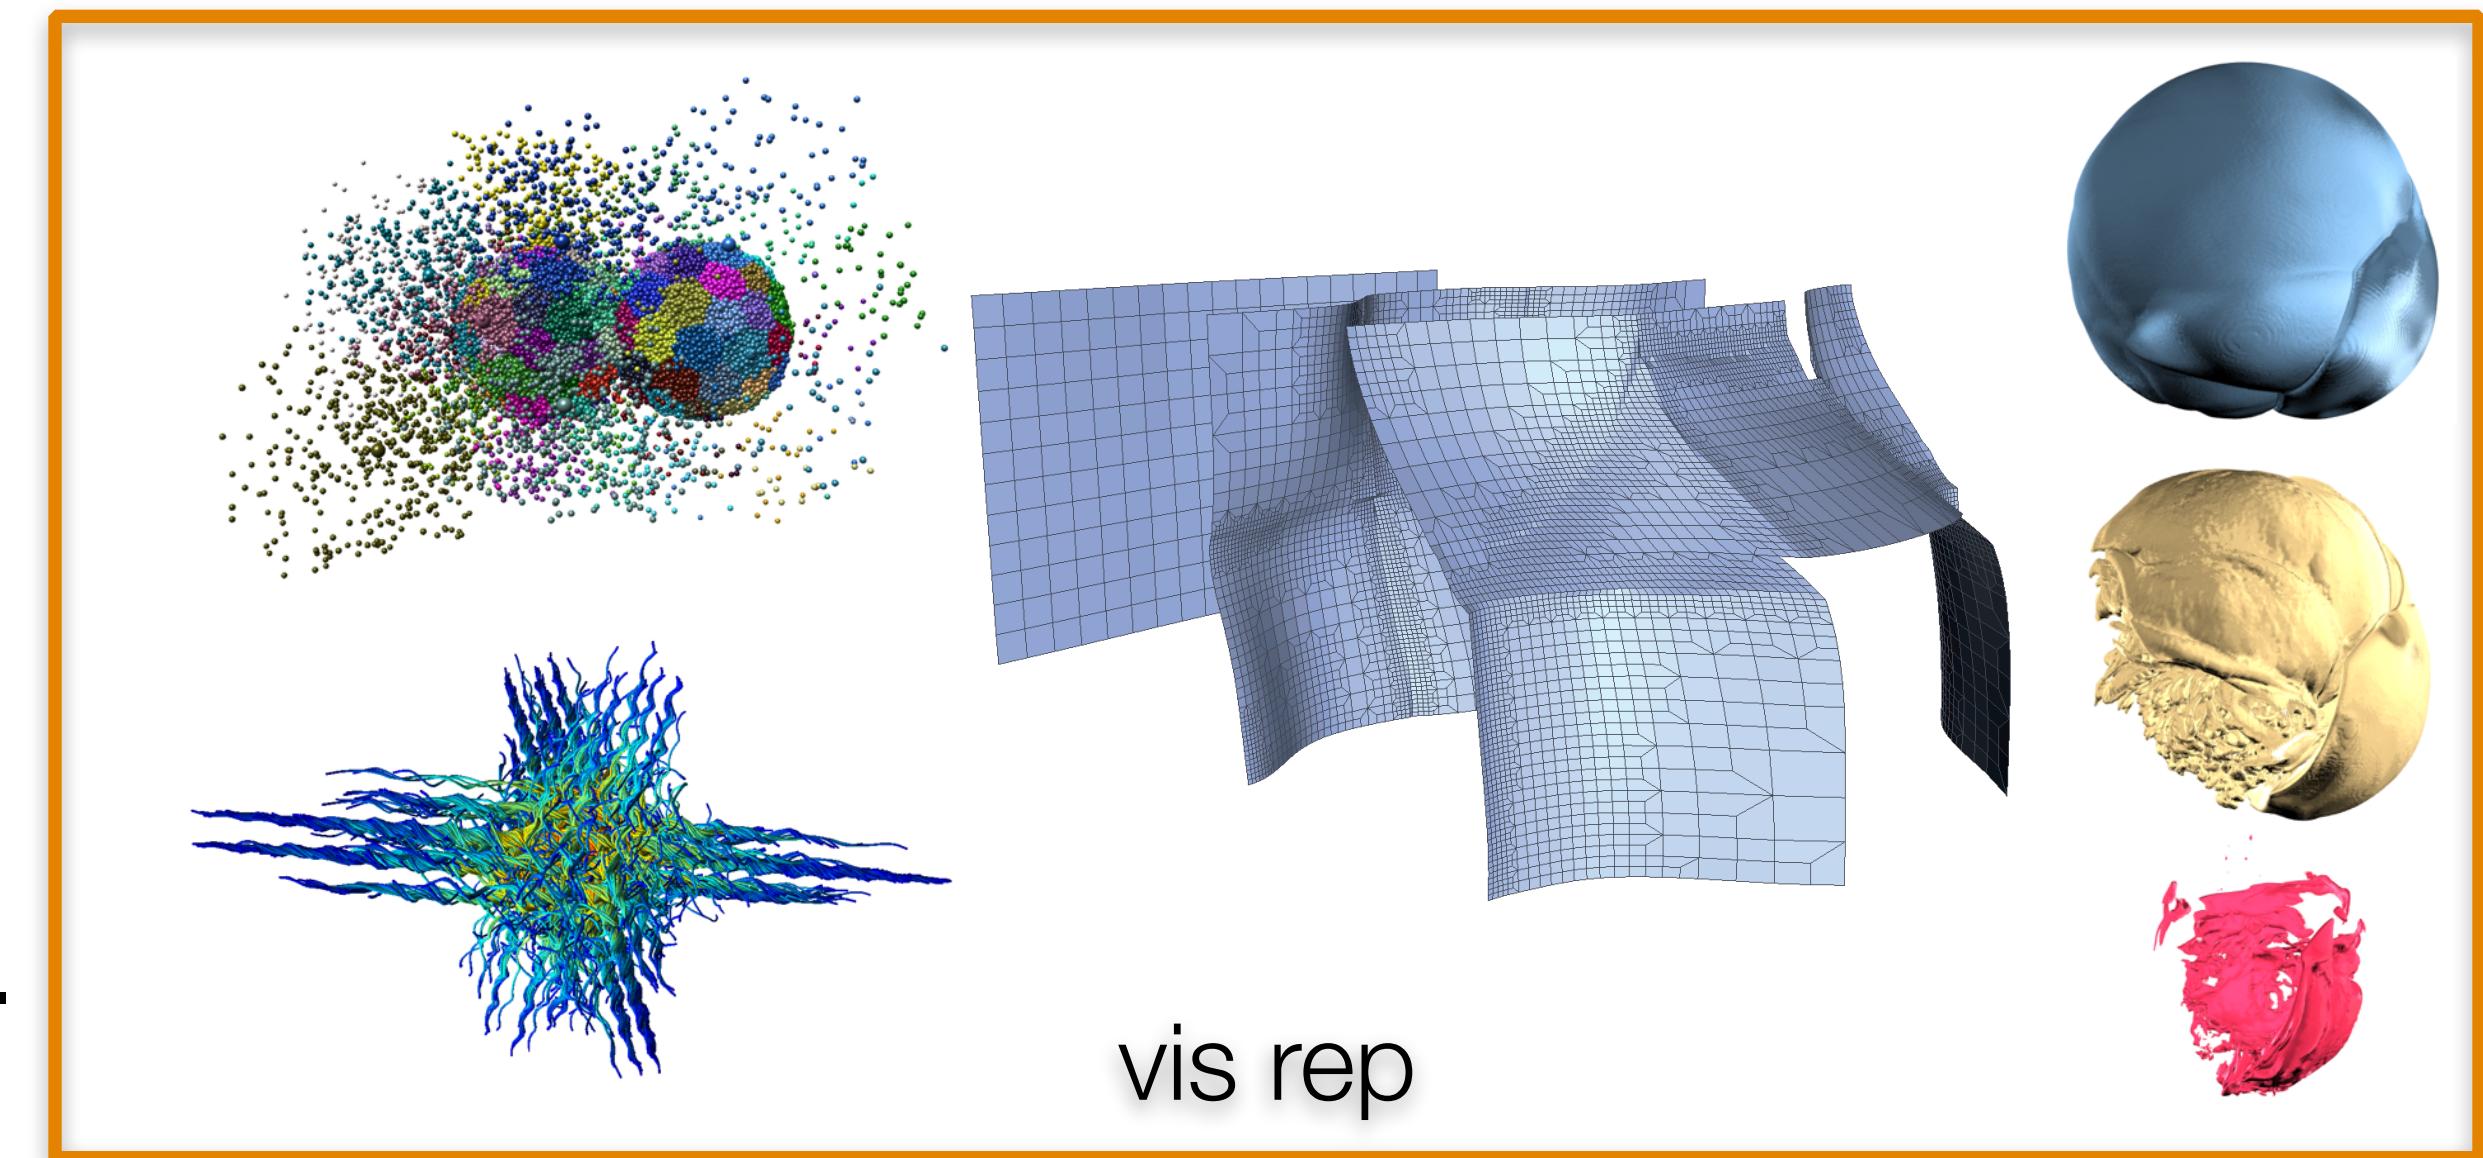
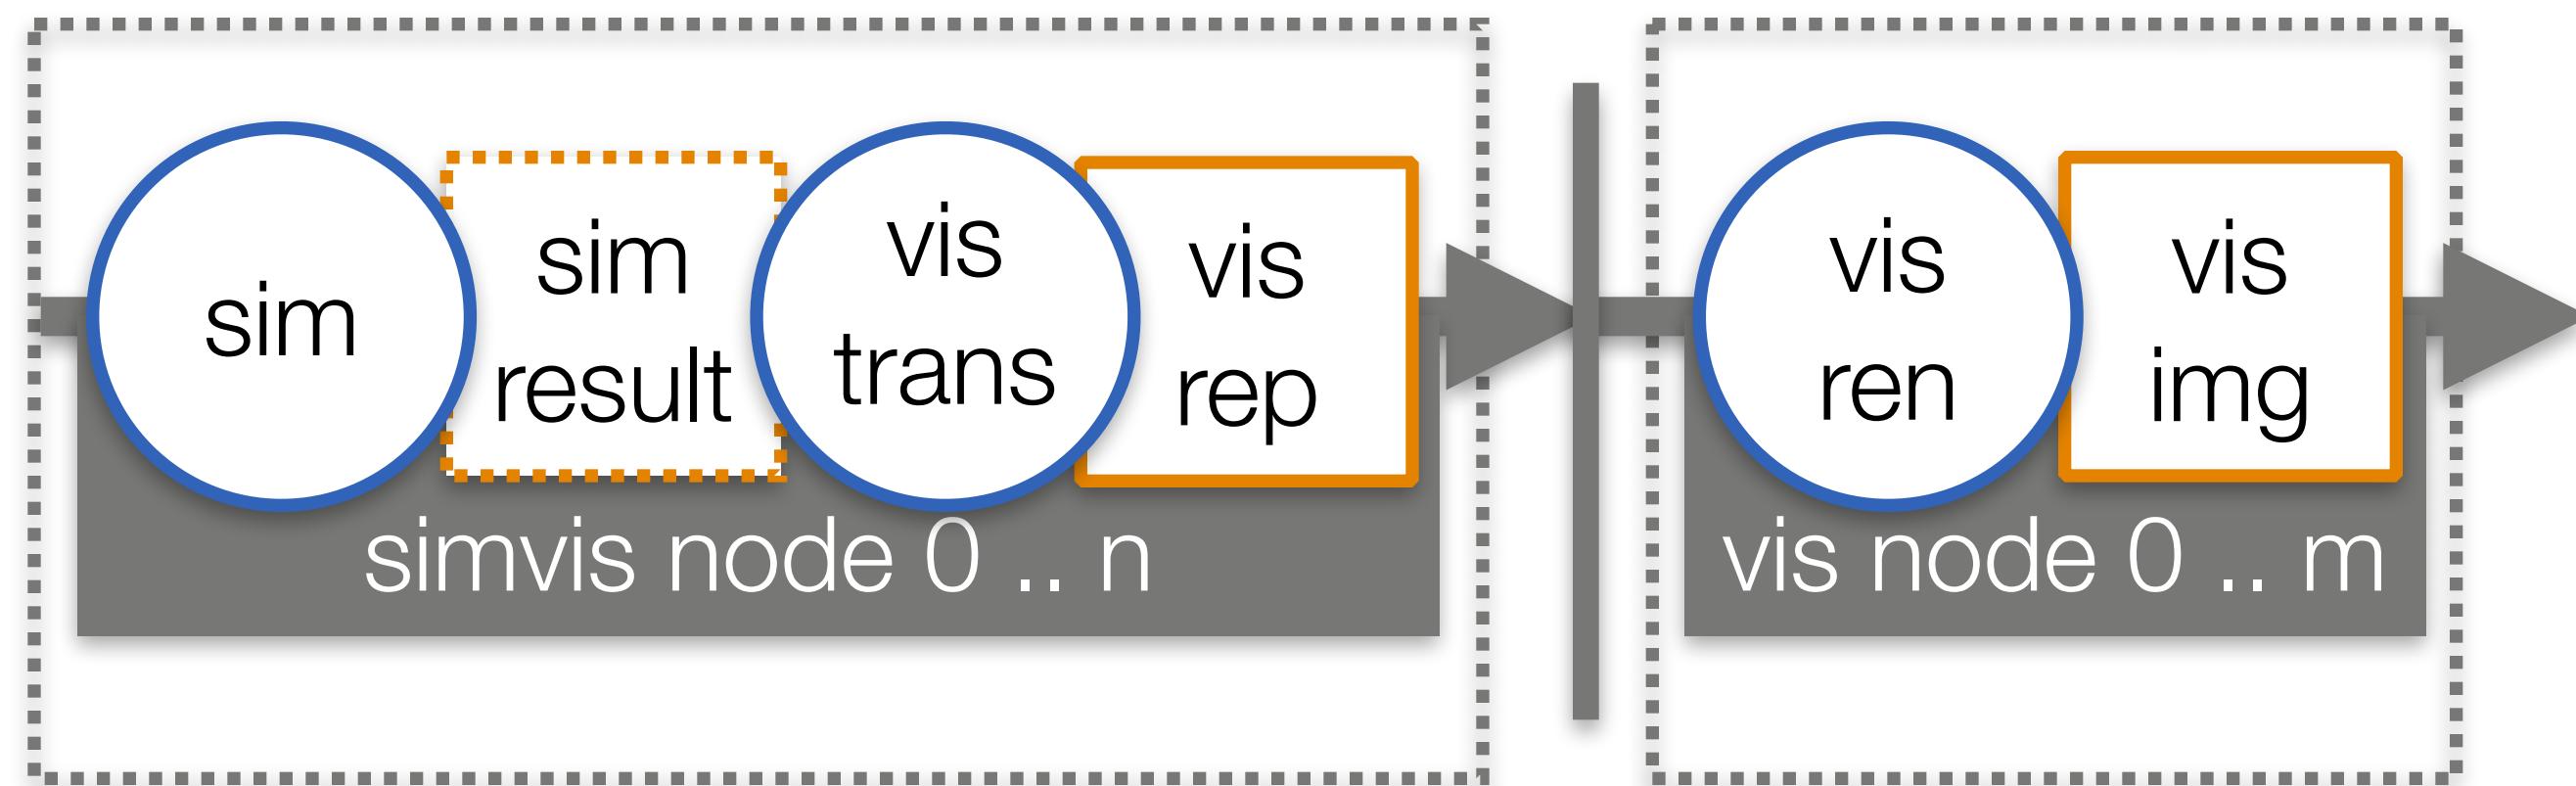
# Simulation-Visualization Coupling (3)

- *tight coupling*
  - visualization runs on the same cluster node as simulation (in-situ)
  - they share ...
    - data structures
      - optimally no transfer or data conversion required
    - space decomposition
    - compute resources
    - ...
  - the result are images generated by the visualization
    - no flexibility a posteriori for (interactive) exploration



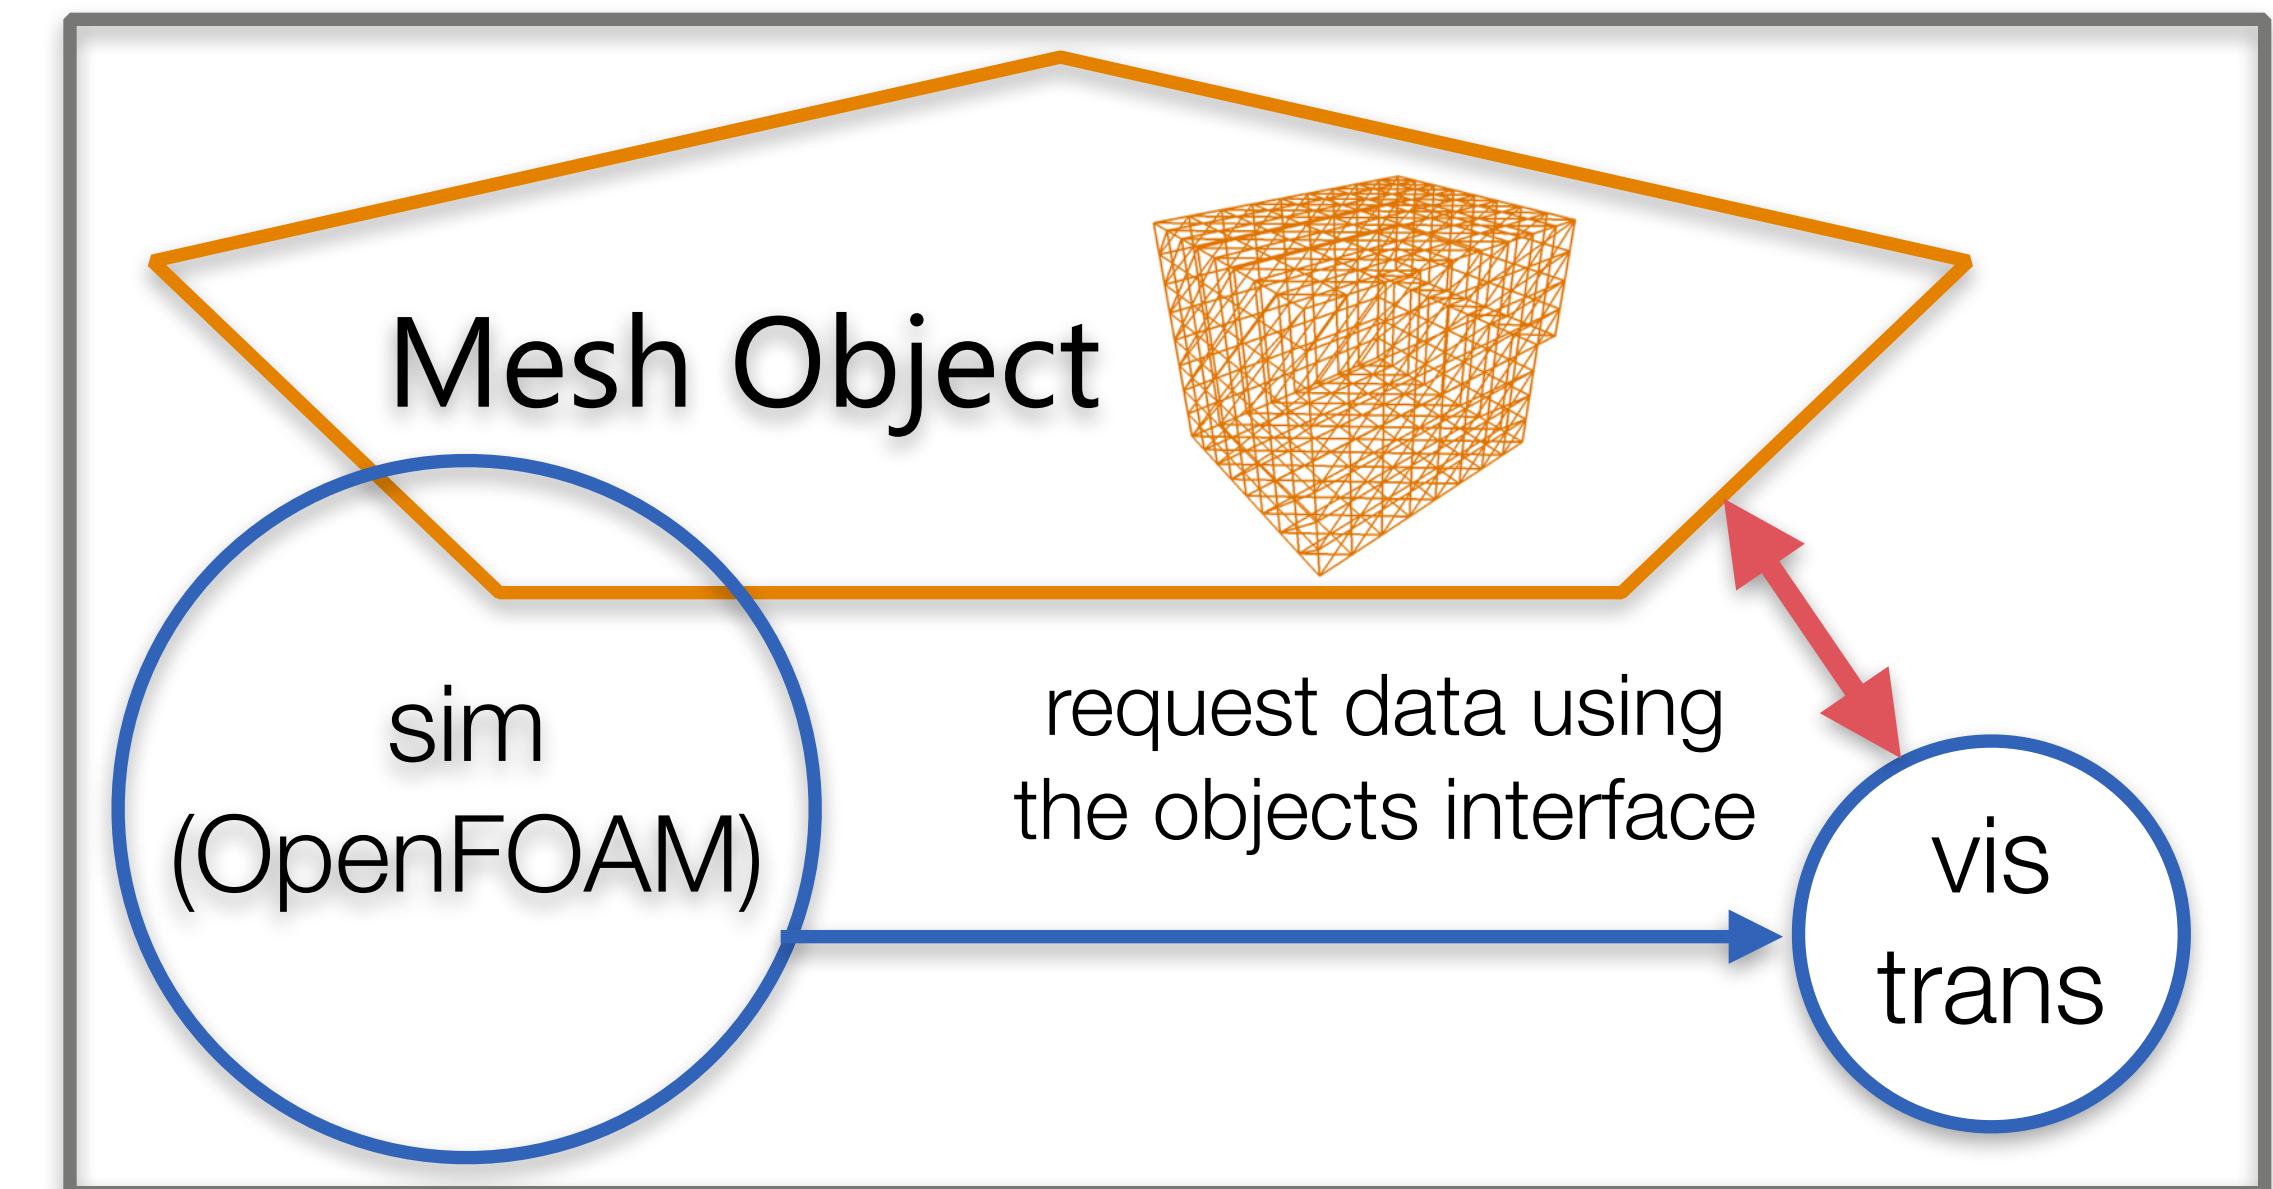
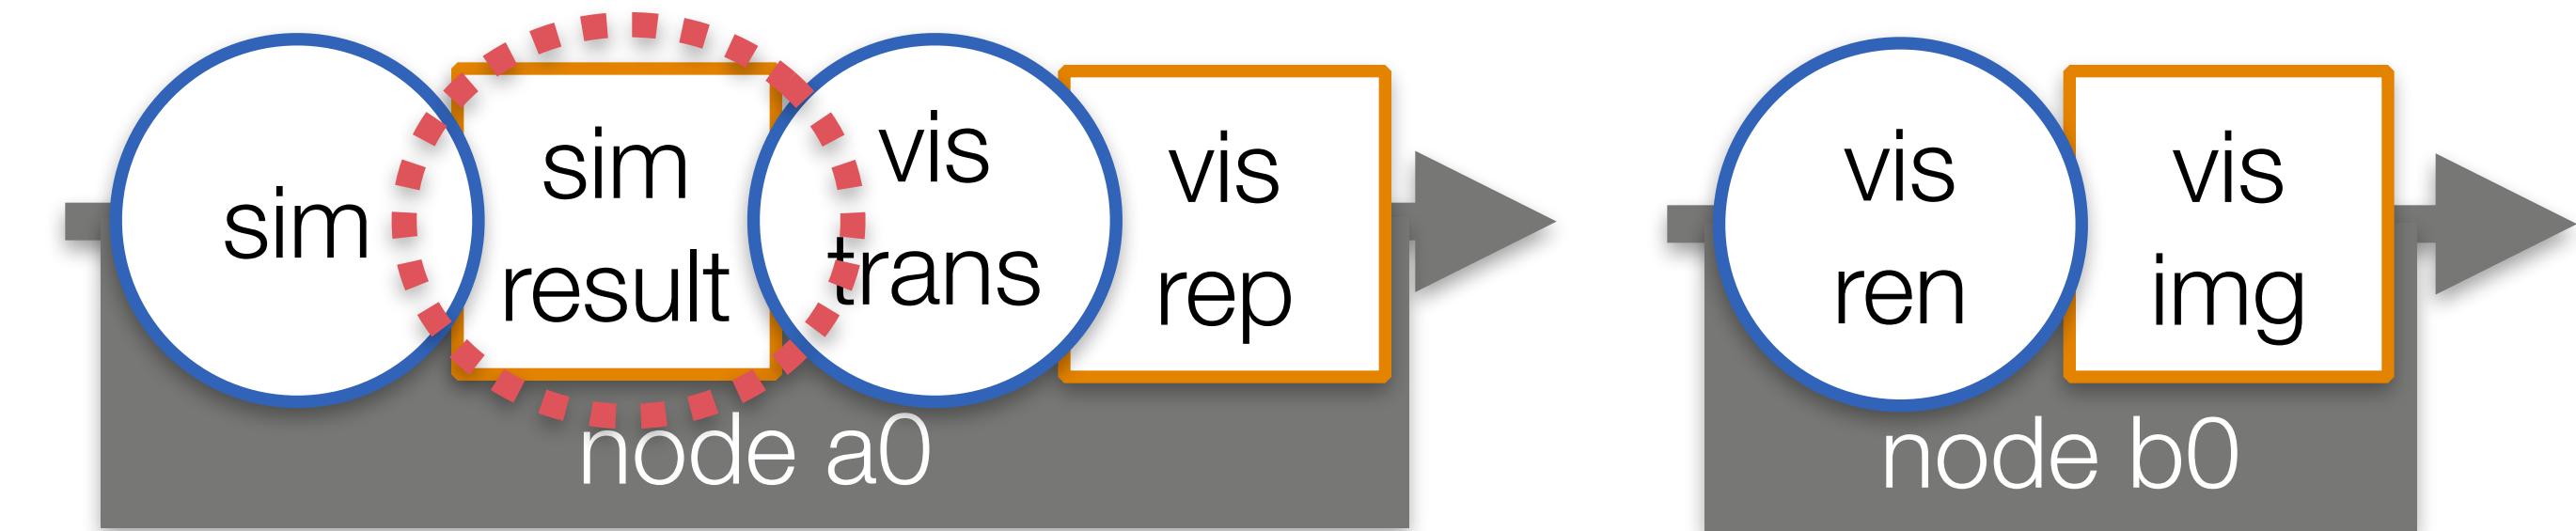
# Simulation-Visualization Coupling (4)

- *hybrid coupling*
  - visualization procedure split into two parts
    1. vis trans: generate intermediate visualization representation (vis rep)
      - tight in-situ coupling
      - reduces data & maintains flexibility
    2. vis ren: render image
      - visualization representation is transferred to vis node
      - generate image from intermediate representation
        - based on parameters that can be adjusted (interactively) by a user
          - e.g., camera position, transfer functions, etc.
  - focus on hybrid coupling in the following



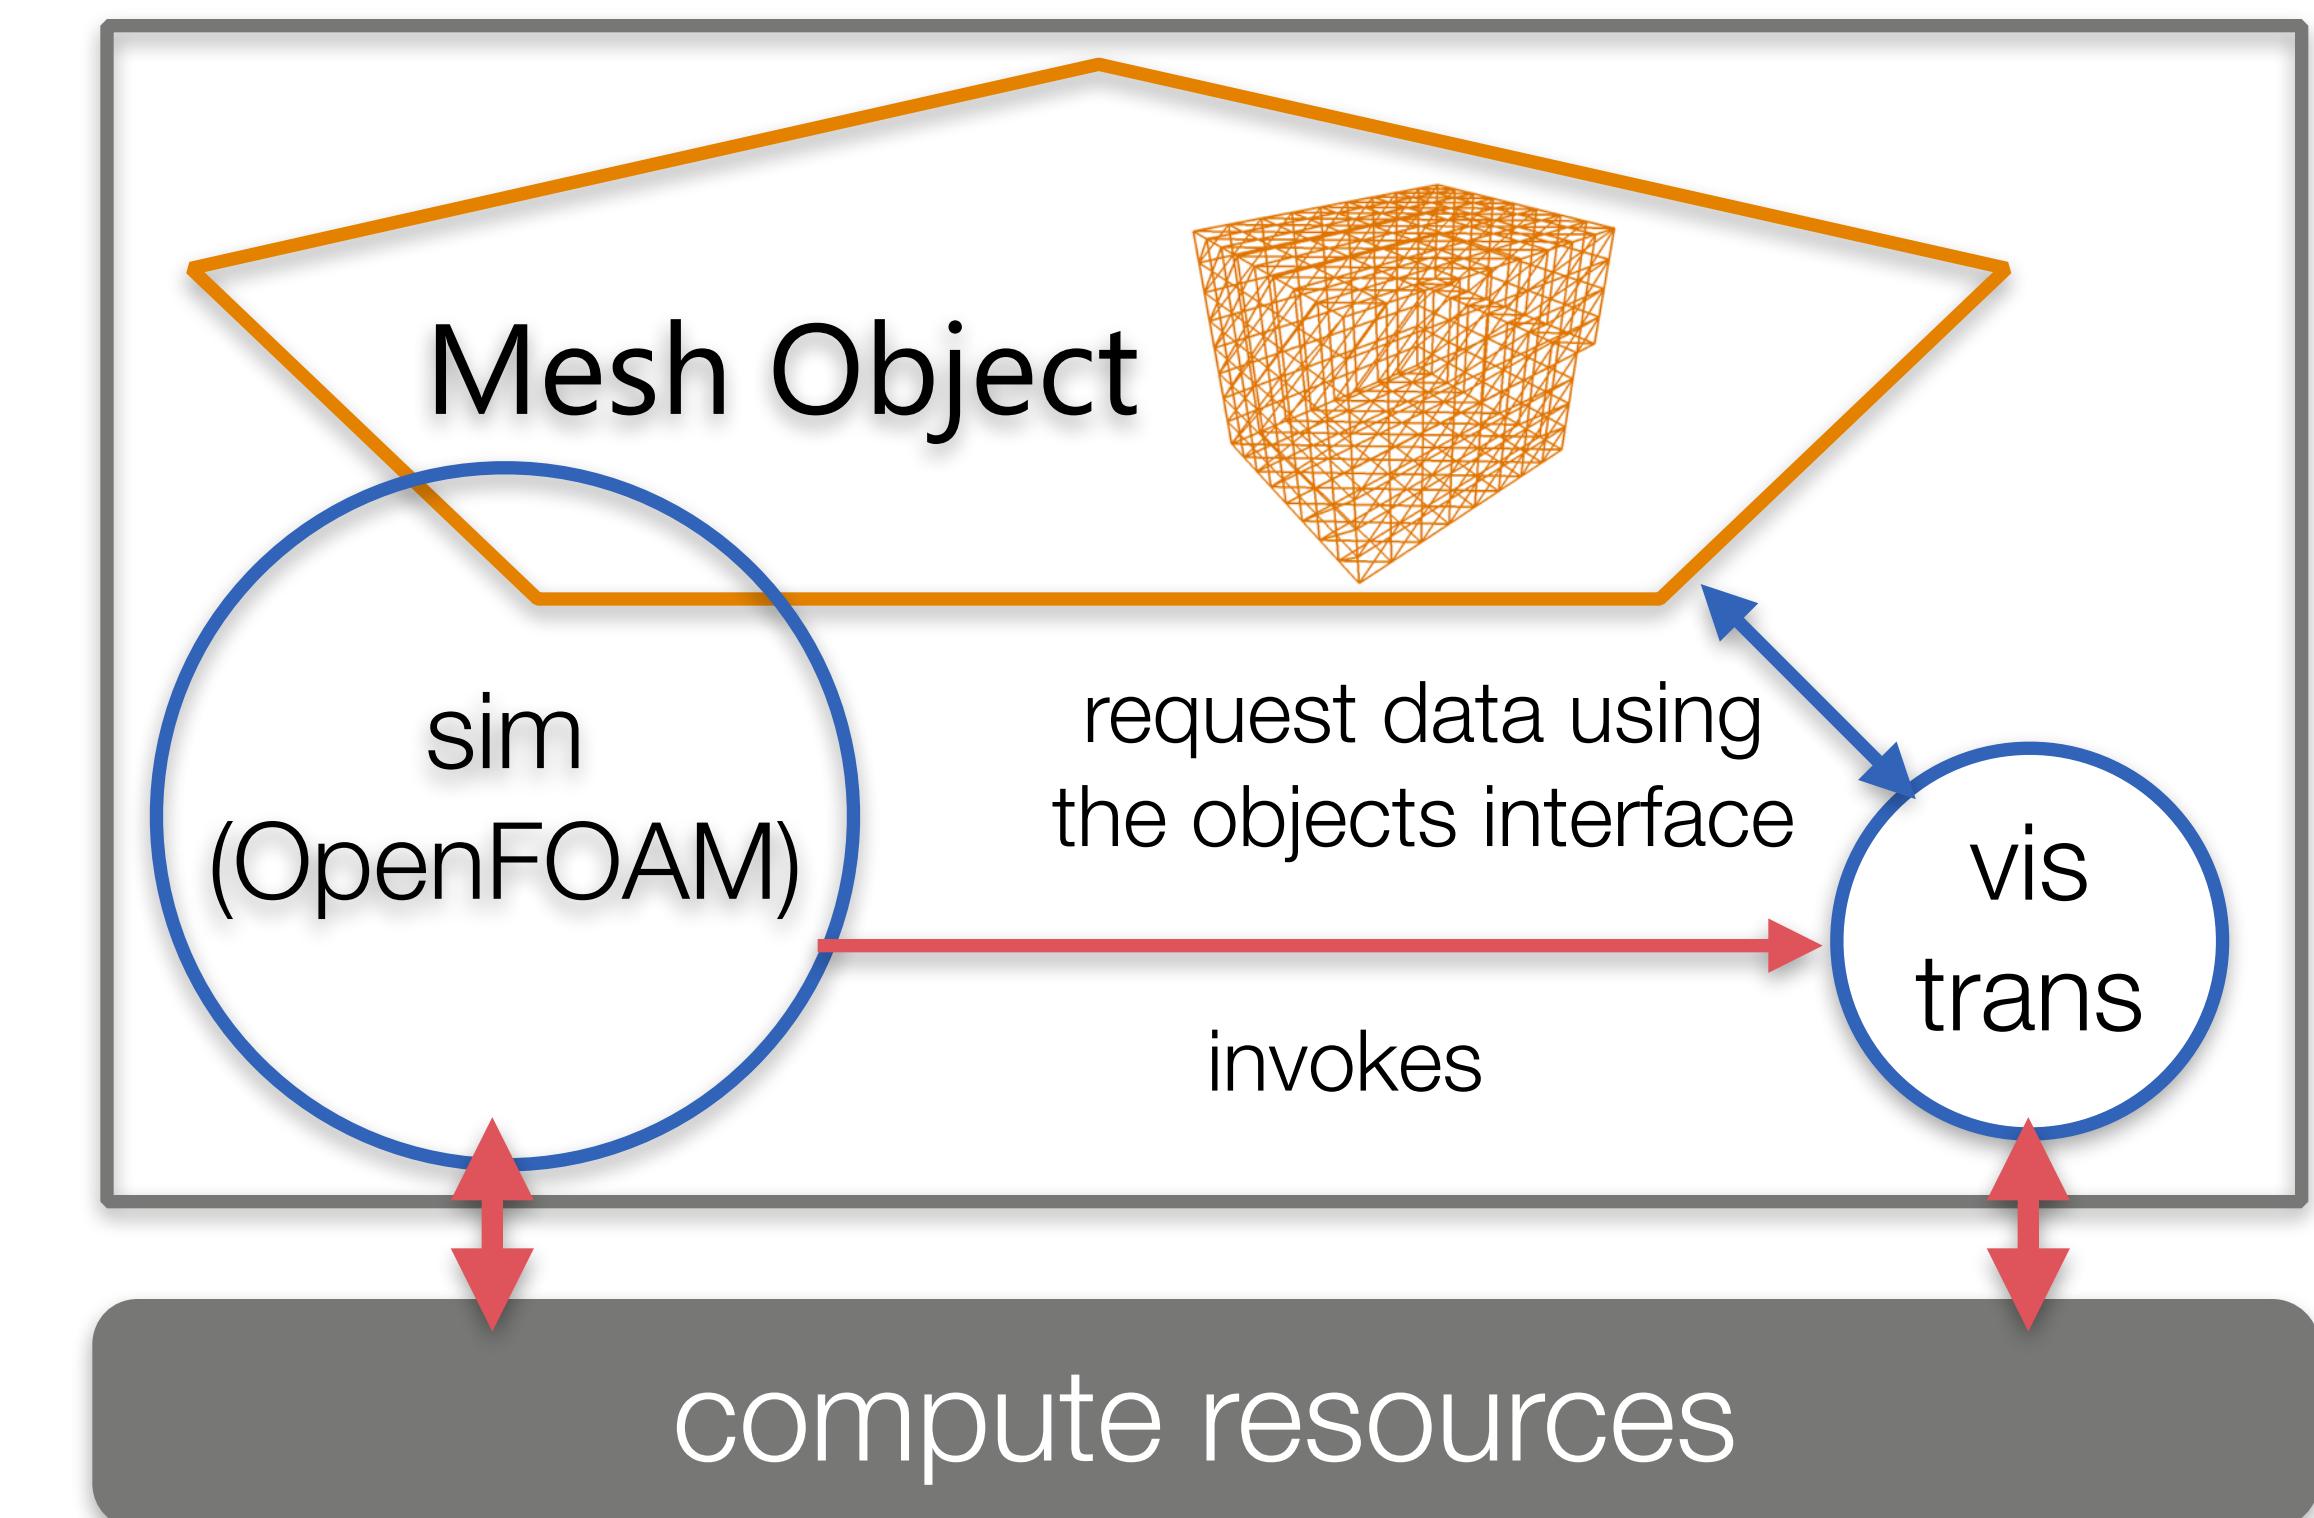
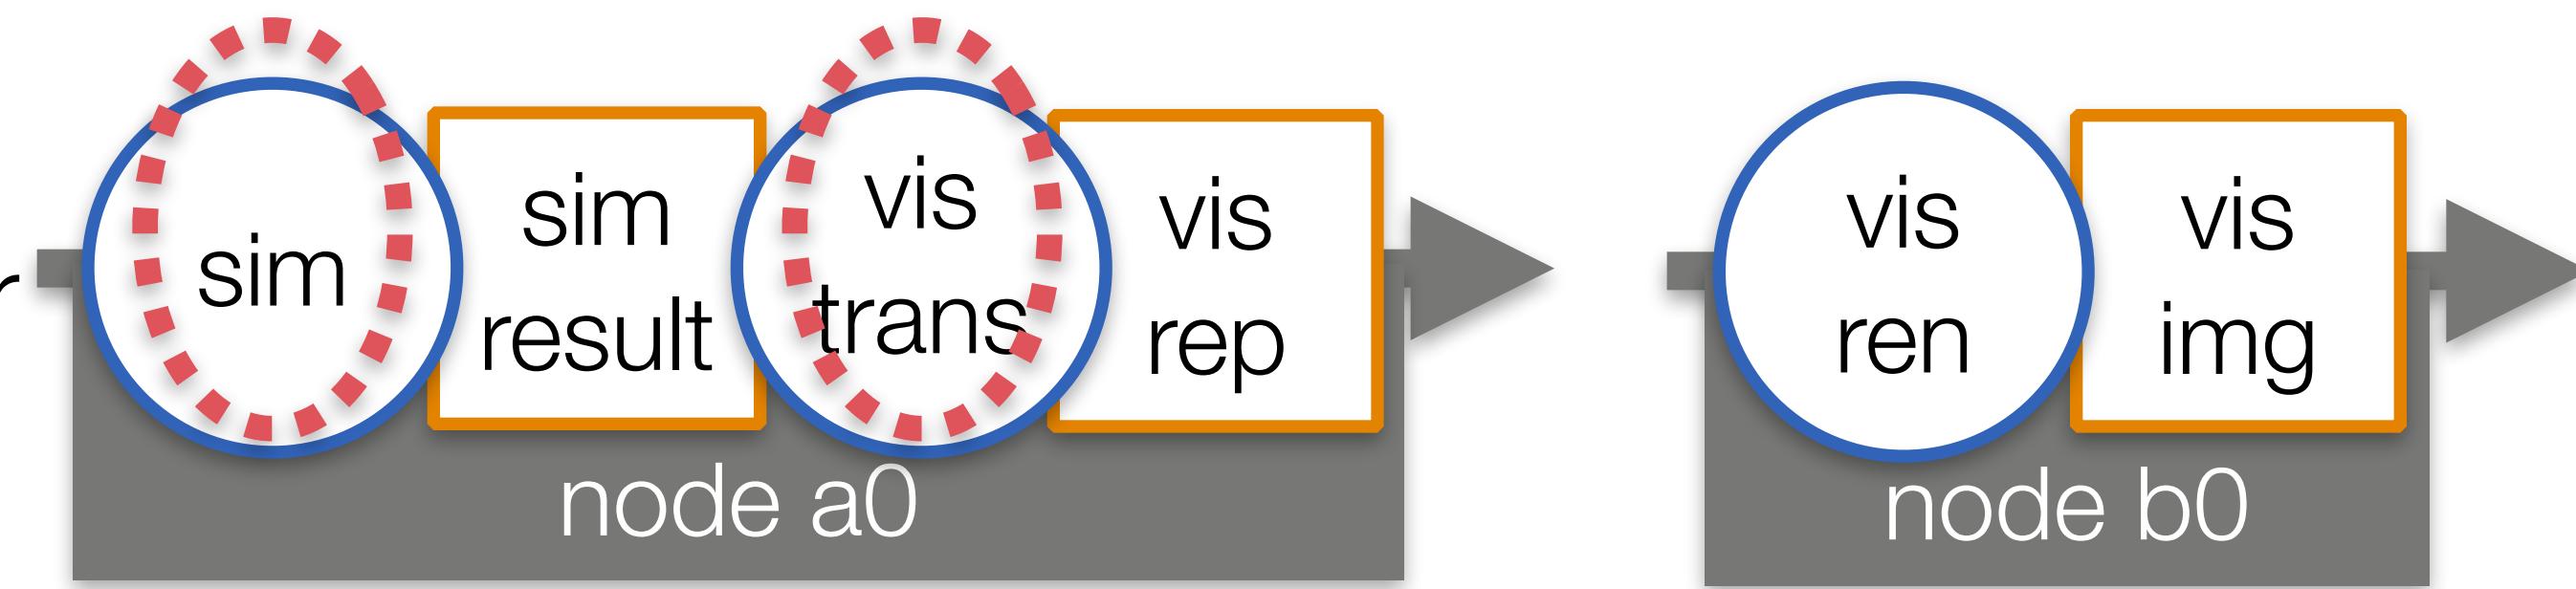
# In-Situ Visualization Challenges (1)

- *simulation-visualization interface*
  - minimize data transfers
  - avoid data conversion wherever possible
  - use optimized data structures from simulation directly for visualization
    - not necessarily the most efficient solution
      - visualization can exhibit very different access characteristics
  - example: raycaster using OpenFOAM
    - simulation and visualization run in the same process
      - visualization embedded as a library
      - data structures and interface access functions of OpenFOAM are directly used
    - i.e., visualization as a library



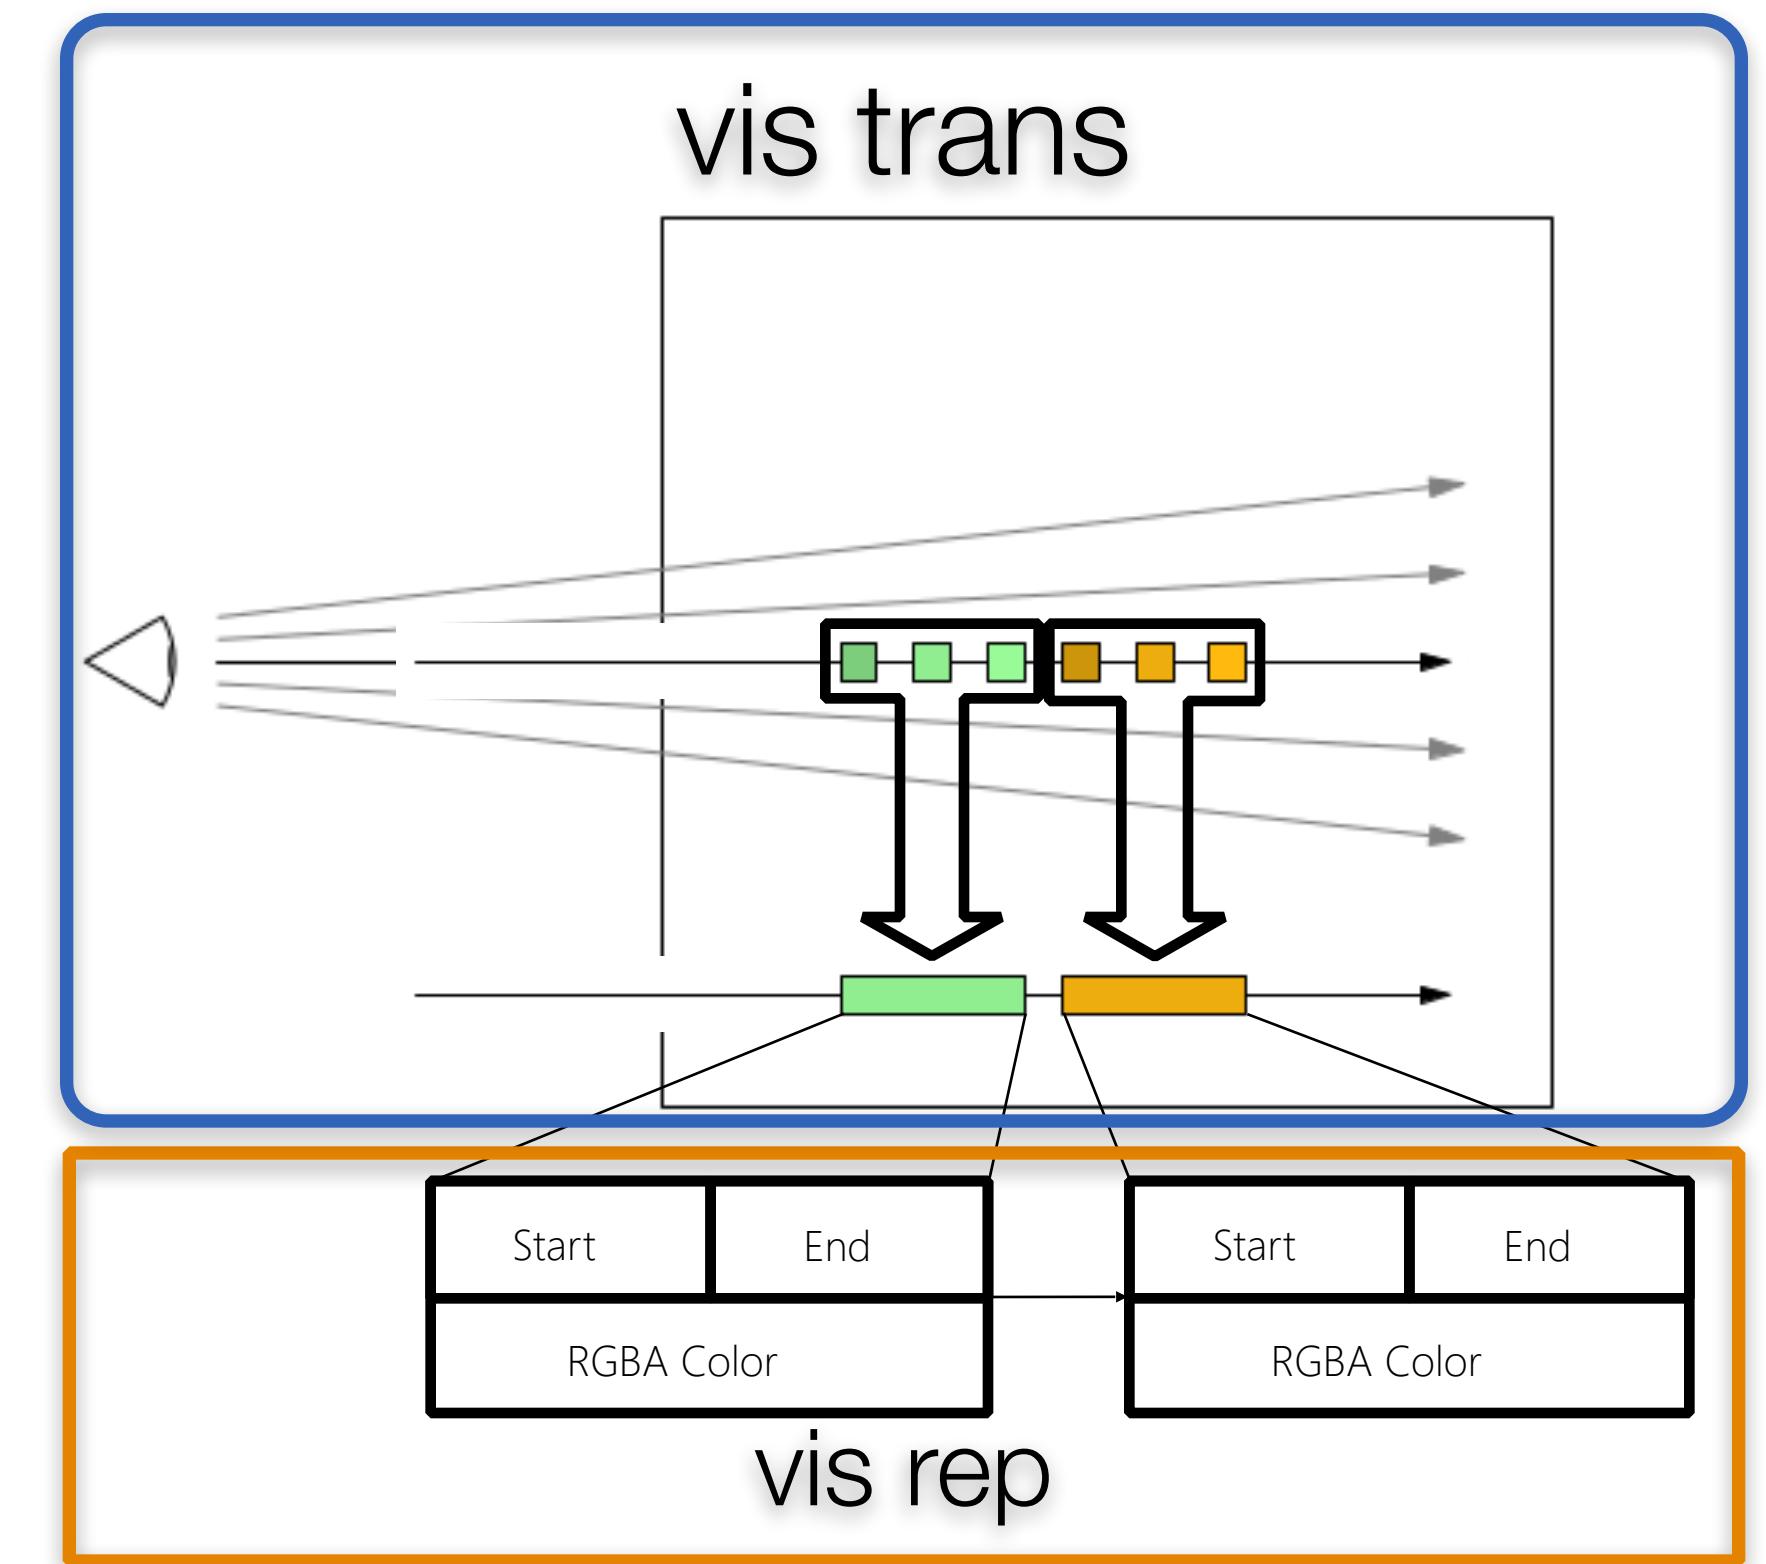
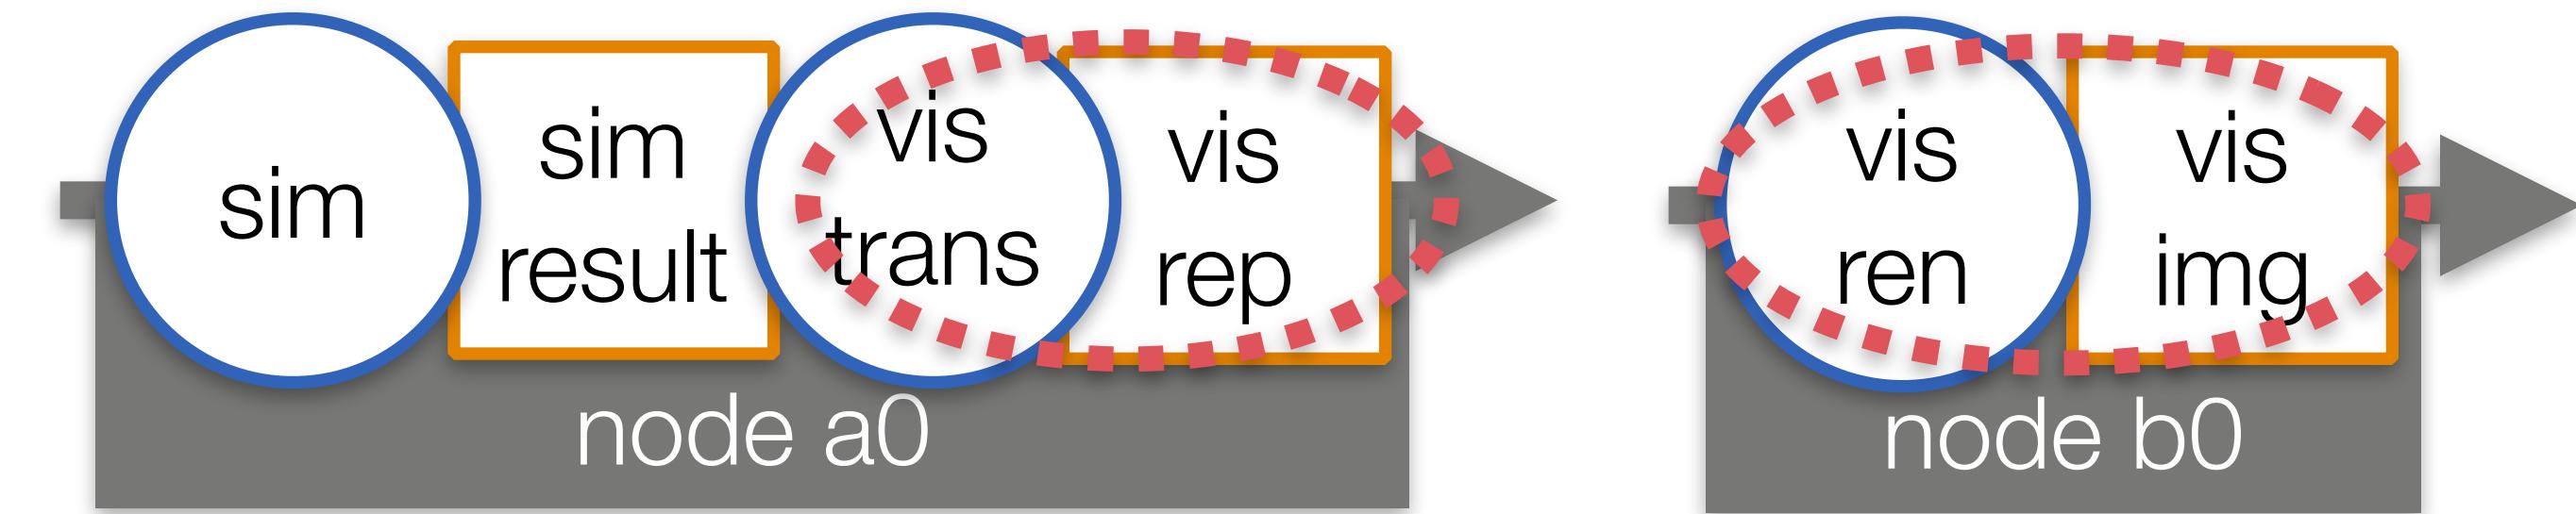
# In-Situ Visualization Challenges (2)

- *resource sharing*
  - visualization competes with simulation for compute performance
  - scheduling
    - vis trans can be deferred
    - utilize open slots due to load-imbalance
      - e.g., determined by simulation scheduler
  - cost control
    - make vis trans flexible
      - quality-resource consumption tradeoff
    - determine acceptable cost for vis trans
  - example
    - simulation invokes visualization
    - resolution of generated vis rep used as parameter for cost vs quality



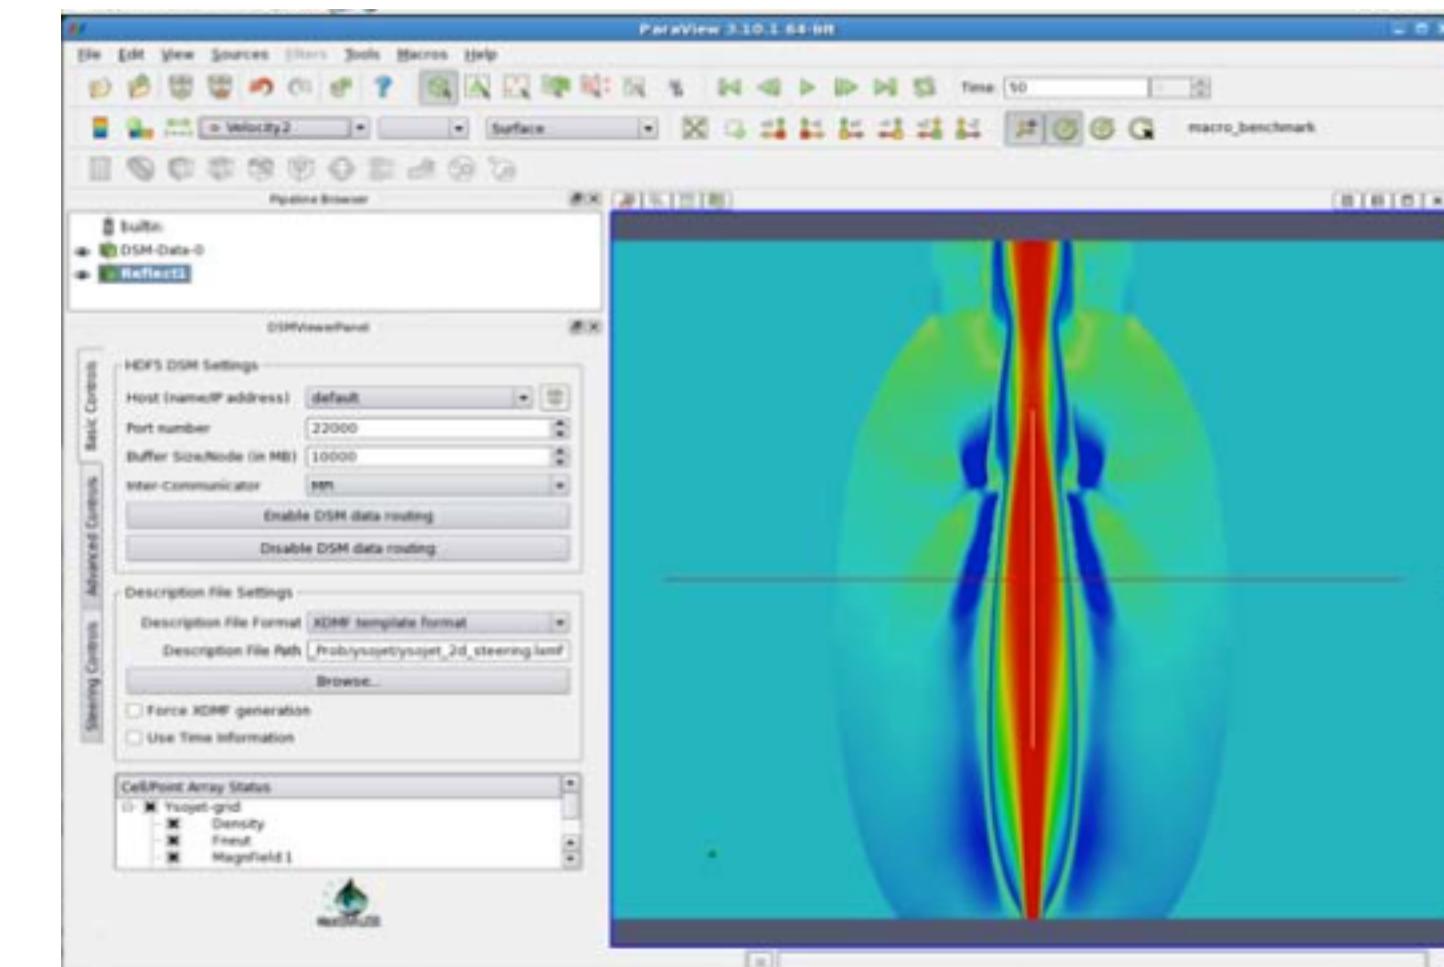
# In-Situ Visualization Challenges (3)

- *reduced/intermediate visualization representation*
  1. should be of small size, ...
  2. fast and efficient to generate, ...
  3. .... , and still represent information of interest
    - with some exploration capabilities preserved
- two main development approaches
  - start from standard visualization approach and reduce/limit degrees of freedom
  - tailor toward specific domain research question
- example (VDI, more detail in the following)
  - fix mapping of color to values (transfer function)
  - optimize for a certain view, but generally allow arbitrary camera movement (impact on quality)
- focus on this in the following

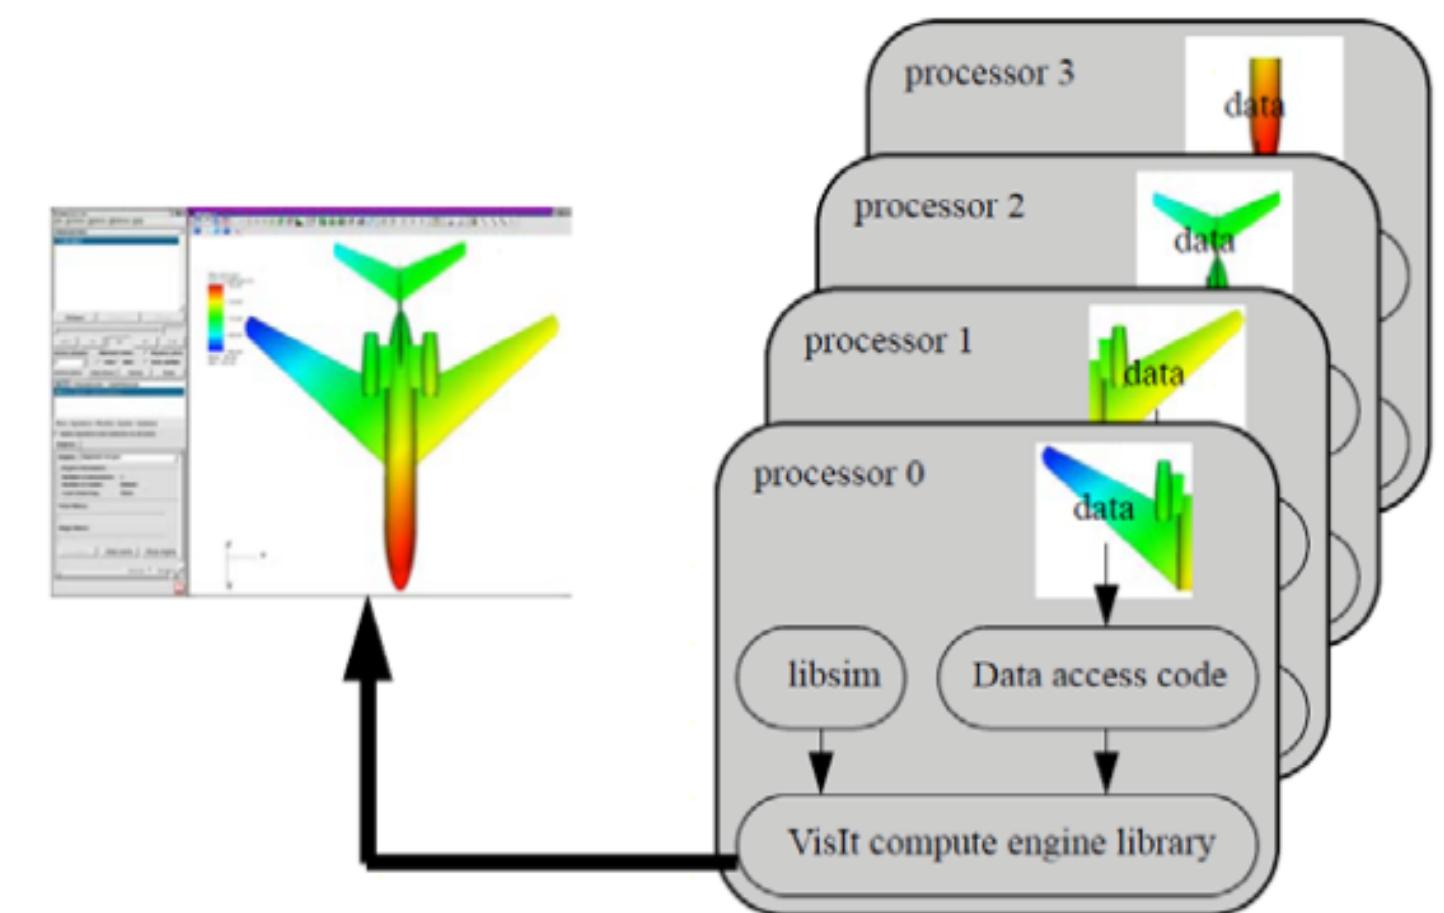


# State of the Art in In-Situ Visualization (1)

- frameworks implementing loose and tight coupling
  - Paraview ICARUS (Initialize Compute Analyze Render Update Steer)
    - loosely coupled and push-driven
      - HDF5 file in shared memory
      - no modification of simulation codes
  - ParaView Co-processing
    - simulation communicates with outside visualization tool that generates images
      - co-processing library is embedded in simulation (tight coupling)
      - graphical configuration tools for co-processing configuration
  - VisIt (libsim)
    - user needs to implement access functions required by VisIt in libsim (tight coupling)
      - simulation-specific
    - provides interfaces to VisIt



ICARUS

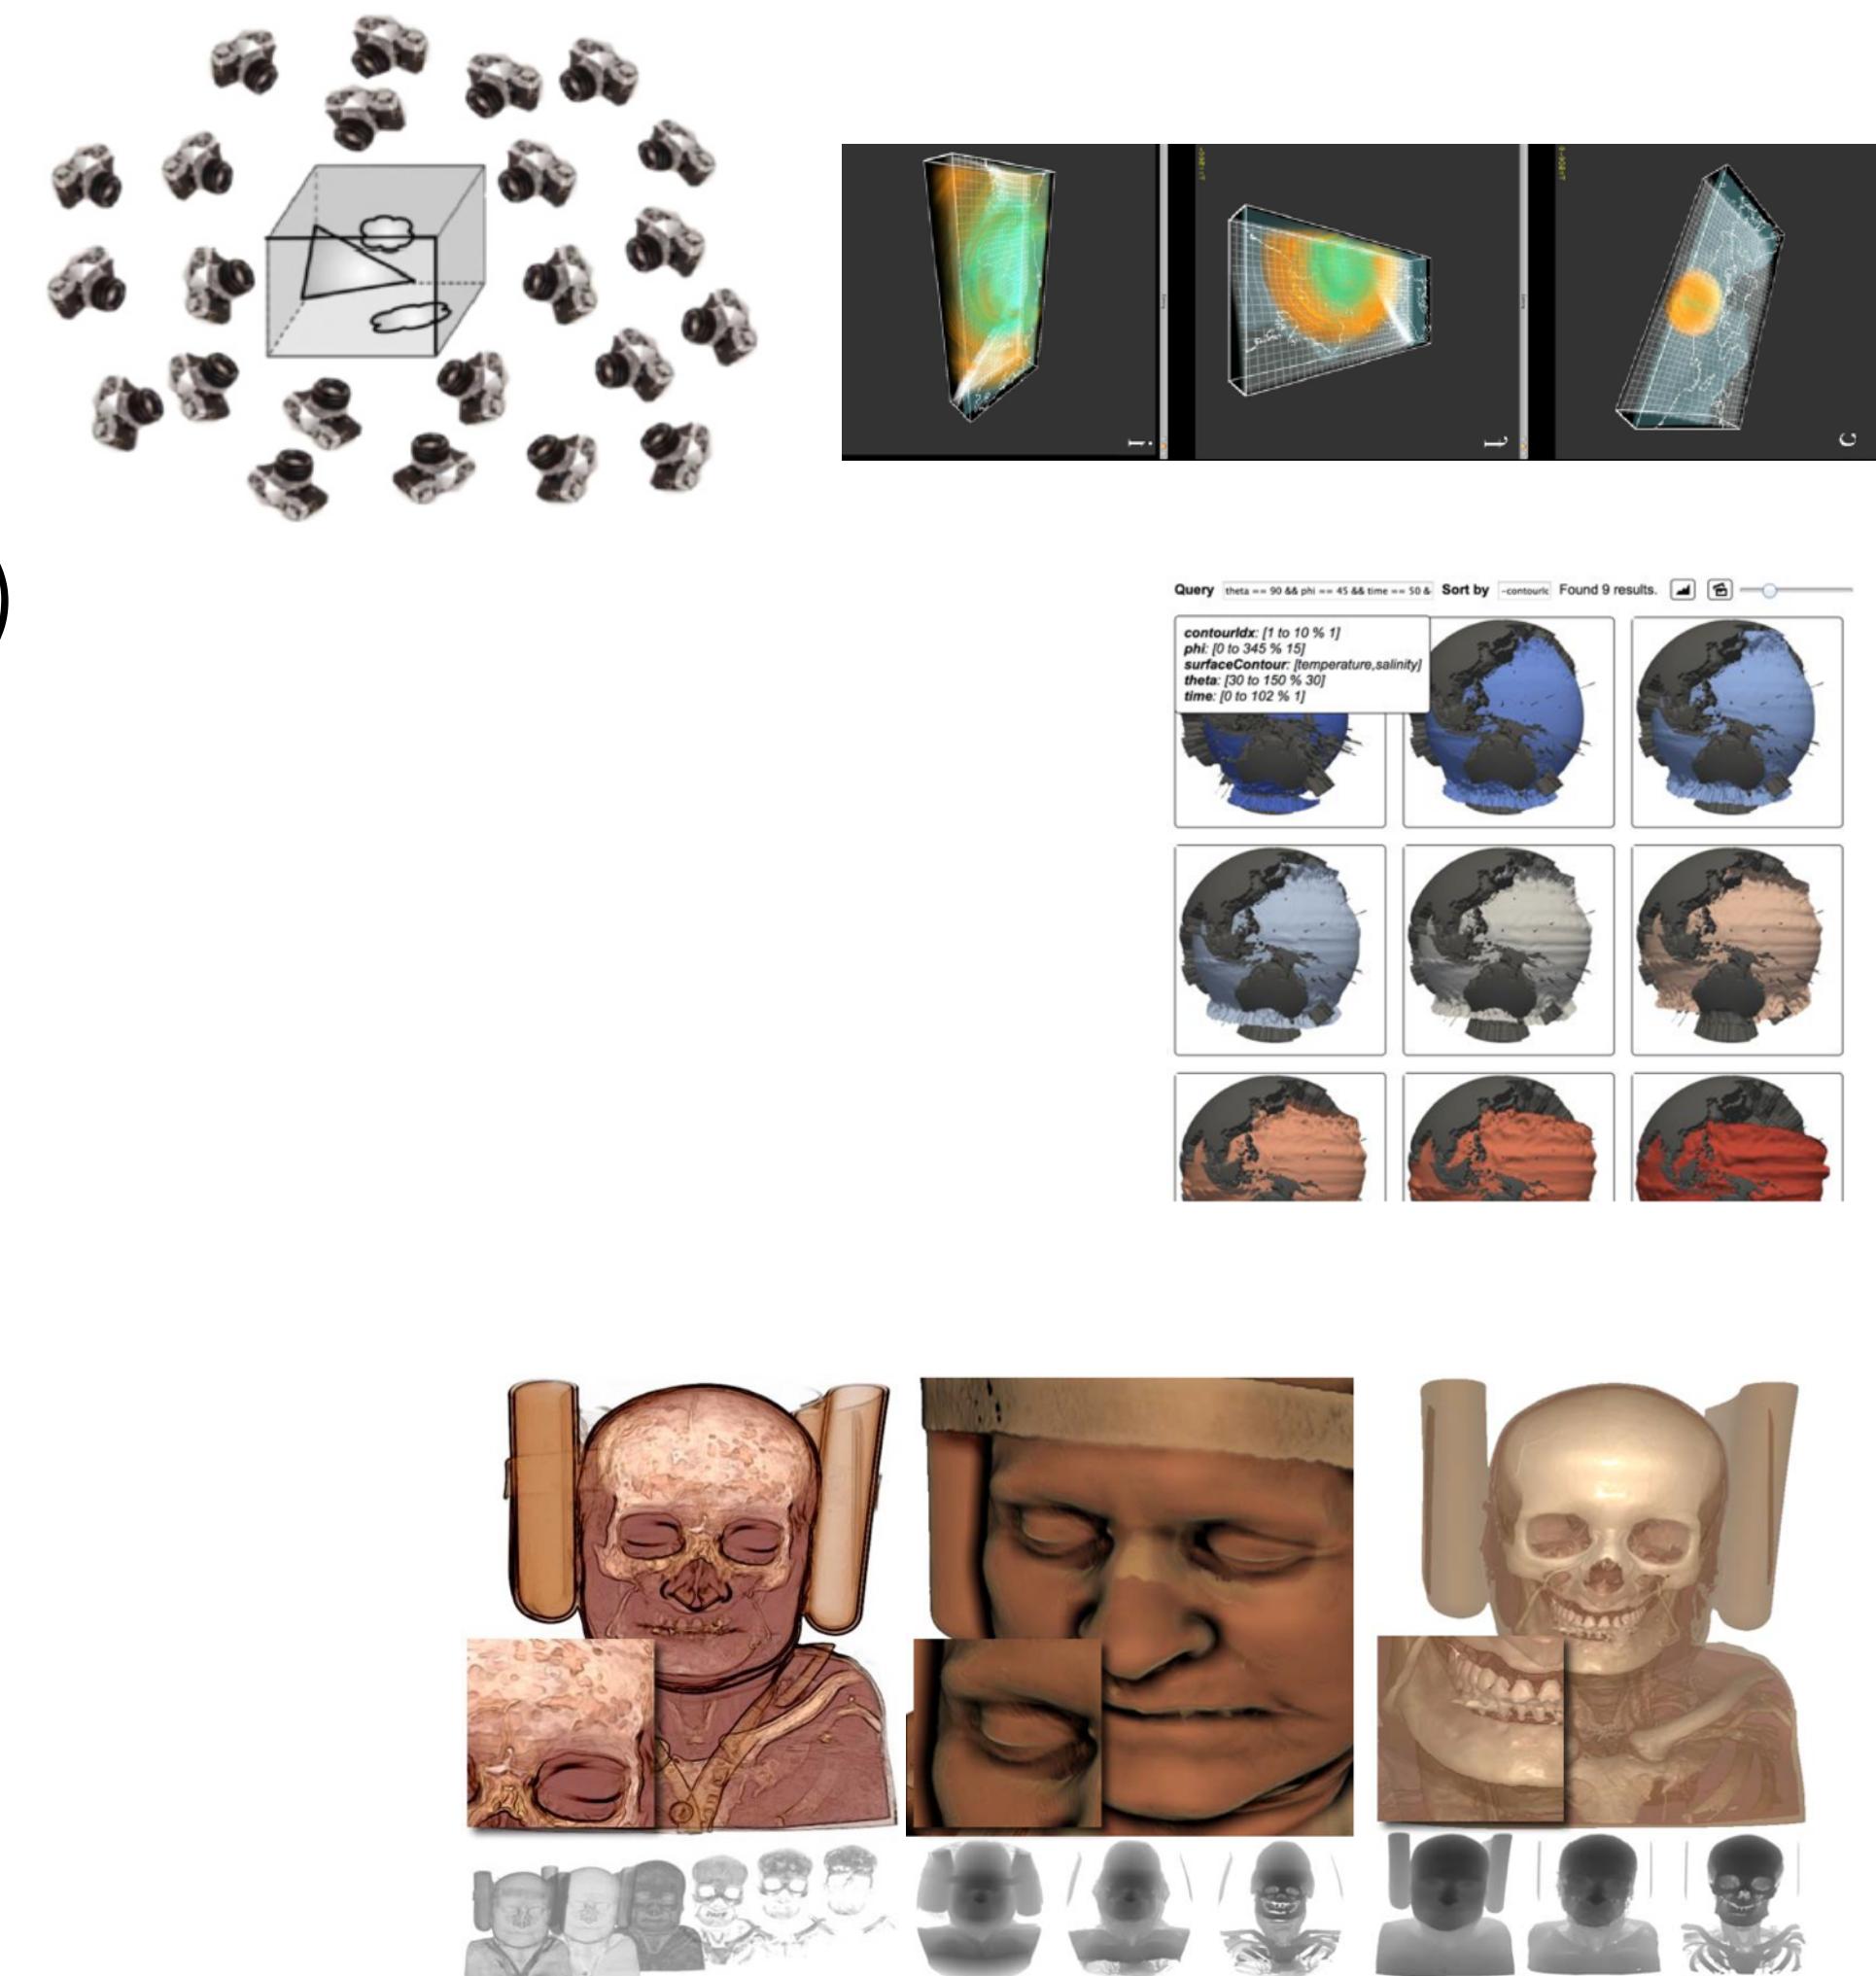


VisIt

screenshots from Rivia et al.

# State of the Art in In-Situ Visualization (2)

- *reduced representations for volume rendering*
  - interactive Viewing [Kageyama and Yamada]
    - *visrep*: record videos from different views
    - *vis*: interactive view of many movies (camera selection)
  - image data bases [Ahrens et al.]
    - *visrep*: store a set of output images directly from the simulation into an image database
    - *vis*: content querying and creation of new images by composing of individual visualization objects
  - proxy images [Tikhonova et al.]
    - *visrep*: collection of proxy images containing different information (depth, different view, etc.)
    - *vis*: mix and match images



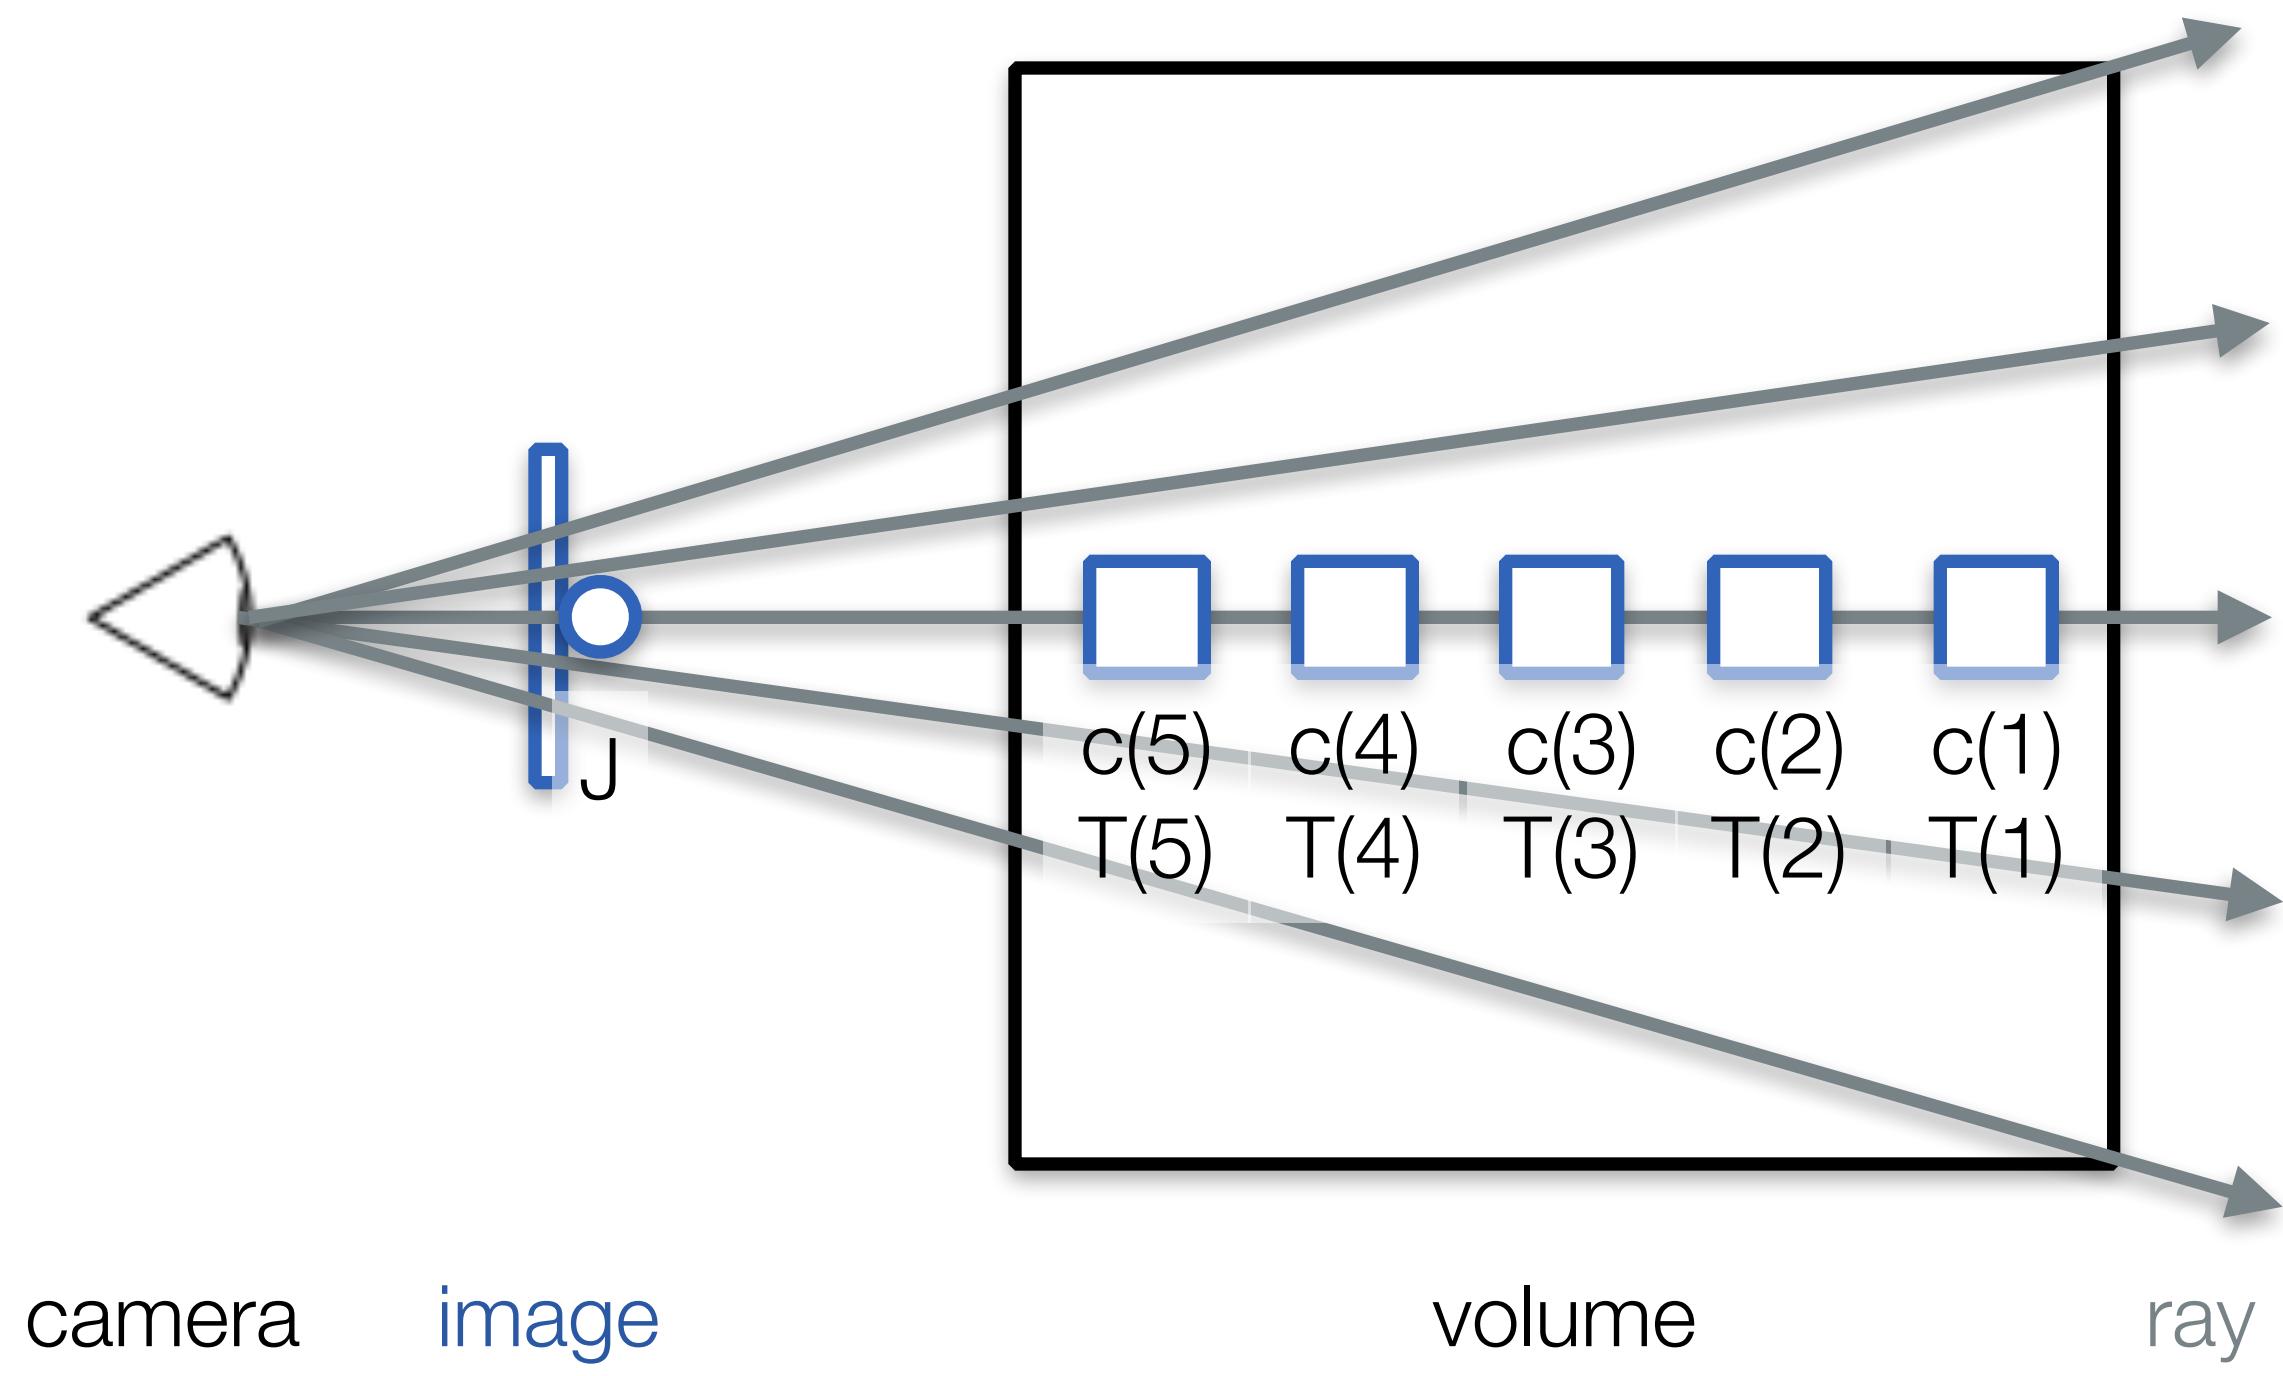
[Tikhonova et al.] Tikhonova, A., Correa, C. D., & Ma, K.-L. (2010). Visualization by Proxy: A Novel Framework for Deferred Interaction with Volume Data. *IEEE Transaction on Visualization and Computer Graphics*, 16(6), 1551–1559.

[Ahrens et al.] Ahrens, J., Jourdain, S., O'leary, P., Pratchett, J., Rogers, D. H., & Petersen, M. (2014). An Image based Approach to Extreme Scale In Situ Visualization. *SC '14 Proceedings*, 424–434.

[Kageyama and Yamada] Kageyama, A., & Yamada, T. (n.d.). An approach to exascale visualization: Interactive viewing of in-situ visualization. *Computer Physics Communications*, 79-85.

# Volumetric Depth Images

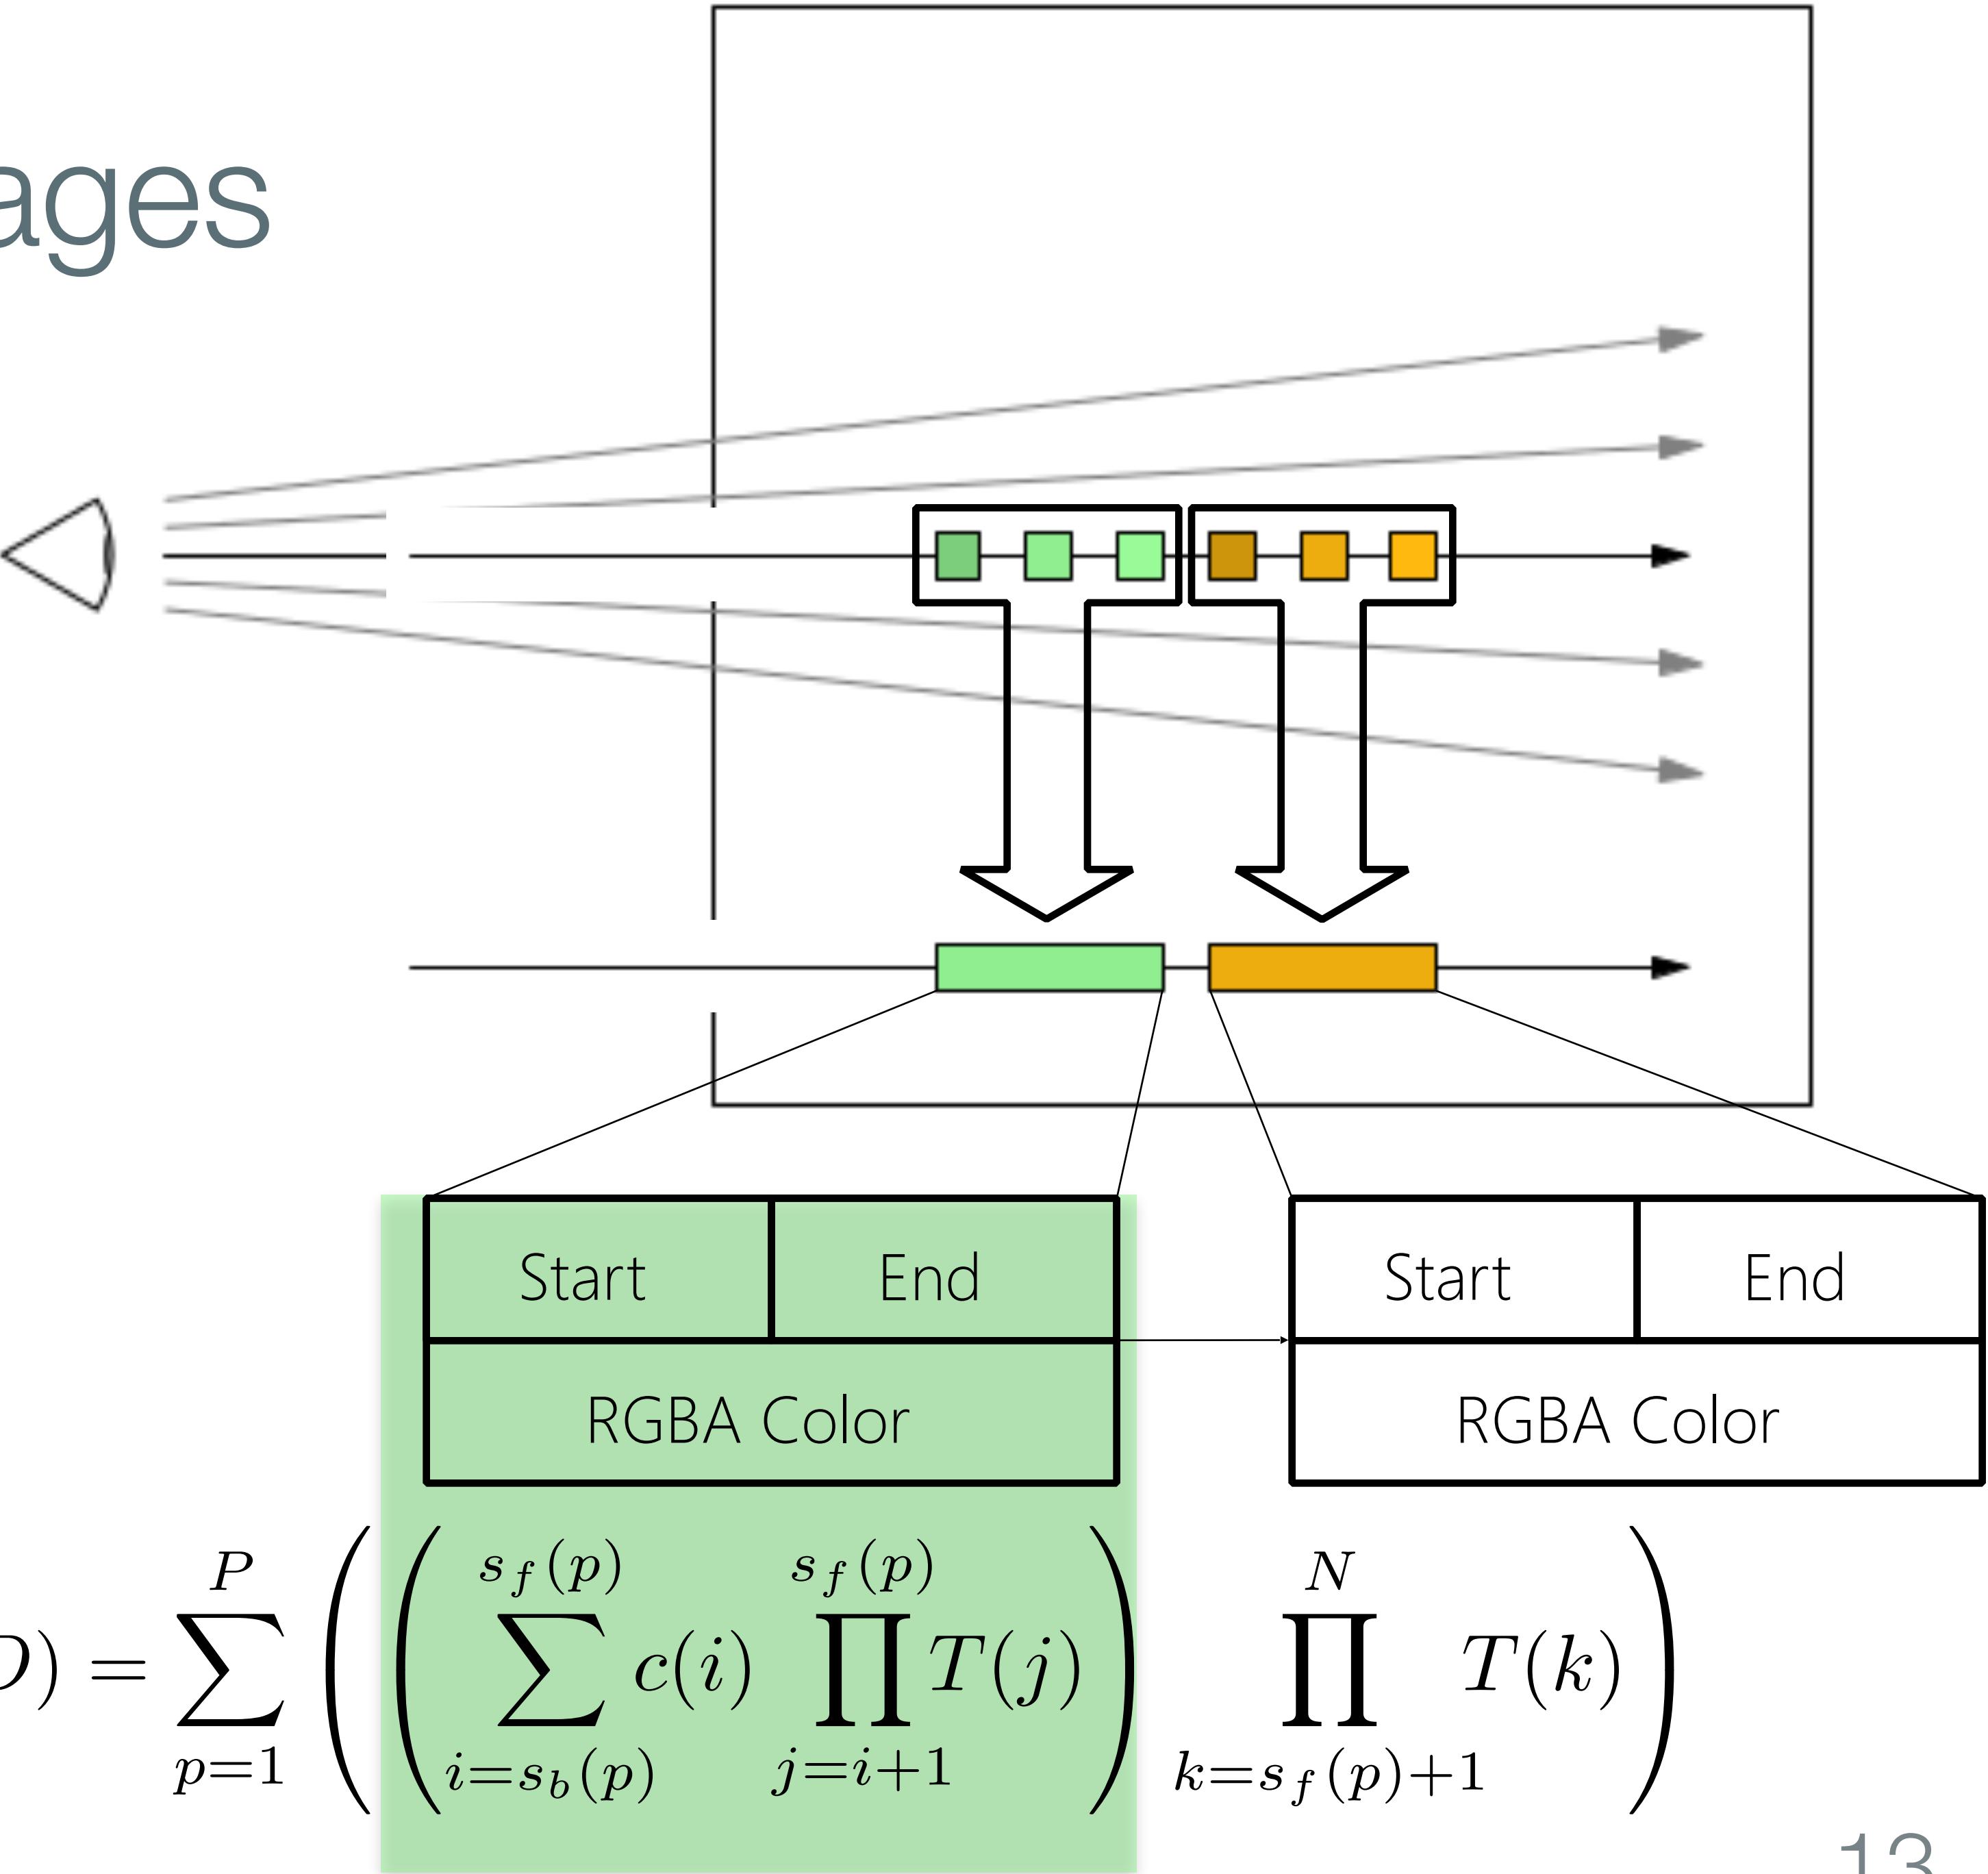
- volumetric depth images (VDIs)
  - directly based on volumetric raycasting
    - standard technique for volume visualization
  - intermediate representation from (partial) volumetric raycasting
- volumetric raycasting
  - send one ray through each pixel
  - sample along ray
  - accumulate color (c) and opacity / transparency (T)



$$J(D) := \sum_{i=1}^N c(i) \prod_{j=i+1}^N T(j),$$

# Volumetric Depth Images

- subsequent samples can be arbitrarily grouped into segments and composited
  - no loss in quality
- Enabling “hierarchical” representations
  - segments represented by RGBA + 2 depth values

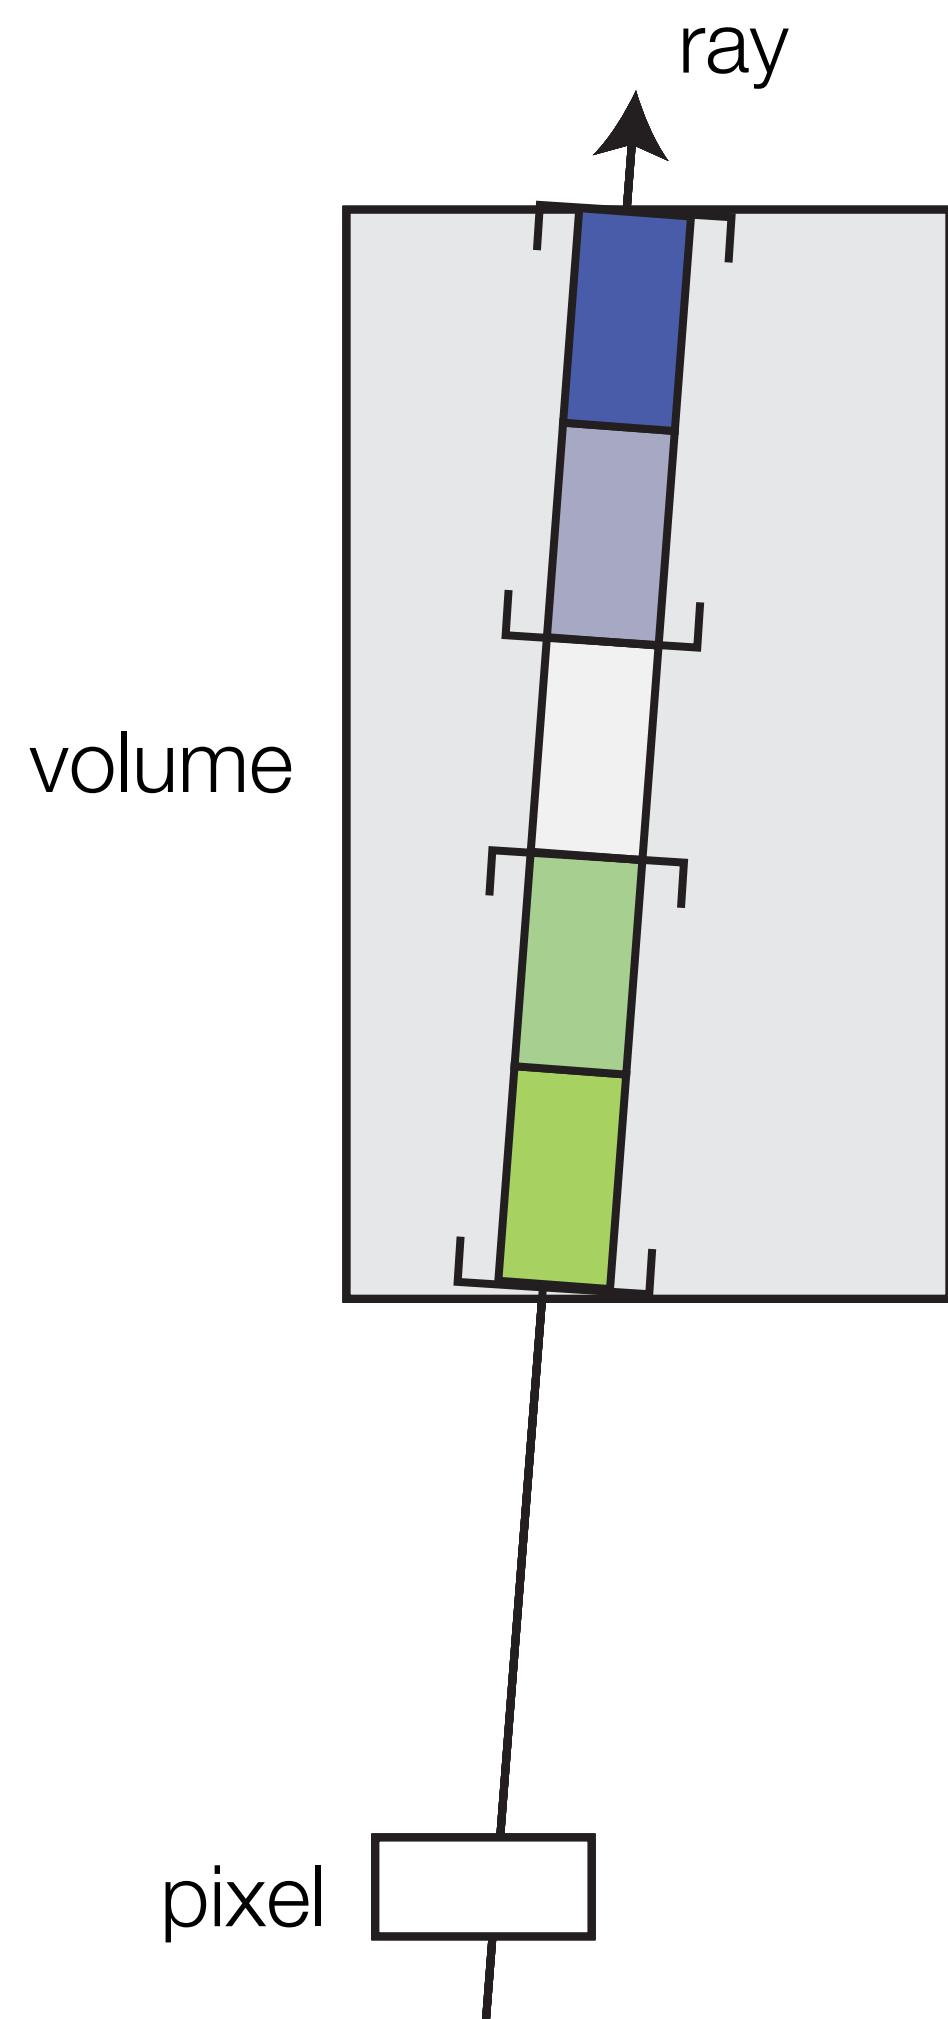


# VDI Generation

- modified front-to-back raycasting procedure
- main adjustment: color and opacity are not composited over the whole ray but only within supersegments
- a raycaster segment is merged into the current segment if ...
  - ... the adjusted color and opacity difference is below the user-provided sensitivity parameter  $\gamma$
  - ... and the sample does not represent empty space

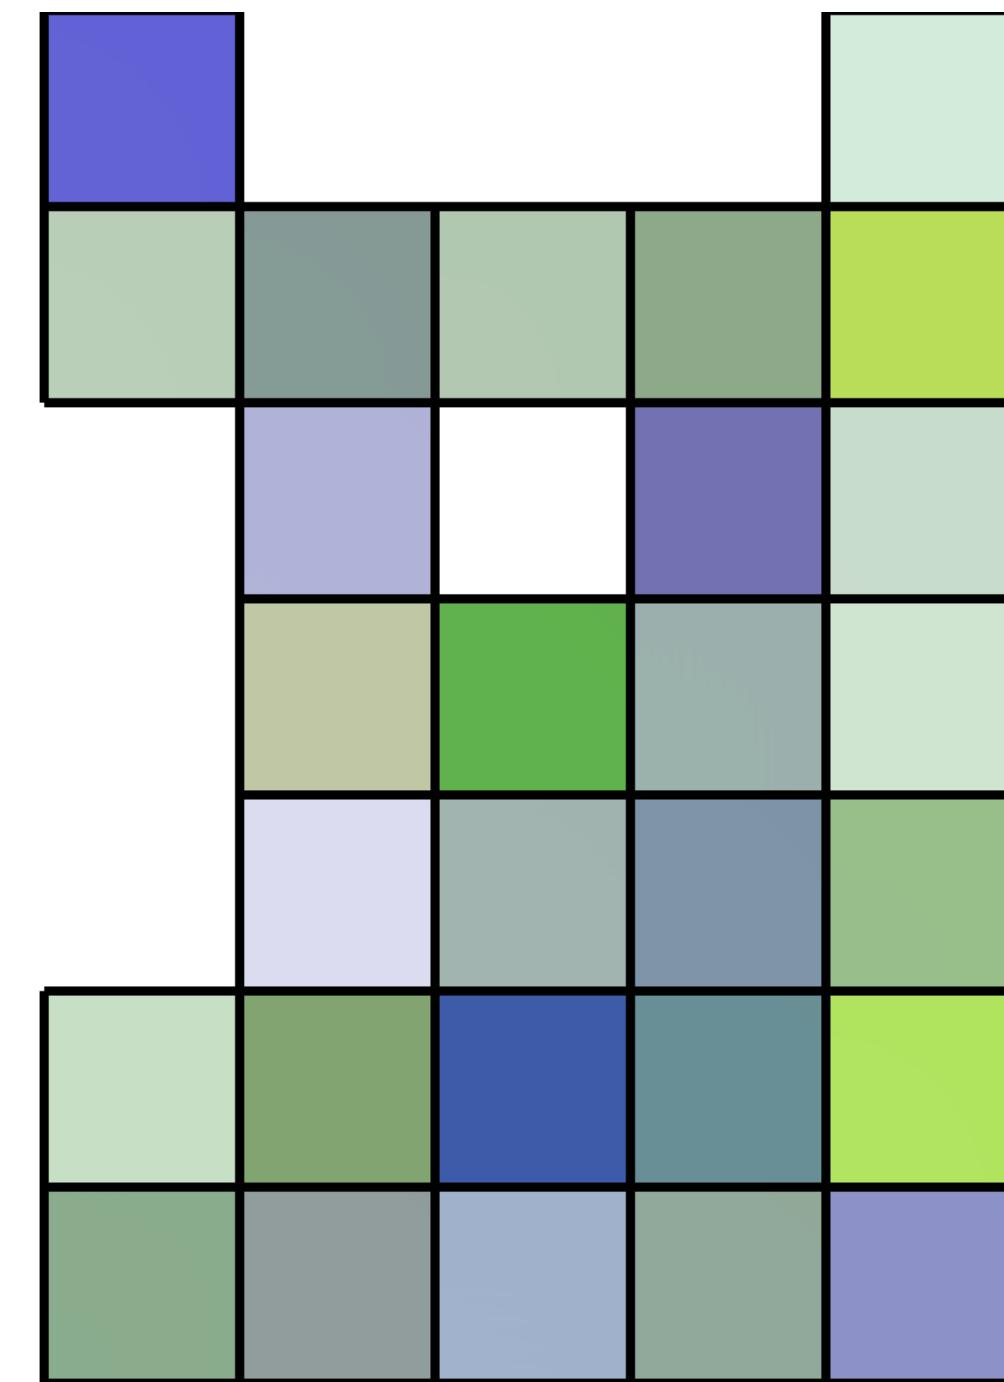
$$\gamma > |(c_{s_f}^{i-1}, \alpha_{s_f}^{i-1}) - (\hat{\alpha}(s_f, i)C(i), \hat{\alpha}(s_f, i))|$$

new segment RGBA                          current supersegment RGBA

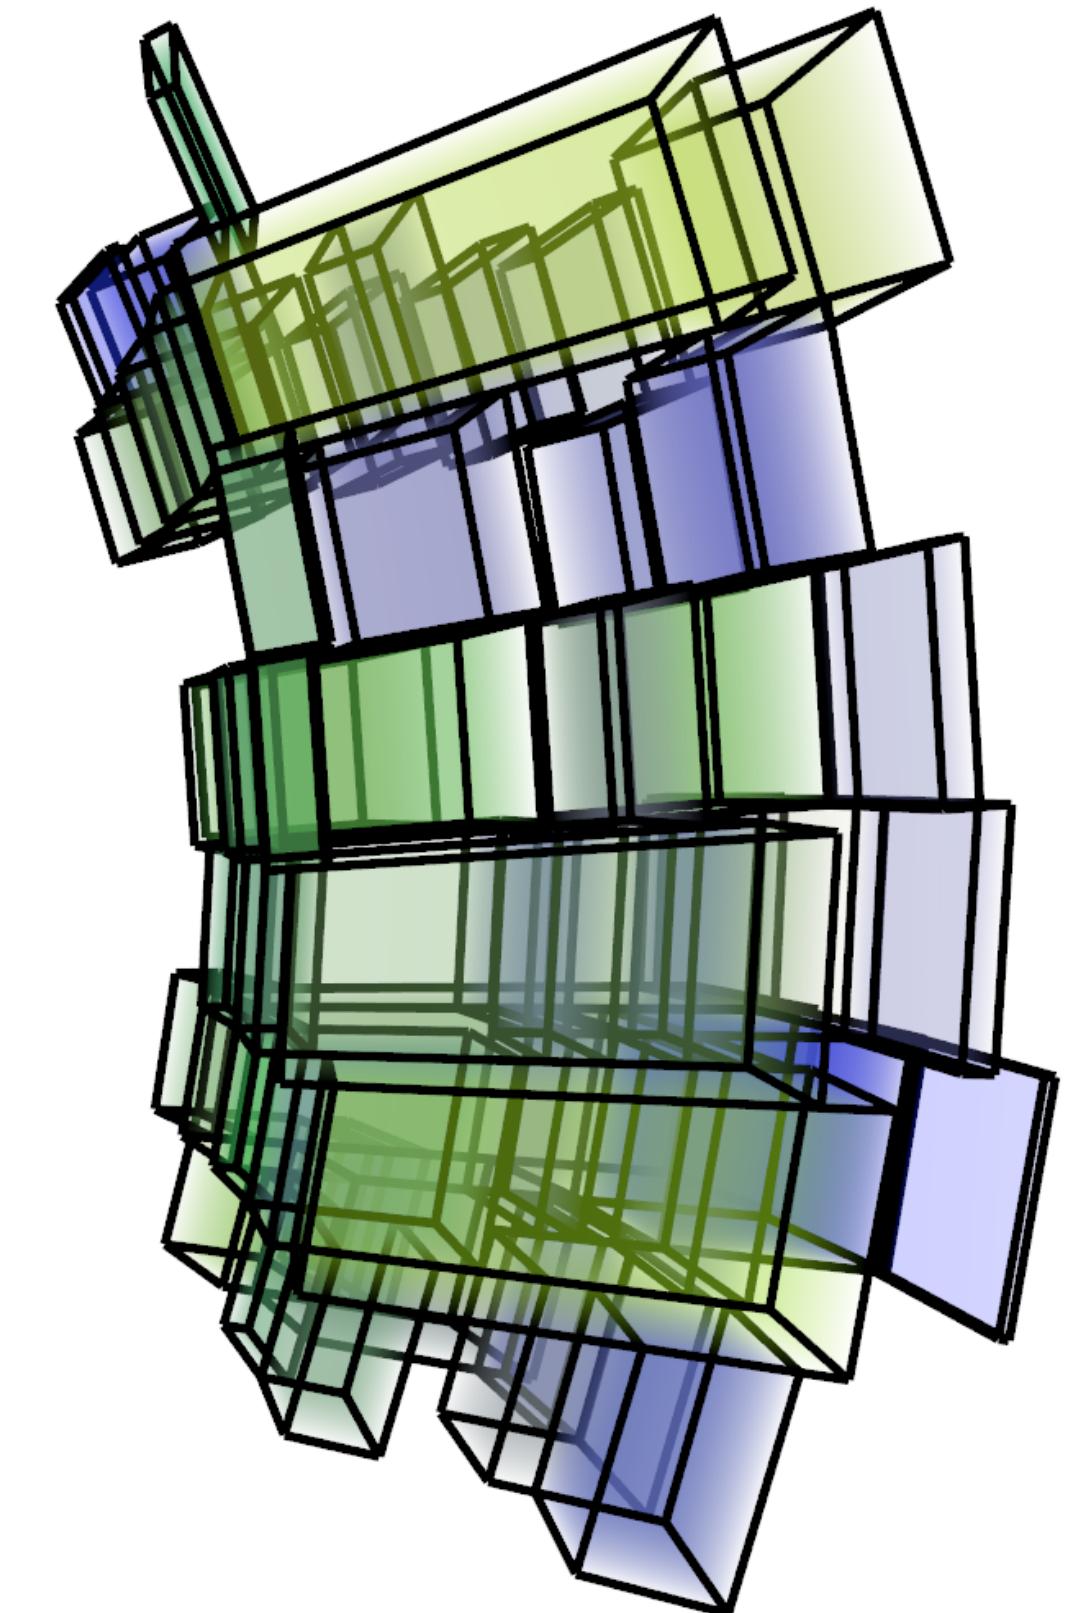


# Rendering Volumetric Depth Images

- create frustums: geometric representation of segments
  - using RGBA+depth information
  - camera configuration of original view
- rendered efficiently using OpenGL & shaders
  - alpha-blending with color and opacity from the segments
  - contribution of each frustum adjusted w.r.t. length of view ray passing through it
  - efficient depth sorting of frustums by exploiting knowledge about their creation

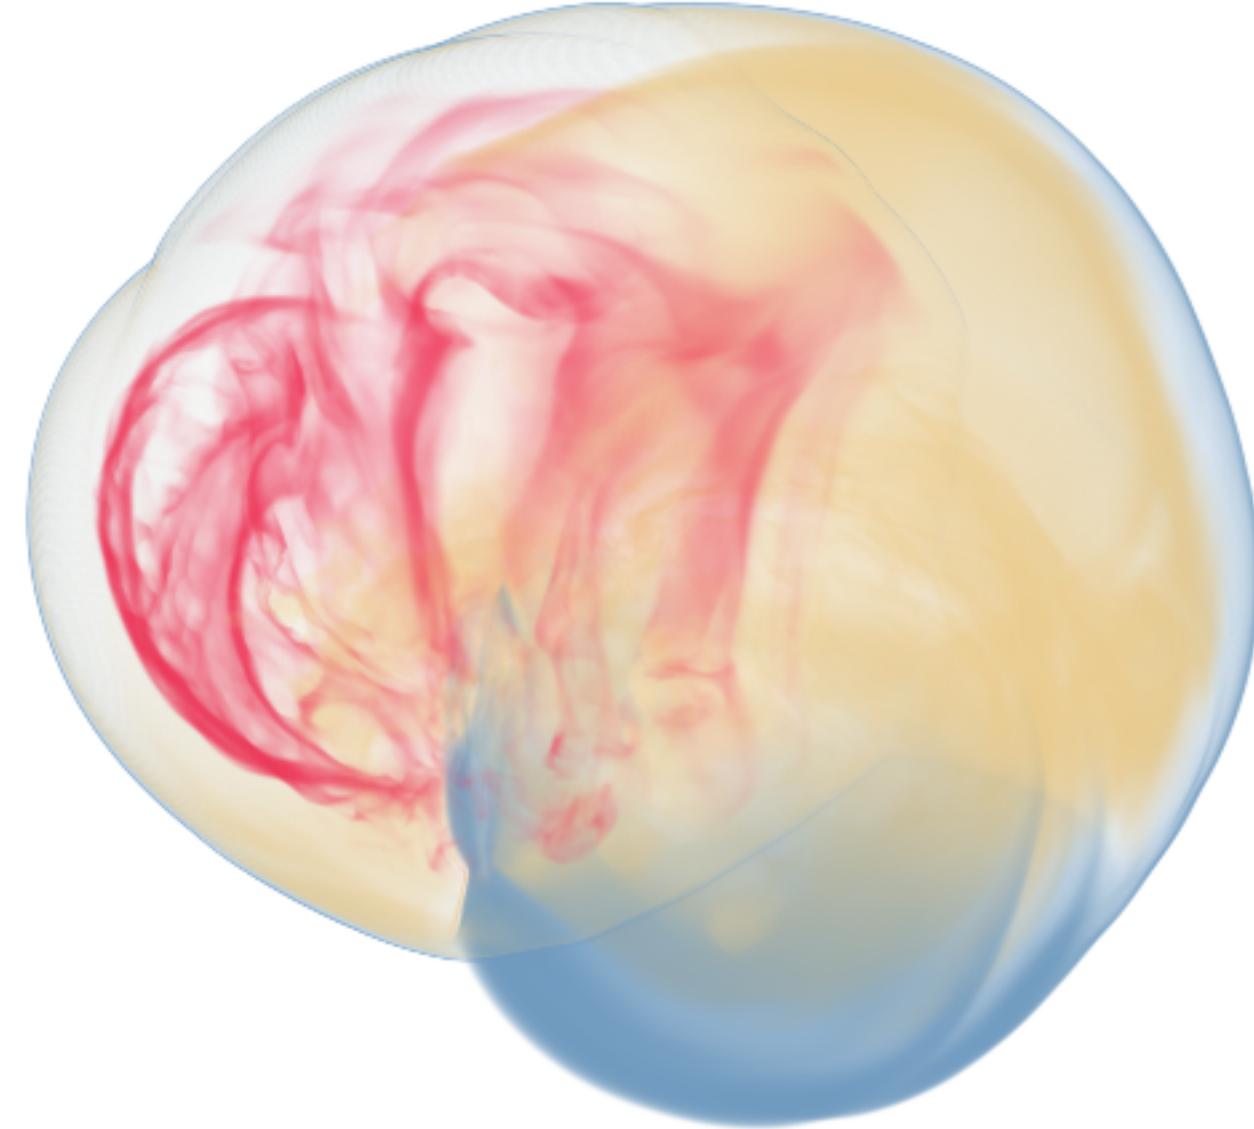


(low resolution) render representation  
(from front and 90° rotation)

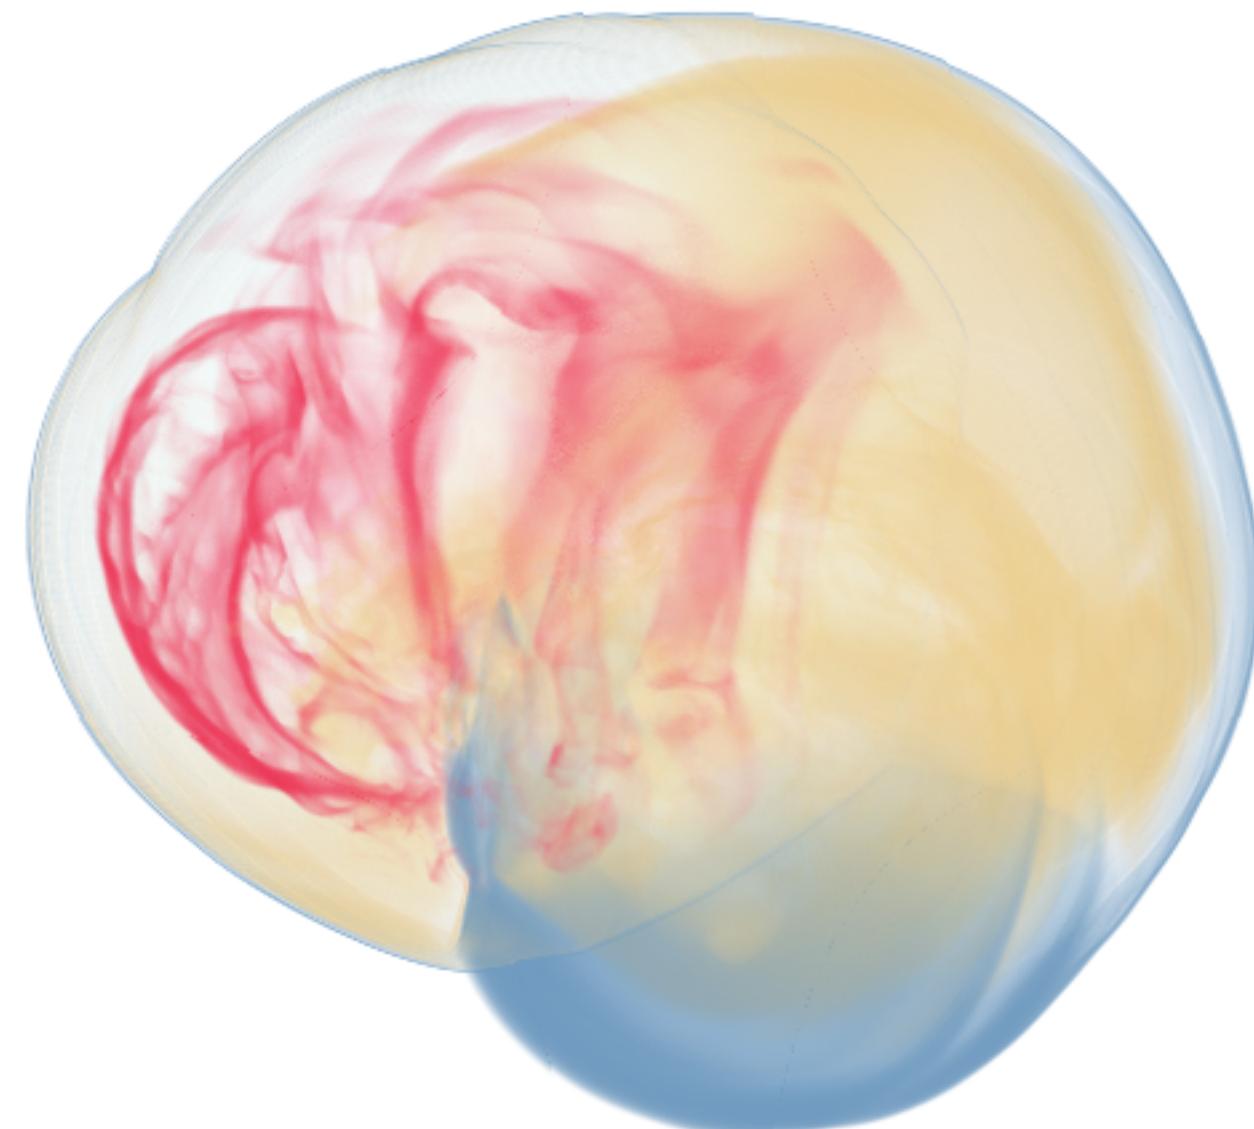


# Volumetric Depth Images

- VDI exactly matches raycasting from initial view
  - for small deviations minor visual difference ....
  - ... that increases with larger view angles
- VDI rendering performance approximately one order of magnitude faster
  - 34ms vs 275ms
- VDI generation fast
  - few milliseconds overhead to standard raycasting

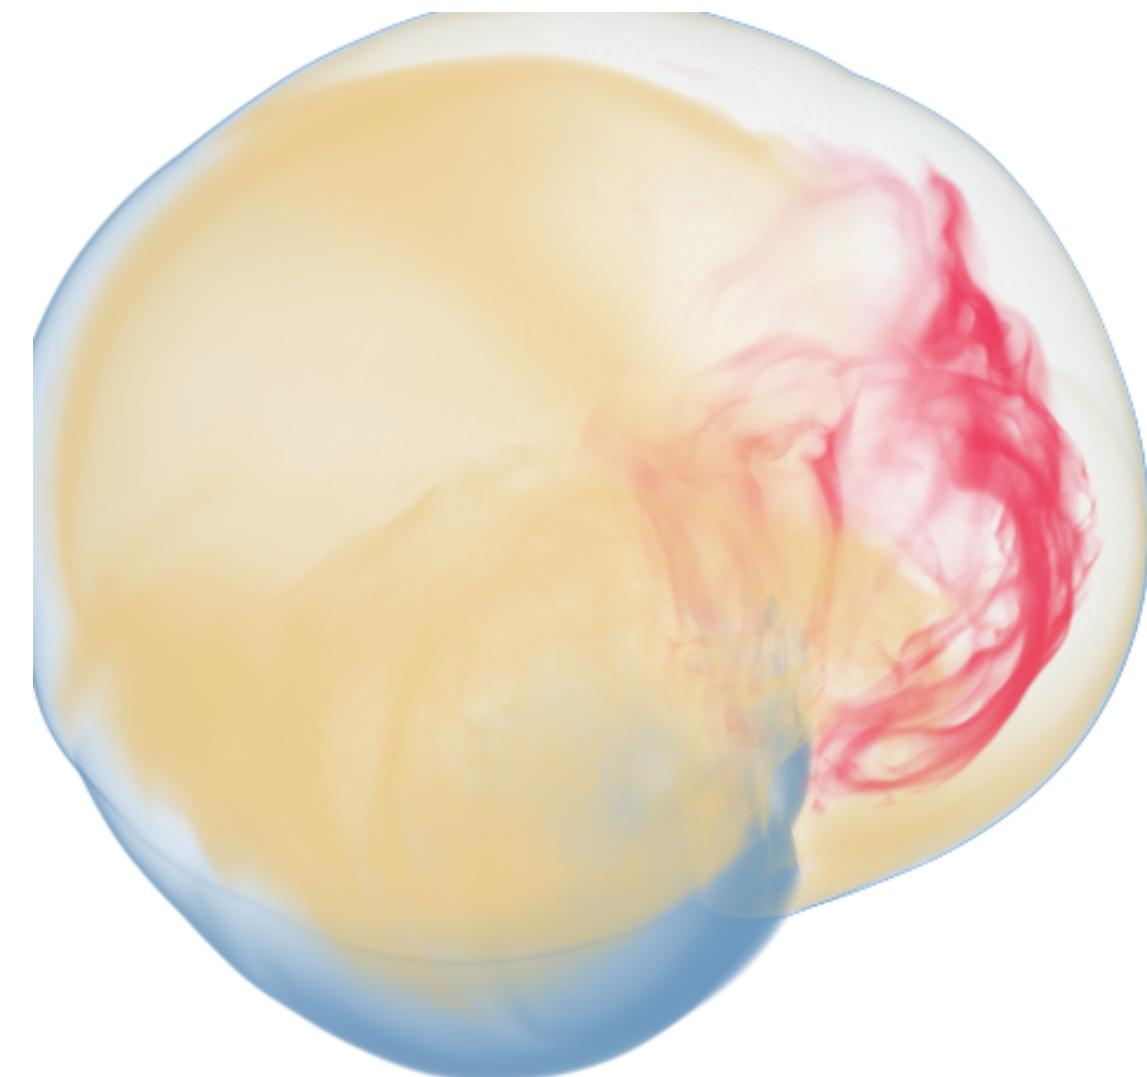
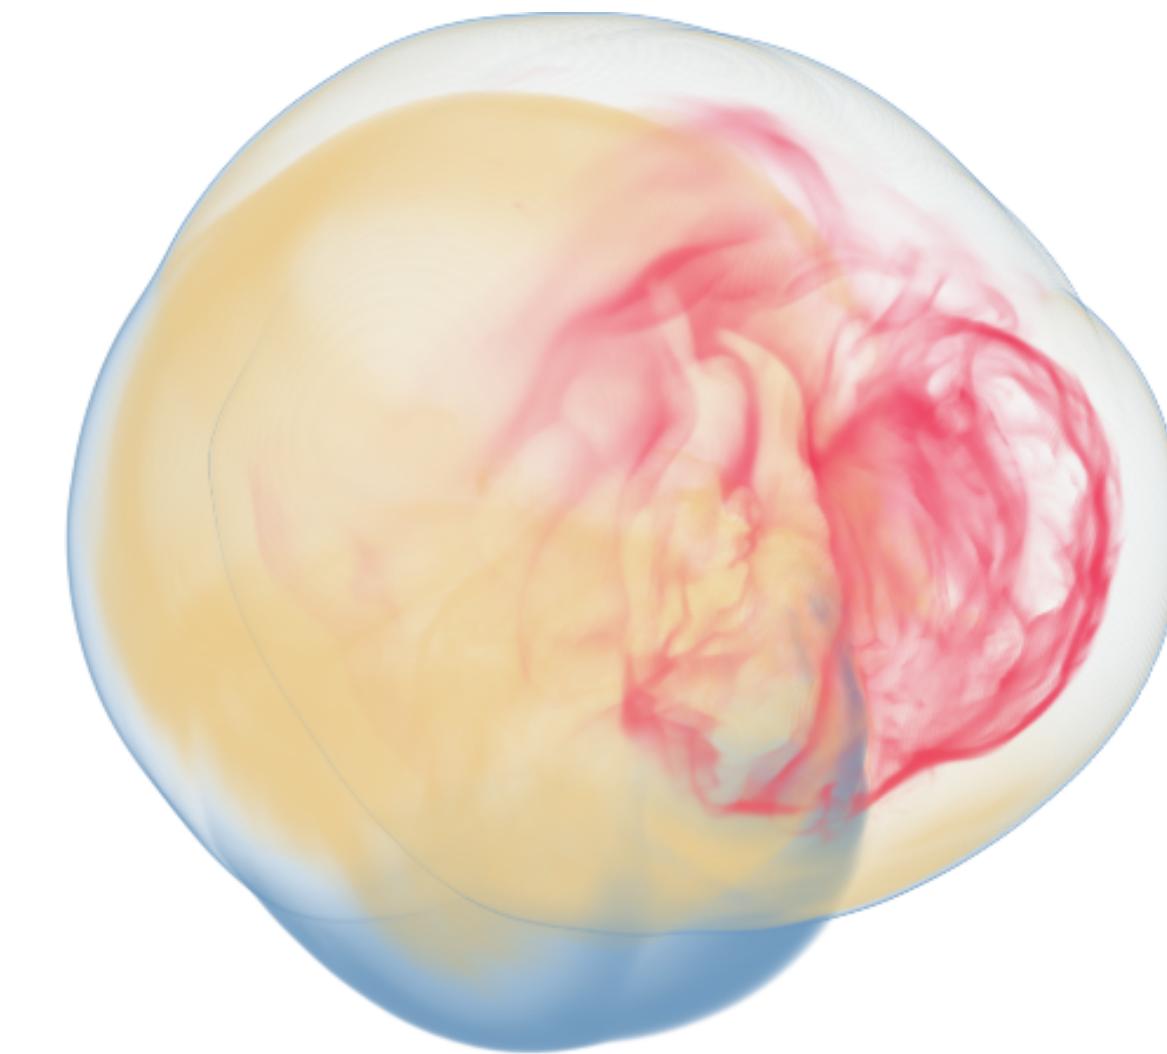
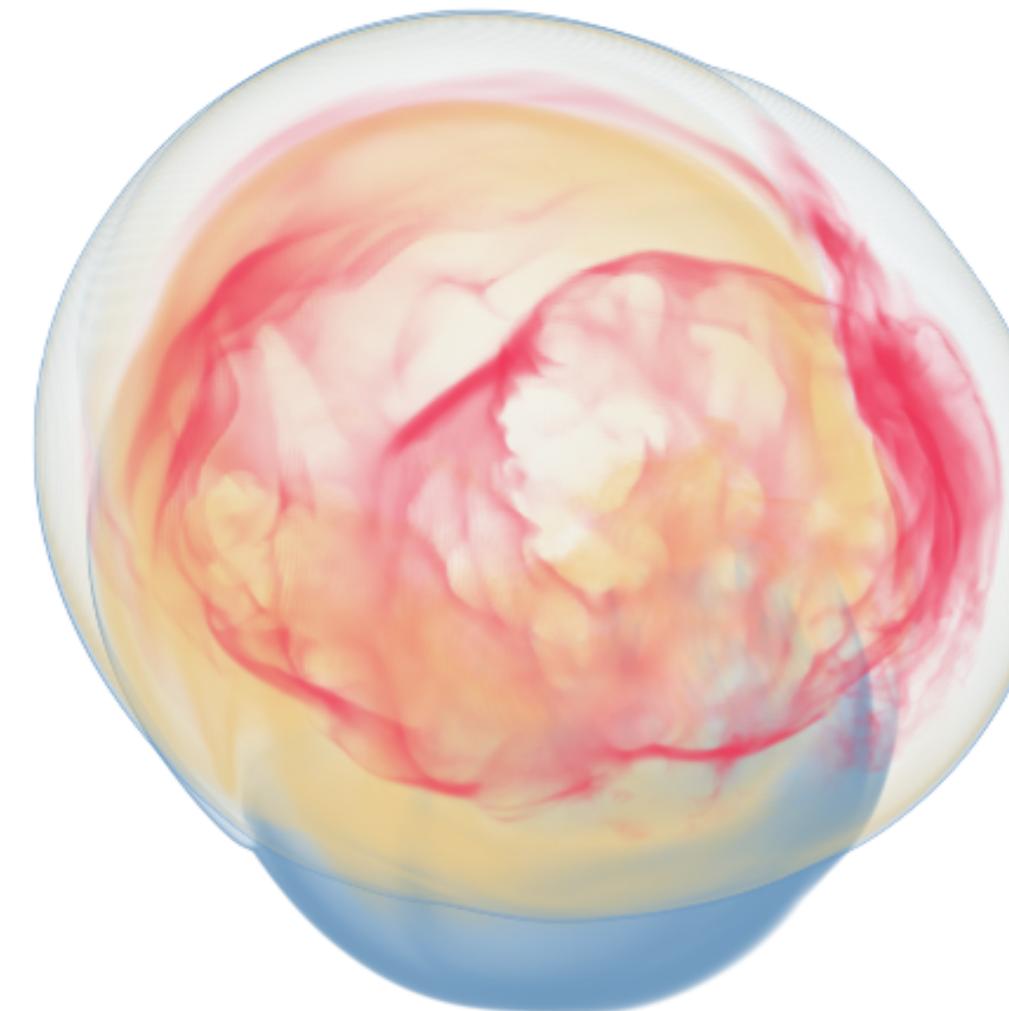
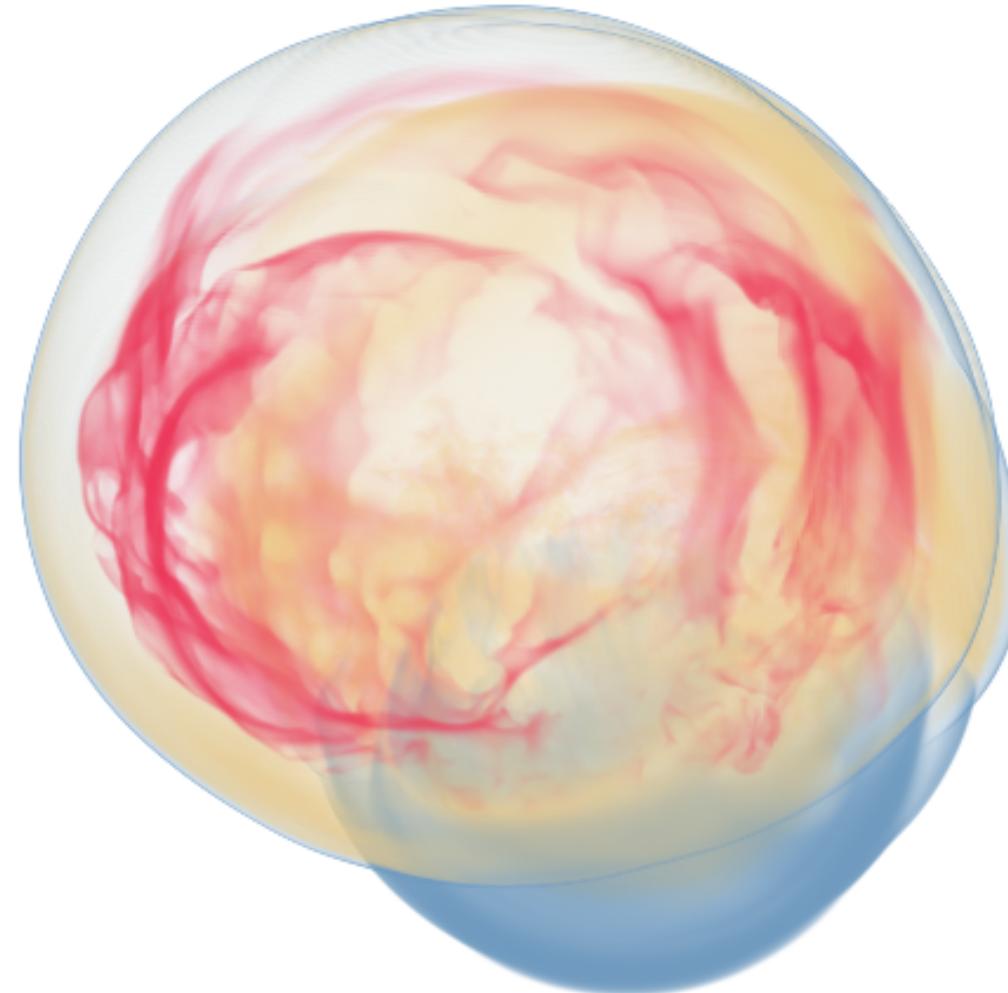


raycasting 80M cells in 275ms per frame

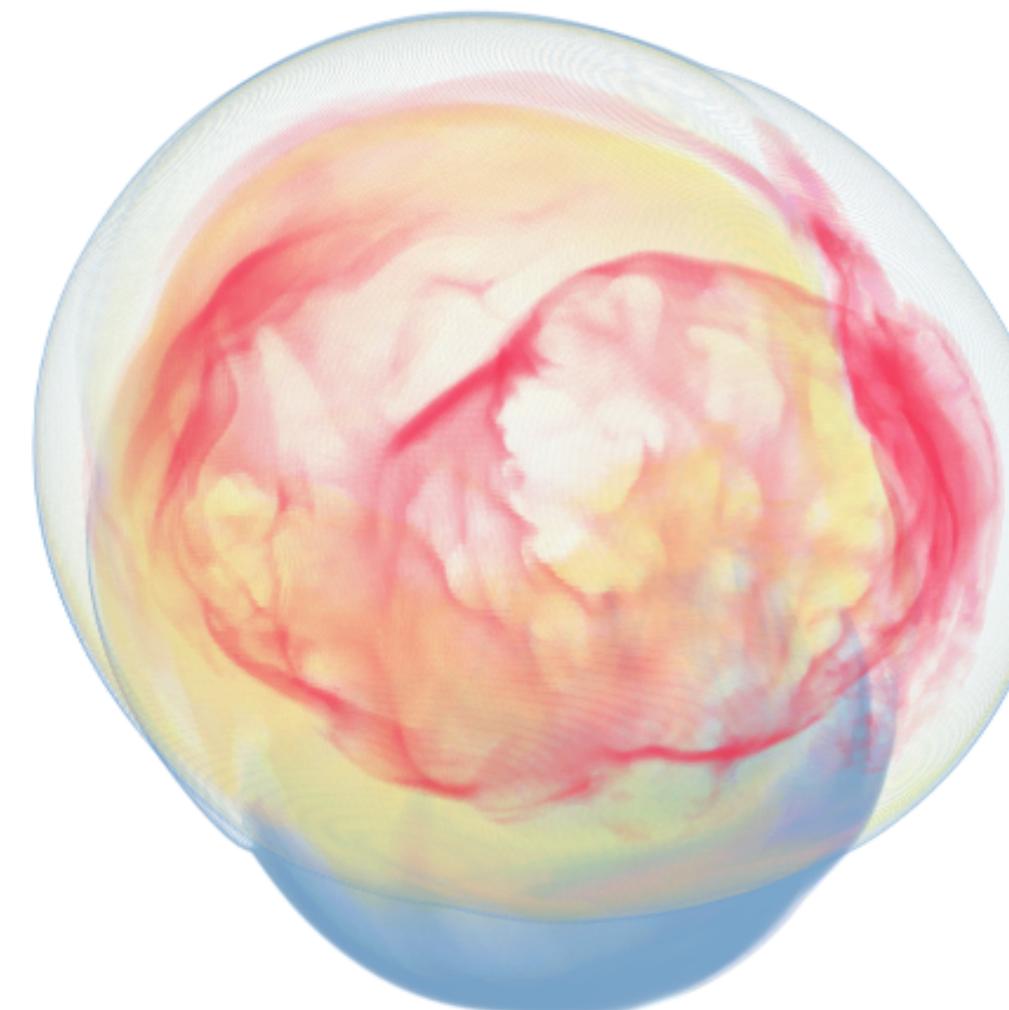
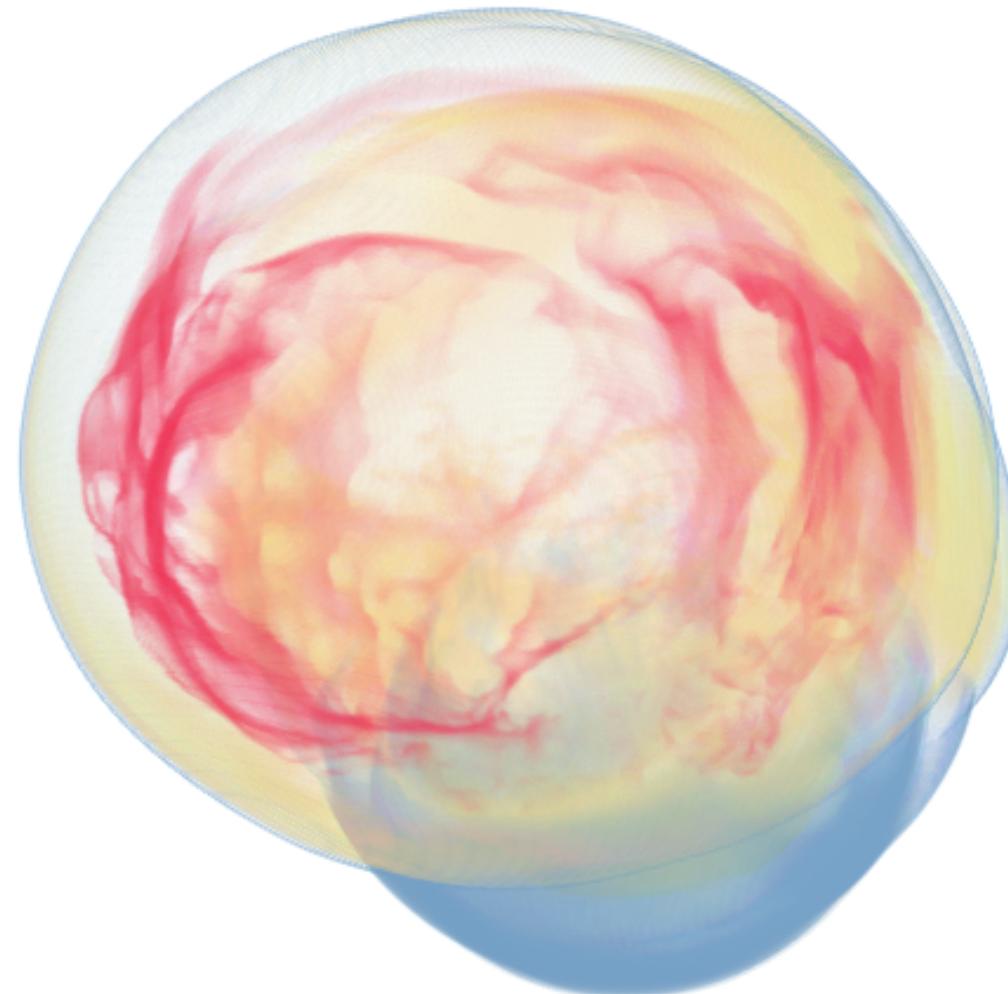


VDI rendering of 1M frustums in 34ms per frame (original view) 16

# Volumetric Depth Images



raycasting



VDI rendering with rotation  $20^\circ$ ,  $40^\circ$ ,  $60^\circ$ , and  $80^\circ$  from original view (previous slide)

# Space-Time Volumetric Depth

- so far: samples only accumulated along rays
- space-time extension: accumulate segments
  - across rays (in image space) and across time
  - store depth values and color more efficiently
    - (spend more time and accept some error for a further reduction in data size)
- procedure
  - determine local neighborhood in space and time
    - segments uniformly arranged (simple depth comparison)
  - defines connectivity graph of all segments
  - cluster regions of color
    - quantize depth values and color values
    - additional delta encoding of depth values
  - lossless compression of results

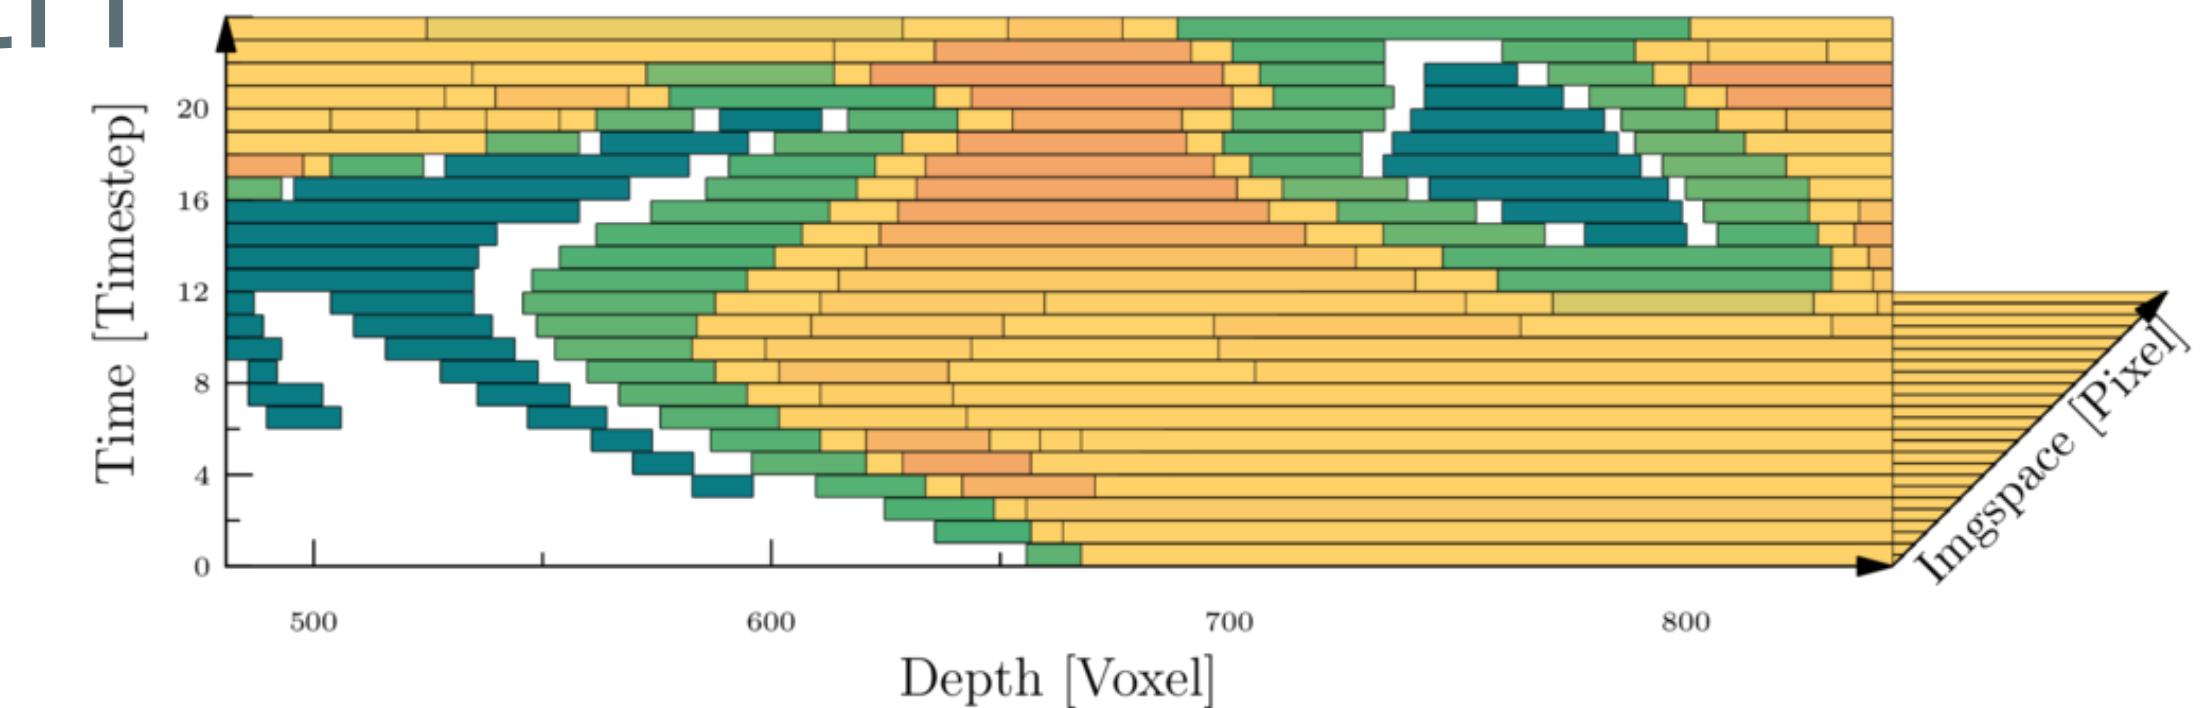
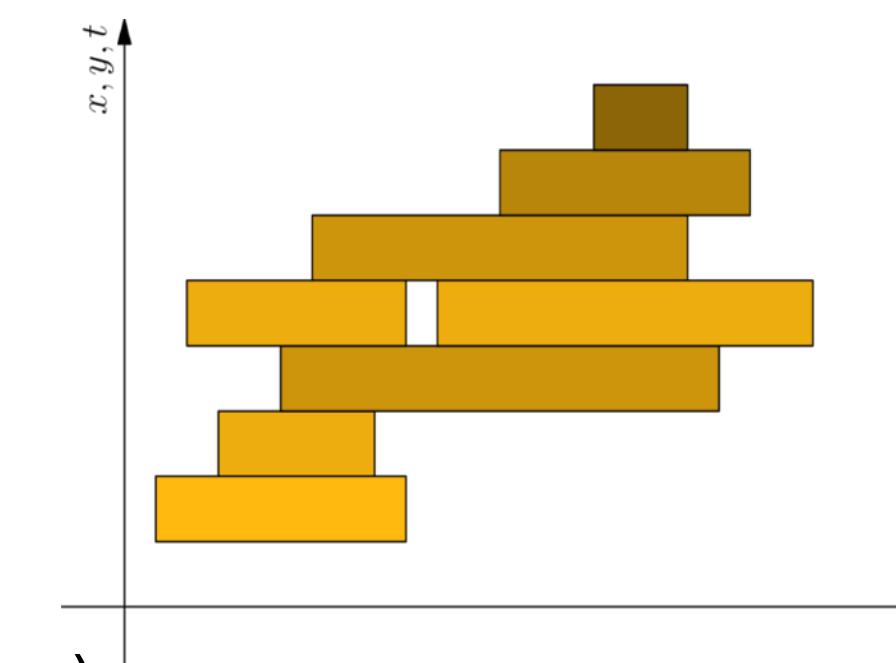
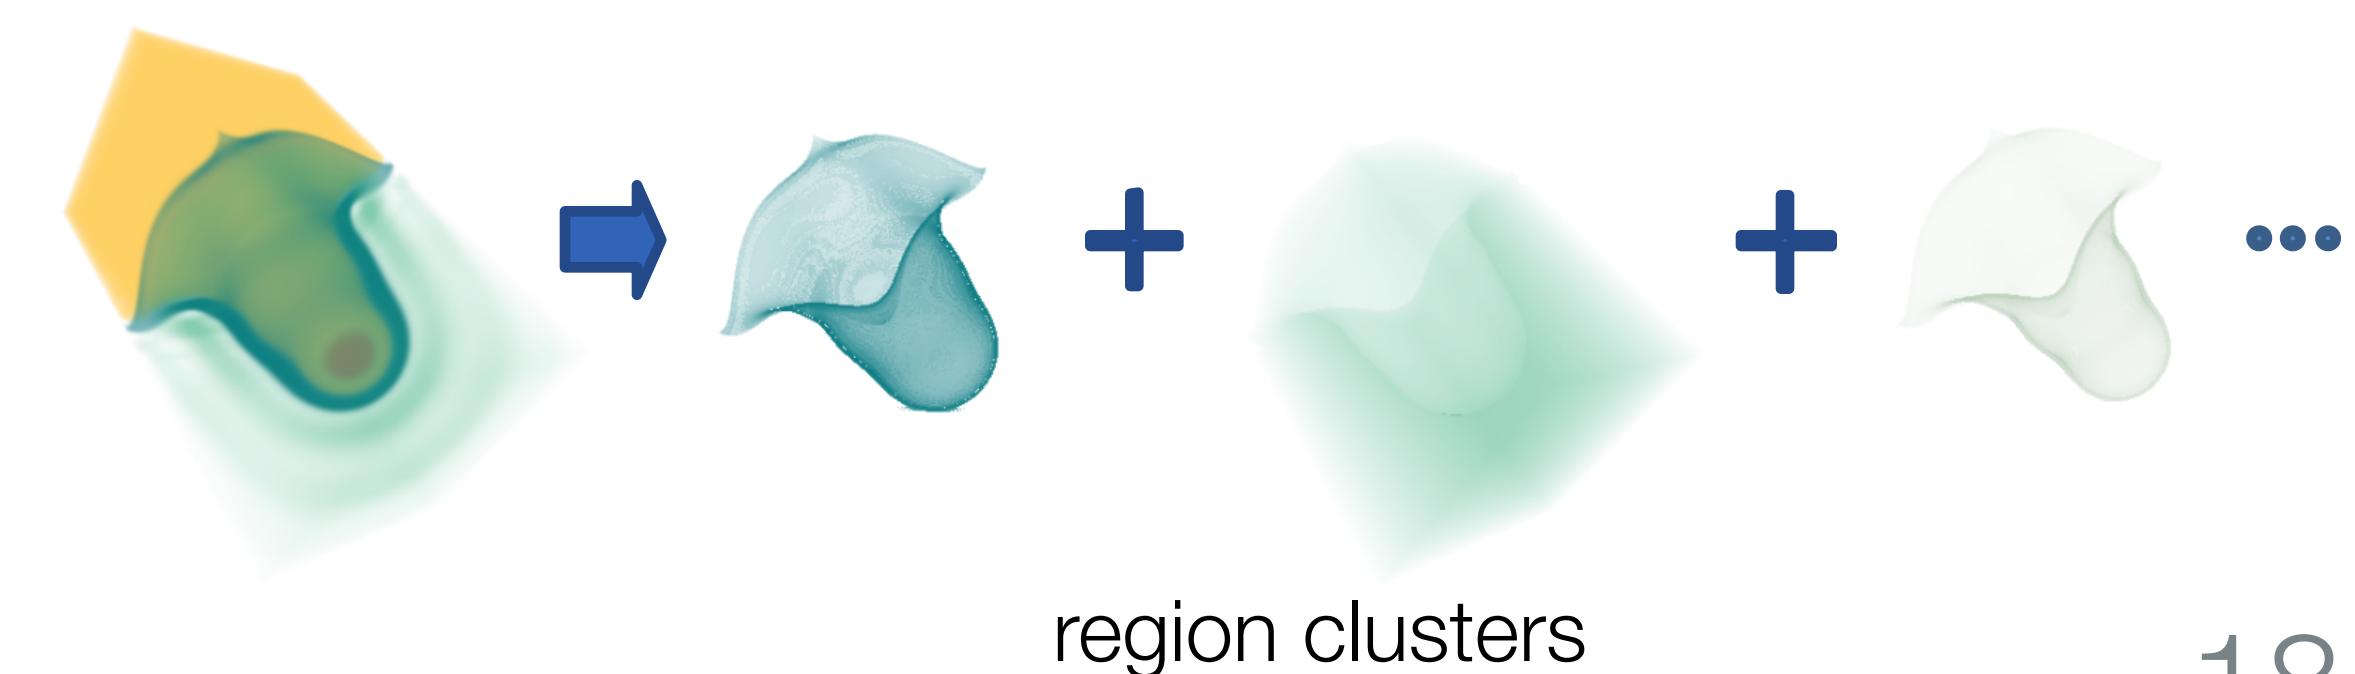
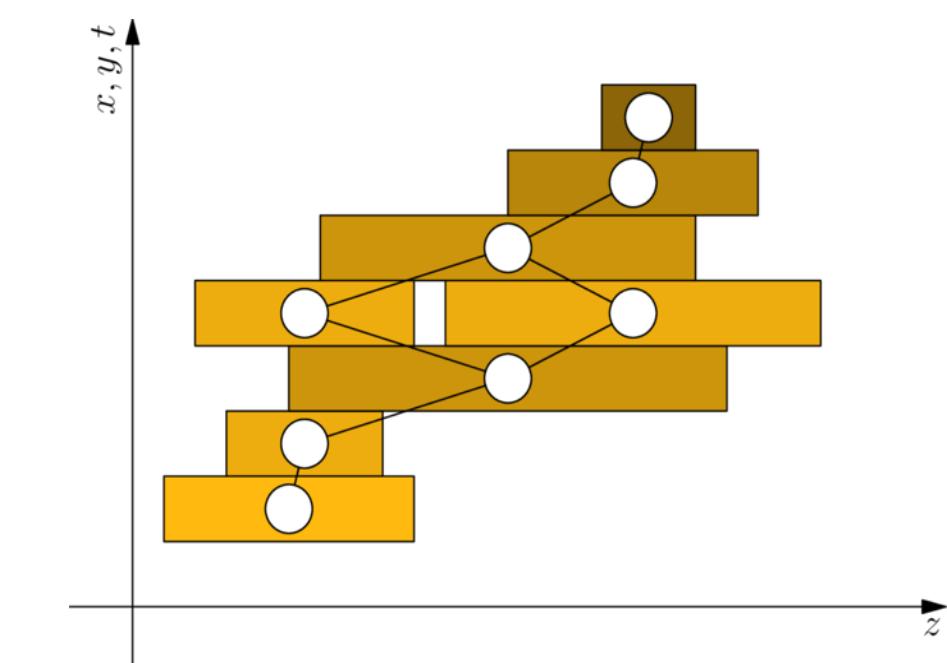


illustration of extracted segments



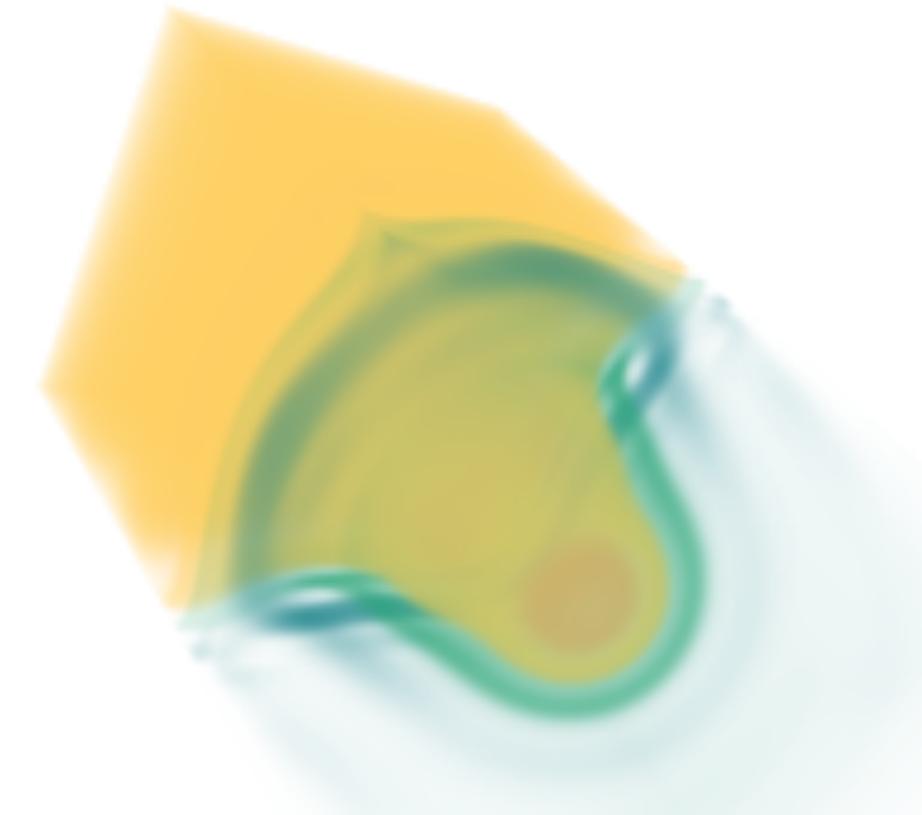
graph of segments



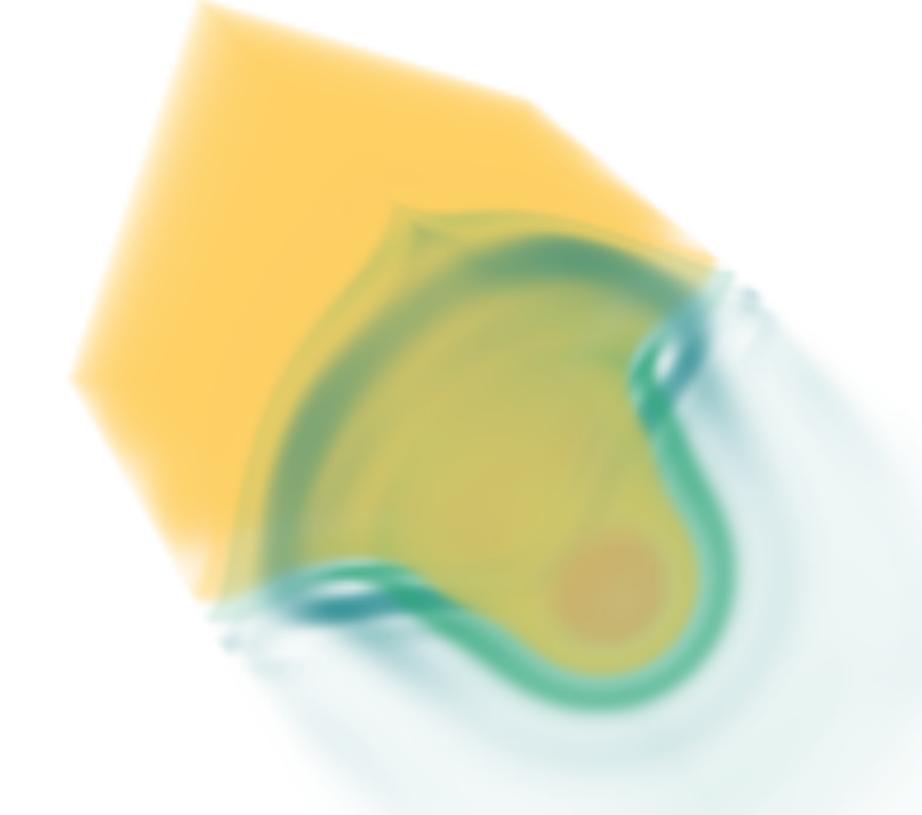
region clusters

# Space-Time VDI Results

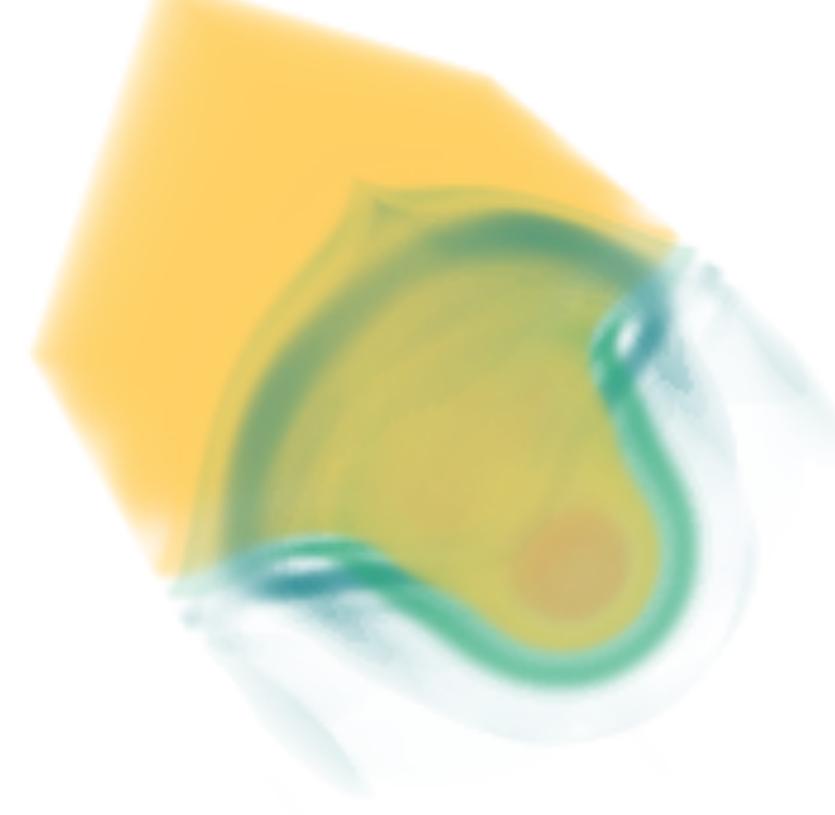
- significant reductions can be achieved over standard VDIs
- how much depends on data set and parameters
  - quantization of depth values
  - color threshold



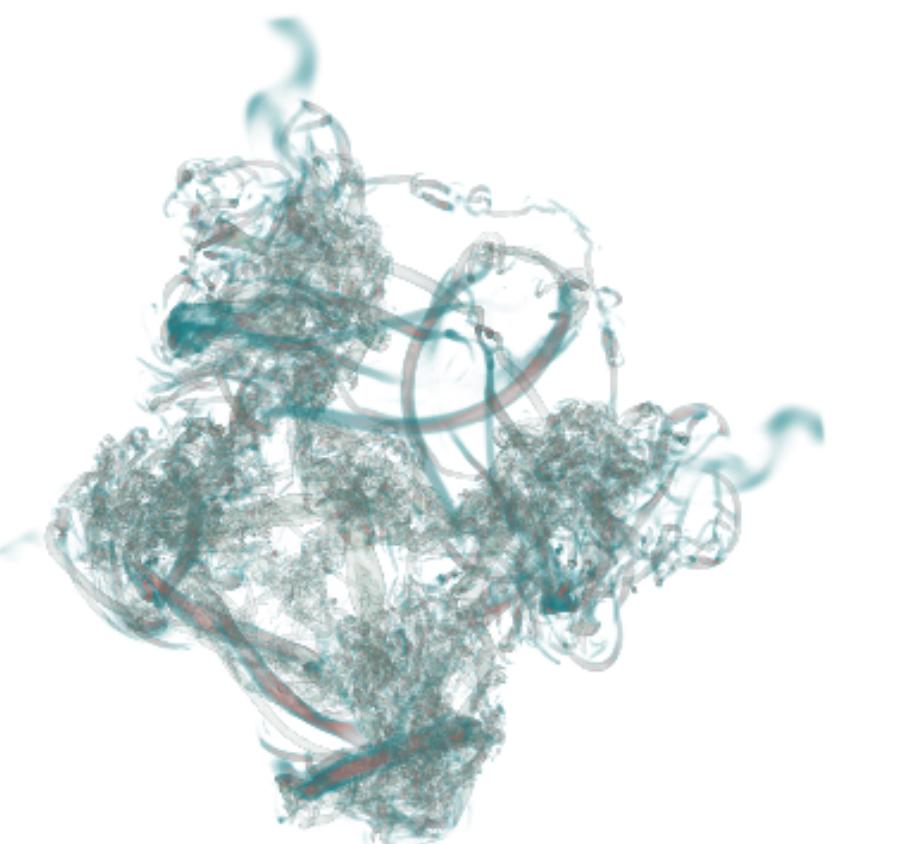
Original size  
~85.0 MB  
(raycasted)



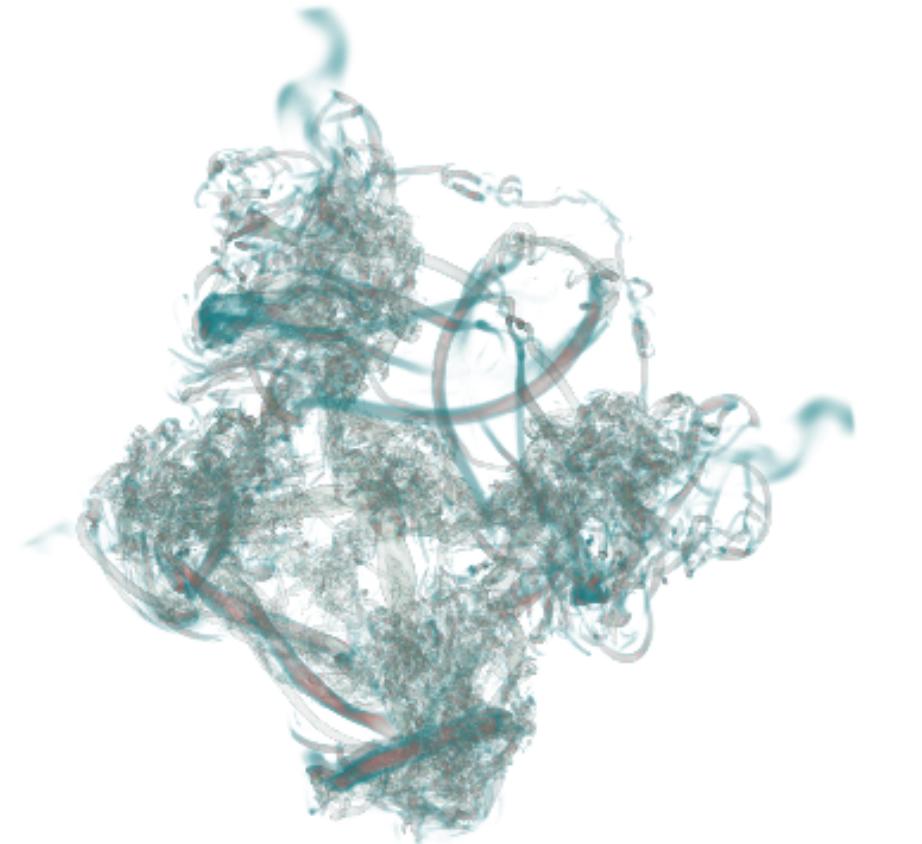
Original VDI  
~40.0 MB  
(lossless enc.)



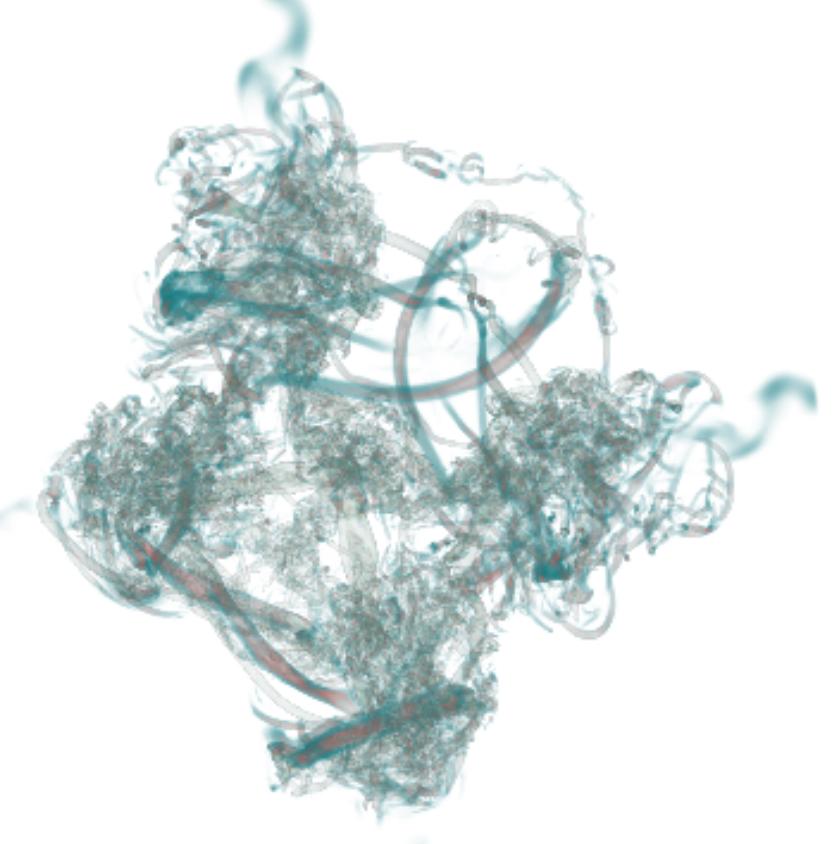
Space-Time VDI  
~5.6 MB  
(lossless enc.)



Original size  
~2130 MB  
(raycasted)



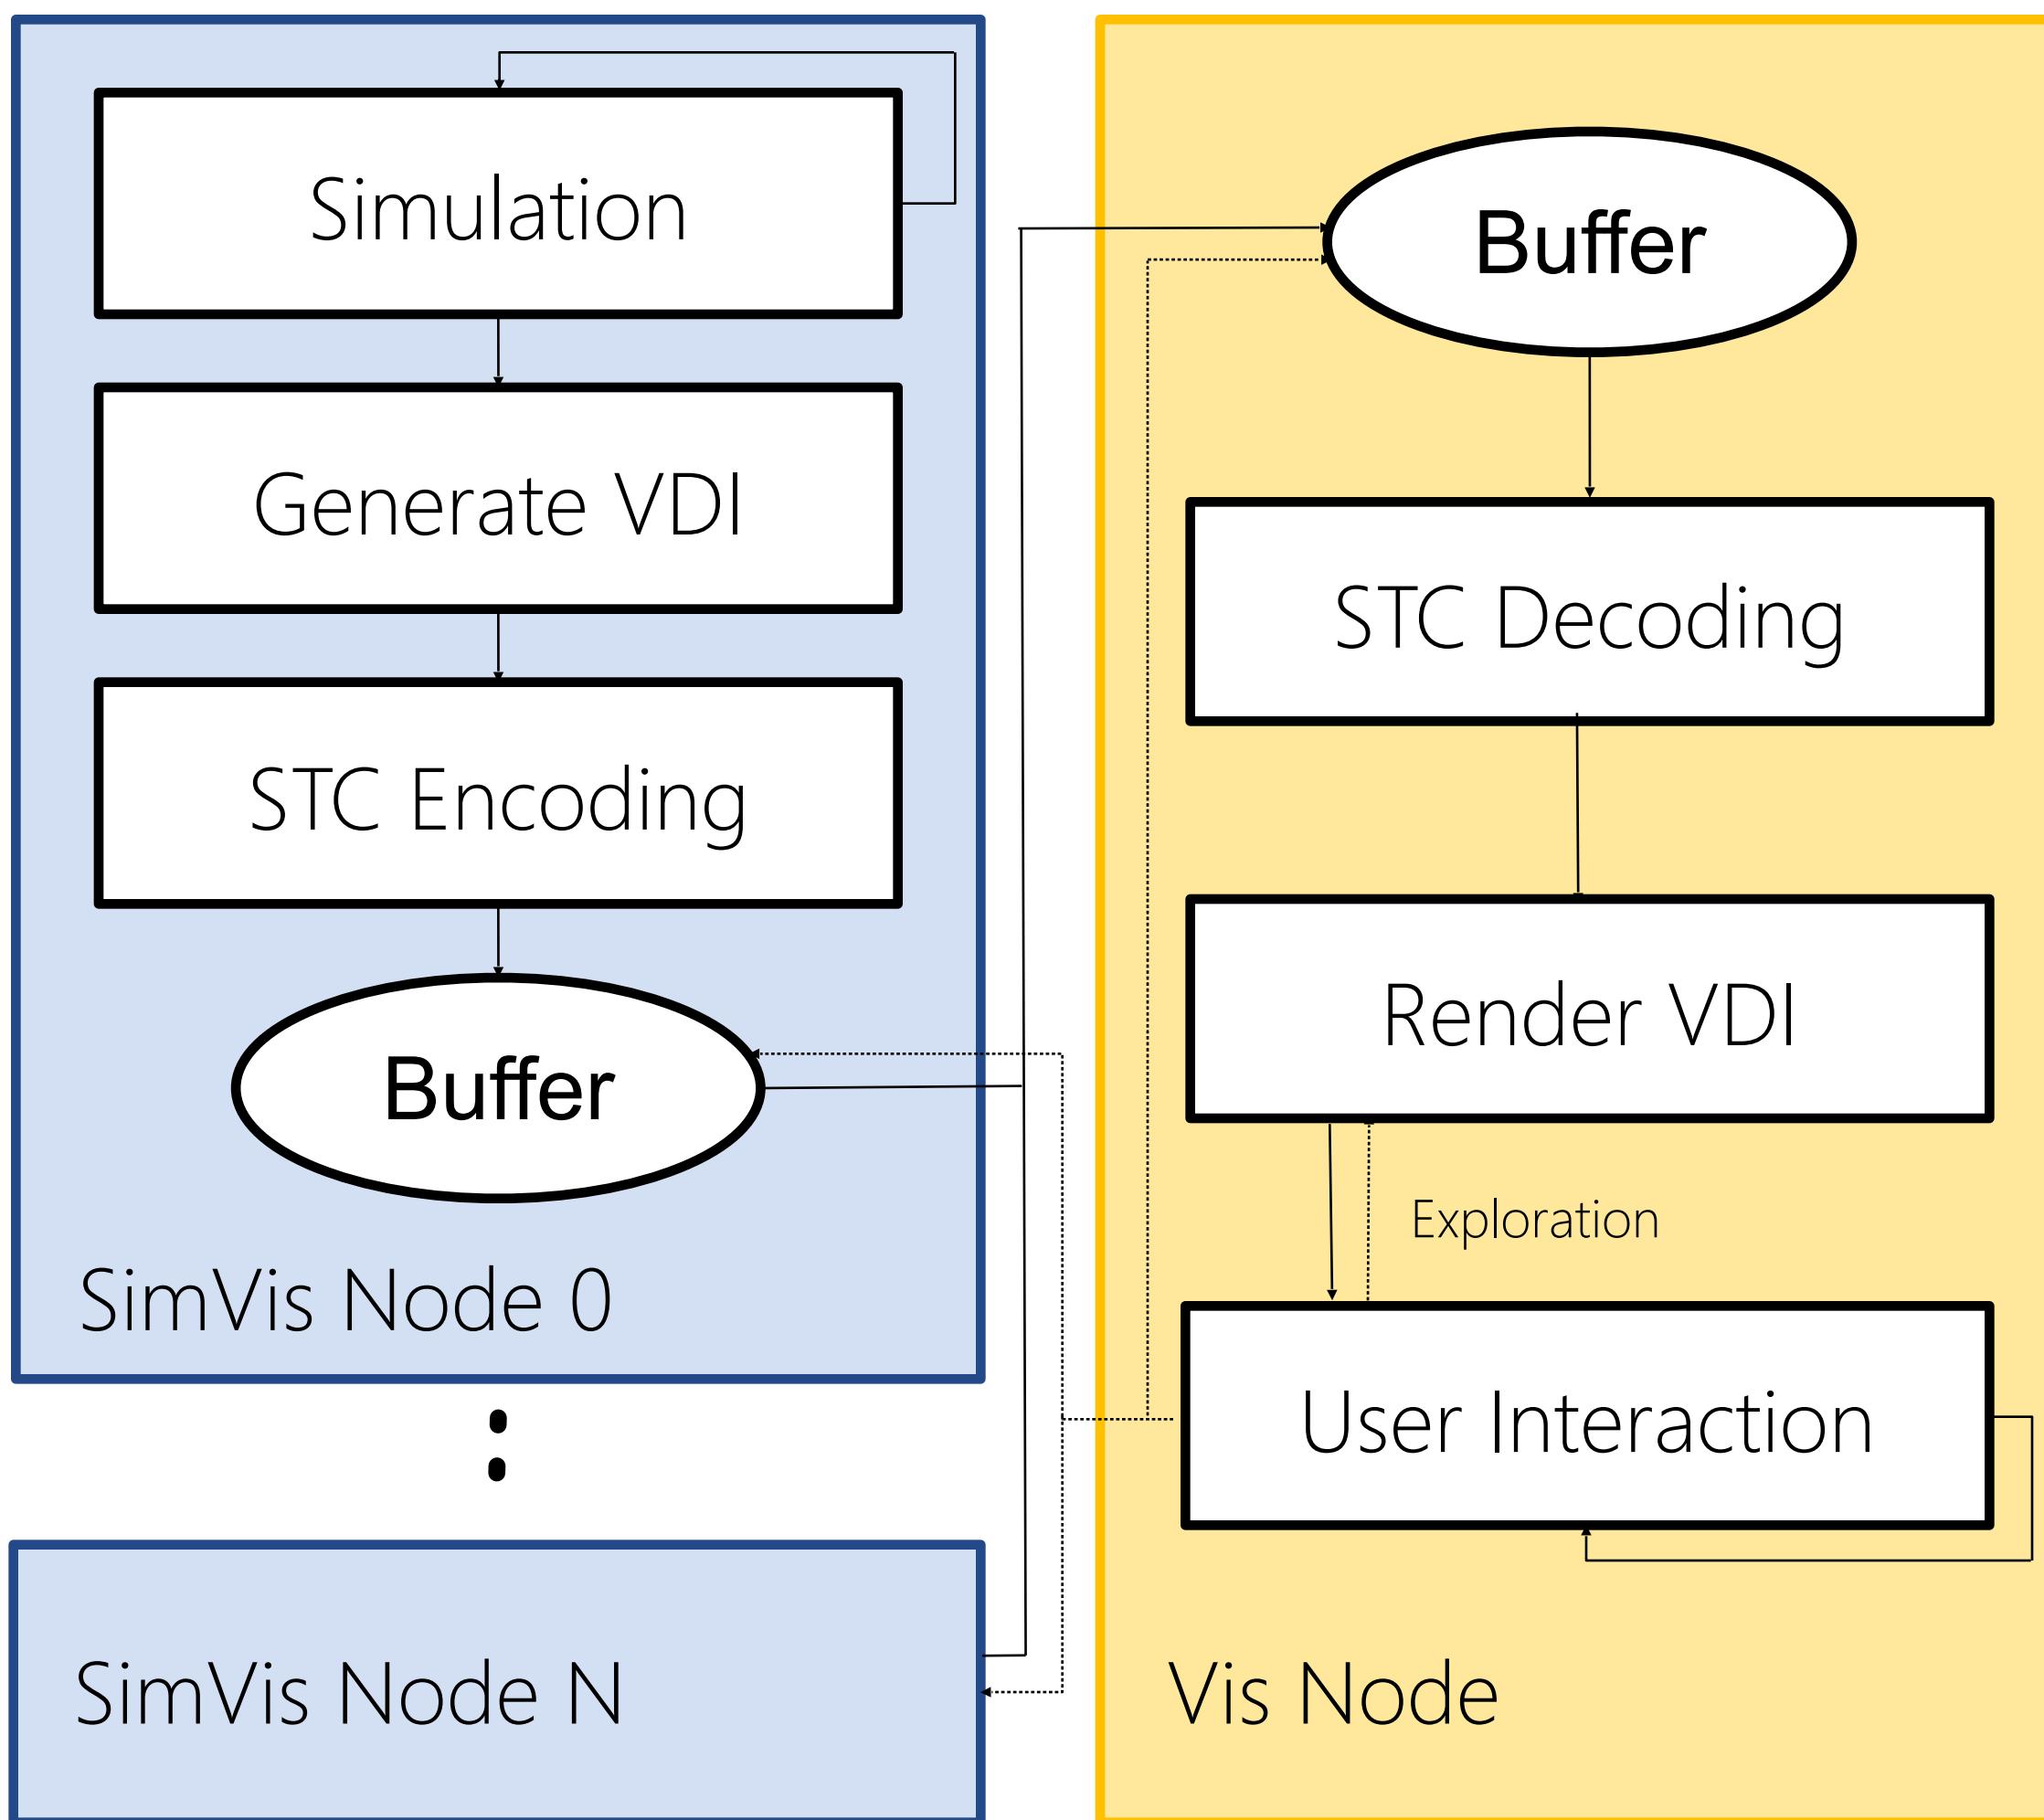
Original VDI  
~375 MB  
(lossless enc.)



Space-Time VDI  
~247 MB  
(lossless enc.)

# Overall Architecture

- simulation-visualization node
  - VDI is generated as soon as simulation generates a time step
  - space-time coding reduces data size
    - continuously updated with new time steps
- visualization node
  - pipeline is executed in reverse order
  - allows for interactive exploration of each time step
  - ... and temporal navigation
  - sliding window spans certain time interval
    - old time steps might have to be discarded if certain limited buffer size is reached



# Conclusion

- in-situ visualization aims to
  - ... minimize the cost for transfer and storage
  - ... with a low computational overhead
  - ... keep the data of interest
  - ... maintain the ability of a user for a posteriori exploration
  - ... be available with minimum delay
- hybrid coupling in-situ architecture can achieve this
  - ... but requires the development of reduced visualization representations
  - research-question and application-specific
  - discussed some of our work in the area
- ongoing challenges
  - simulation-visualization coupling and resource sharing considering software architecture and runtime characteristics of simulation
  - interdisciplinary choice and development of visualization representation

